

# Discovery of Approximate Lexicographical Order Dependencies

Yifeng Jin, Zijiang Tan<sup>✉</sup>, Jixuan Chen, and Shuai Ma<sup>✉</sup>

**Abstract**—Lexicographical order dependencies (LODs) specify orders between list of attributes, and are proven useful in optimizing SQL queries with order by clauses. To discover hidden dependencies from dirty data in practice, approximate dependency discoveries are actively studied, aiming at automatically discovering dependencies that hold on data with some exceptions. In this paper we study the discovery of approximate LODs. (1) We adapt two error measures, namely  $g_1$  and  $g_3$ , to LODs. We prove their desirable properties, present efficient algorithms for computing the measures and related lower and upper bounds, and study the relationship between the two measures. (2) We present an efficient approximate LOD discovery algorithm that is well suited to the two error measures, with a set of pruning rules, optimization techniques and ranking functions. (3) We study techniques for estimating  $g_1$  by sampling, with high accuracy and far less time. (4) We conduct extensive experiments to verify the effectiveness and scalability of our methods, using both real-life and synthetic data.

**Index Terms**—Data profiling, data dependency, metadata

## 1 INTRODUCTION

LEXICOGRAPHICAL order dependencies (LODs) are proposed in [30], [32], which state *lexicographical* ordering specifications. Different from traditional dependencies that are defined on *sets* of attributes, e.g., functional dependencies (FDs) and denial constraints (DCs) [4], LODs are defined on *lists* of attributes. LODs lend themselves to wide applicability, since sorting is one of the most important database operations, and sorting in the *SQL order-by* clause is lexicographical on lists of attributes. Indeed, LODs have been proven useful in query optimizations concerning sorting [30], [31], [32].

**Example 1:** Consider Table 1, concerning employees in a company (now suppose  $t_3.Salary = 4500$  and  $t_6.Salary = 8500$ ). We see an employee with a higher salary is at a higher level, or stays at the same level for more years. This is denoted by  $Salary \mapsto LevelYear$  in the notation of LODs [30], [32]: the ascending order on *Salary* guarantees the ascending order on *Level*, and the ascending order on *Year* within each *Level* group. The lexicographical ordering specification on the left-hand-side (LHS) (resp. right-hand-side (RHS)) of the LOD is consistent with the *SQL* clause *ORDER BY Salary ASC* (resp. *Level ASC, Year ASC*).

- Yifeng Jin, Zijiang Tan, and Jixuan Chen are with the School of Computer Science, Fudan University, Shanghai 200433, China. E-mail: {18210240103, zjtan, 20210240341}@fudan.edu.cn.
- Shuai Ma is with SKLSDE Lab, School of Computer Science and Engineering, Beihang University, Beijing 100191, China. E-mail: mashuai@buaa.edu.cn.

Manuscript received 22 May 2021; revised 10 Oct. 2021; accepted 21 Nov. 2021. Date of publication 23 Nov. 2021; date of current version 7 Mar. 2023. This work was supported in part by the National Key R&D Program of China under Grant 2018YFB1700403 and in part by the NSFC under Grants 62172102, 61572135, and 61925203. (Corresponding author: Zijiang Tan.) Recommended for acceptance by F. Rusu. Digital Object Identifier no. 10.1109/TKDE.2021.3130227

LODs are defined on *lists of attributes*, and can have multiple attributes on both LHS and RHS. Observe the following unique features of LODs.

(1) The order of attributes in a list is relevant. For example,  $Salary \mapsto YearLevel$  does not hold on  $r$ .

(2) The attributes in the LHS and RHS list usually *cannot* be separated. For example,  $Salary \mapsto Year$  does not hold on  $r$ .

To avoid the error-prone and labor-intensive process of designing dependencies manually, dependency discovery techniques are actively studied; see e.g., [1], [27] for surveys. Recently, LOD discoveries have received an increasing attention [5], [11], [17]. The LOD discovery problem is necessarily very difficult. It has a search space (the total number of candidate LODs) factorial in the number of attributes since LODs concern lists of attributes, in contrast to the search space of FDs that is exponential in the number of attributes.

Data in practice are often dirty, and hence, some of the discovered (exact) dependencies *cannot* correctly express the characteristics of data. It is known that discovered constraints on dirty data may *overfit* [16], [21], [23]. This means too many LHS attributes used in FDs, or too many predicates specified in DCs. Intuitively, this is because the discovered constraints have to be more “specialized” to “tolerate” errors in data. When it comes to LOD discovery, the problem is even more involved. We find LODs discovered from dirty data can not only overfit, but also *underfit*. By *underfit*, we mean too few attributes are specified on the RHS, which implies that the order specification is not fully established. Indeed, the following example shows that a single LOD can overfit on the LHS, and simultaneously, underfit on the RHS.

**Example 2:** Recall Table 1. Now suppose that there are errors in the data:  $t_3.Salary = 5800$ ,  $t_6.Salary = 8000$ .

We see  $Salary \mapsto LevelYear$  no longer holds. Specifically, (a)  $t_4$  is before  $t_3$  by *Salary ASC*, but  $t_3$  is before  $t_4$  by *Level ASC, Year ASC*. (b) The order of  $t_5, t_6$  is *unspecified* by *Salary ASC*, but  $t_5$  is before  $t_6$  by *Level ASC, Year ASC*.

TABLE 1  
Relation Instance  $r$

	Name	Salary	Level	Year	Age
$t_1$	Alan	4200	1	2	30
$t_2$	Mark	4300	1	3	31
$t_3$	Jack	4500→5800	2	2	38
$t_4$	William	5600	2	3	30
$t_5$	Steven	8000	3	2	35
$t_6$	Thomas	8500→8000	4	5	39

Consequently, an exact LOD discovery algorithm may find  $\overrightarrow{\text{SalaryAge}} \mapsto \overrightarrow{\text{Level}}$  that holds on the dirty data. Compared with the original (desirable) LOD  $\overrightarrow{\text{Salary}} \mapsto \overrightarrow{\text{LevelYear}}$ , it *overfits* on the LHS and *underfits* on the RHS. The attribute *Year* on the RHS is removed for resolving the violation incurred by  $t_3, t_4$ , and the attribute *Age* is included for resolving the violation caused by  $t_5, t_6$ .

Dirty data in practice motivate the quest for discovering *approximate* dependencies that hold on data with some exceptions. Although desirable, approximate dependency discovery is usually more challenging and expensive than the exact counterpart. Intuitively, exact dependency discoveries concern the *decision* problem of whether a dependency holds or not, while approximate dependency discoveries concern the *counting* problem of measuring the error rate of a dependency. The increasing complexity is well demonstrated in recent studies for approximate FDs [16] and DCs [21], [23].

**Contributions.** In this paper, we study the discovery of approximate lexicographical order dependencies.

(1) We adapt two error measures to LODs. For a LOD, measure  $g_1$  concerns the number of violating tuple pairs, while  $g_3$  concerns the minimum number of tuples that must be removed such that the LOD is satisfied. We show that both measures have desirable properties and can be efficiently computed. We study the lower and upper bounds for the measures, to enable efficient pruning in approximate LOD discoveries. We also investigate the relationship between the two measures. (Section 5).

(2) We provide an algorithm to discover the complete set of minimal and valid approximate LODs. It traverses the search space of candidate approximate LODs, and computes the error measures and lower/upper bounds of the candidates. A set of pruning rules, optimization techniques and ranking functions are also incorporated into the algorithm, for improving effectiveness and efficiency. (Section 6).

(3) We study techniques for estimating  $g_1$  with high accuracy, by sampling a small portion of tuple pairs. Our sampling strategy is efficient and widely applicable, since the sample size is independent of the instance size, and no assumption is made about data distributions. (Section 7).

(4) Using a host of real-life and synthetic data, we conduct an extensive experimental evaluation. The results demonstrate the effectiveness of approximate LOD discovery in recalling correct LODs from dirty data, and the effectiveness of pruning rules, optimizations, ranking functions and estimation by sampling. (Section 8).

## 2 RELATED WORK

This work extends our conference version [10]. (1) We have added (a) the study of the relationship between measures  $g_1$  and  $g_3$  (Section 5.3), (b) functions to rank the discovered approximate LODs (Section 6.3), (c) the investigation of the estimation of error measure  $g_1$  by sampling (Section 7), and (d) experimental evaluations concerning the ranking functions and error measure estimation (Exp-6 and Exp-7 in Section 8). (2) We have presented more theoretical results underlying our approach (Propositions 4, 7, 8, 9, and 11). (3) We have also provided the details of *state compaction* used in Algorithm 5 (Section 6.2).

**Theoretical Foundations of Order Dependencies.** Unidirectional and bidirectional LODs are proposed in [30], [32], and proven useful in optimizing queries with *order-by* clauses [31]. Different from the list-based LODs, two classes of *set-based* order dependencies (ODs) are also discussed. Set-based *canonical* ODs are proposed in [28], [29]. They are shown to generalize LODs, in the sense that each LOD can be mapped to a set of canonical ODs and any relation instance satisfies the LOD iff it satisfies all canonical ODs in the set. Another set-based ODs [7], [8], known as *pointwise* ODs, further generalize canonical ODs. There is no one-to-one relationship between a set-based OD and a LOD, and to our best knowledge, no techniques exist for transforming set-based ODs to LODs.

**Exact Order Dependency Discoveries.** There has been an increasing interest in OD discovery techniques. They are studied in [5], [11], [17] for LODs, and in [28], [29], [33] for set-based ODs, aiming to automatically find ODs that hold on the data without exceptions.

In this paper, we study the discovery of approximate LODs, which significantly differs from prior works on exact LOD discovery. Exact LOD discovery algorithms test the satisfaction of each candidate LOD, and a LOD does not hold if a single violation is identified. In contrast, approximate LOD discoveries need to measure the error rate for each candidate approximate LOD. We adapt two error measures to LODs, with both theoretical analyses and efficient computation methods. We also present a discovery algorithm that is well suited to the error measures, with novel pruning rules and optimization techniques for improving efficiency.

It is highly non-trivial, if not impossible, to extend exact LOD discovery methods to approximate LODs. [5] is based on the observation that each LOD can be divided into an FD and an *order compatibility dependency* (OCD) (refer to Section 5), and the LOD holds iff both FD and OCD hold. This does not apply to approximate LODs. Intuitively, we do not have the “embedded” FD and OCD in a LOD both hold with exceptions if the LOD holds with exceptions. [11] is experimentally verified to be very efficient, by first discovering LODs on a small subset of data, and then refining LODs on full data in an iterative way. The rationale behind [11] is that any exact LOD that holds on data must hold on any subset of it. This does not apply to approximate LODs. Indeed, the error rate of a LOD may increase after removing some tuples from data, and hence, an approximate LOD valid on full data may be invalid on subsets of it.

**Approximate Dependency Discoveries.** To cope with dirty data in practice, approximate dependency discoveries are studied for, e.g., FDs [16], CFDs [24], DCs [21], [23] and

set-based canonical ODs [12], [29]. Based on the definition of approximation using notions from information theory, [14] studies implication for approximate dependencies, and discoveries of approximate multi-valued dependencies and then acyclic schemes are investigated in [13].

In this paper, we study the discovery of approximate LODs. The computation of error measures significantly depends on the dependency types, and a completely different strategy is required to generate list-based candidate approximate LODs compared with those set-based dependencies.

A different notion, referred to as *approximate band conditional* LOD, is proposed in [18], [19]. Band LODs relax themselves to hold approximately and conditionally on subsets of data. [18], [19] are complementary to this work. For an approximate LOD discovered by us, the method of [18], [19] can be employed to split the data instance into contiguous segments such that the LOD holds on each segment.

### 3 PRELIMINARIES

In this section, we review the definition of bidirectional lexicographical order dependencies [30], [32].

**Basic Notations.**  $R(A, \dots)$  denotes a relation schema,  $r$  denotes an instance of  $R$ , and  $t, s$  denote tuples of  $r$ . We use *marked attribute*, written as  $\bar{A}$ , to model the order specifications.  $\bar{A}$  is either  $\overline{A}$  or  $\underline{A}$ , for  $A$  asc or  $A$  desc respectively.  $t_A$  denotes the value of attribute  $A$  in  $t$ , and  $t_{\bar{A}} = t_A$ .

**Attribute List.**  $X$  denotes a list of marked attributes, i.e.,  $[A_1, \dots, A_k]$ , and  $\mathcal{X}$  denotes the set of attributes (without directions) in  $X$ . Given a tuple  $t$ ,  $t_X$  denotes the list of attribute values on  $X$ , i.e.,  $[t_{A_1}, \dots, t_{A_k}]$ .

A non-empty list  $X$  can be denoted as  $[\bar{A}_i | Y]$ , where *head*  $\bar{A}_i$  is a single marked attribute, and *tail*  $Y$  is the remaining list. For  $X = [\bar{A}_1, \dots, \bar{A}_k]$ ,  $prefix(X)$  denotes the set of all possible prefixes of  $X$ , i.e.,  $[\bar{A}_1, \dots, \bar{A}_i]$  for any  $i < k$ .

**Lexicographical Ordering.** For a marked attribute  $\bar{A}$  and tuples  $t, s$ , we write  $t \prec_{\bar{A}} s$  iff (a)  $\bar{A} = \overline{A}$  and  $t_A < s_A$ ; or (b)  $\bar{A} = \underline{A}$  and  $t_A > s_A$ . We write  $t =_{\bar{A}} s$  iff  $t_A = s_A$ .

Given  $X = [A_1, \dots, A_k]$ , we write  $t \preceq_X s$  iff (a)  $X = []$ ; or (b)  $t \prec_{A_1} s$ ; or (c)  $X = [A_1 | Y]$  such that  $t =_{A_1} s$  and  $t \preceq_Y s$ . We write  $t =_X s$  iff  $t \preceq_X s$  and  $s \preceq_X t$ , i.e.,  $t =_{A_i} s$  for all  $i \in [1, k]$ . We write  $t \prec_X s$  iff  $t \preceq_X s$  but  $s \not\preceq_X t$ .

**Bidirectional Lexicographical Order Dependency** [32]. Given two lists  $X, Y$ ,  $\gamma = X \mapsto Y$  denotes a *bidirectional lexicographical order dependency* (abbreviated as LOD). A relation instance  $r$  satisfies  $\gamma$  iff for any two tuples  $t, s \in r$ ,  $t \preceq_Y s$  if  $t \preceq_X s$ . If  $r$  satisfies  $\gamma$ , then we say  $\gamma$  holds on  $r$ .

**Remarks.** (1) Along the same setting as [11], [17], in the sequel we consider *completely non-trivial* LODs whose LHS and RHS attribute lists (neglecting direction) are disjoint. (2) Each LOD  $\gamma$  has a *symmetry* LOD  $\gamma'$  by reversing all directions [11]. For example,  $\overline{A} \mapsto \overline{B} \overline{C}$  is the symmetry of  $\overline{A} \mapsto \overline{B} \overline{C}$ :  $\overline{A}$  (resp.  $\overline{B} \overline{C}$ ) is the reverse order of  $\overline{A}$  (resp.  $\overline{B} \overline{C}$ ), and hence,  $\overline{A} \mapsto \overline{B} \overline{C}$  holds iff  $\overline{A} \mapsto \overline{B} \overline{C}$  holds. W.l.o.g., we consider LODs with asc on the leftmost attribute in the RHS attribute list, e.g.,  $\overline{A} \mapsto \overline{B} \overline{C}$ .

### 4 PROBLEM FORMULATION

In this section we give the definition of approximate LODs, and formalize the approximate LOD discovery problem.

**Error measures.** We use function  $g$  to measure the error rate of LODs. Specifically,  $g(\gamma, r)$  takes as inputs a LOD  $\gamma$  and a relation instance  $r$ . The smaller  $g$  value is, the fewer errors *w.r.t.*  $\gamma$  are on  $r$ . We first present four criteria for judging whether an error measure function  $g$  makes sense, and will study the details of two error functions in Section 5.

(1)  $g(\gamma, r)$  is in the range of  $[0, 1]$  for any LOD  $\gamma$  on any relation instance  $r$ , and  $g(\gamma, r) = 0$  iff  $\gamma$  holds on  $r$ .

(2)  $g(XA \mapsto Y, r) \leq g(X \mapsto Y, r)$ : appending an attribute to the LHS never leads to more errors.

(3)  $g(X \mapsto Y, r) \leq g(X \mapsto YY', r)$ : appending an attribute to the RHS never removes any errors.

(4)  $g(\gamma, r)$  can be efficiently computed, e.g., in polynomial time. This is necessary for a practical setting.

Observe that criteria (2), (3) are consistent with the implication of exact LODs. It is proven in [30], [32] that  $X \mapsto Y$  logically implies  $XA \mapsto Y$ , and  $X \mapsto YY'$  logically implies  $X \mapsto Y$ . Recall that a dependency  $\delta$  logically implies  $\gamma$  in the sense that every instance that satisfies  $\delta$  must satisfy  $\gamma$ .

**Approximate LOD.** Given an error measure function  $g$  and an error threshold  $\epsilon$ , we say that a LOD  $\gamma$  is an *approximate* LOD (abbreviated as AOD) valid on  $r$  iff  $g(\gamma, r) \leq \epsilon$ .

Note that the introduction of a given threshold is a typical setting adopted in approximate dependency discoveries, e.g., [12], [16], [21], [23], [24], [29].

It is better to discover *minimal* valid dependencies rather than all valid ones [16], [21], [22], [23]. In the sequel we establish the minimality of AODs. Intuitively, an attribute list  $X$  is *not minimal* if part of it already imposes the same ordering specification. Inspired by the *reduce order* procedure in [25], we have the following result, and hence the definition of minimal attribute list.

**Lemma 1.** For a list  $X$ , a marked attribute  $\bar{B}$ , a subset  $\mathcal{Y} \subseteq \mathcal{X}$  and two tuples  $t, s$ , (a)  $t \prec_{X\bar{B}} s$  if  $t \prec_X s$ , and (b) if  $FD \mathcal{Y} \rightarrow B$  holds, then  $t =_{X\bar{B}} s$  if  $t =_X s$ .

**Minimal Attribute List.** An attribute list  $X$  is *minimal*, iff there do not exist (a) a subset  $\mathcal{Y}$  of  $\mathcal{X}$  and (b) an attribute  $B$  in  $X$  that is after all attributes in  $\mathcal{Y}$ , where  $\mathcal{Y} \rightarrow B$  holds.

**Example 3.** If  $\overline{AB} \rightarrow C$  holds, then we know neither of  $\overline{A} \overline{B} \overline{C}$ ,  $\overline{B} \overline{D} \overline{A} \overline{C}$ ,  $\overline{A} \overline{B} \overline{C}$  is minimal; we have the same ordering specification after removing  $\overline{C}$  ( $\overline{C}$ ). Note that the LHS attributes of the FD are not required to be a *sublist* (continuous), and directions of attributes are irrelevant.  $\square$

**Implication of AODs.** As noted earlier, an error measure function  $g$  should guarantee that  $g(XA \mapsto Y, r) \leq g(X \mapsto Y, r)$  and  $g(X \mapsto Y, r) \leq g(X \mapsto YY', r)$ . Hence, on any instance  $r$ ,  $XA \mapsto Y$  is a valid AOD if  $X \mapsto Y$  is a valid AOD, and  $X \mapsto Y$  is a valid AOD if  $X \mapsto YY'$  is a valid AOD.

We are now ready to define *minimal AODs*.

**Minimal AODs.** An AOD  $X \mapsto Y$  is *minimal*, iff

(1)  $X$  and  $Y$  are minimal attribute lists; and

(2) for any  $X' \in prefix(X)$ ,  $X' \mapsto Y$  is not a valid AOD; and

(3) for any non-empty list  $Y'$ ,  $X \mapsto YY'$  is not a valid AOD.

**Example 4.** If  $\overline{A} \mapsto \overline{B} \overline{C}$  is valid, then  $\overline{A} \mapsto \overline{B}$  is not minimal.  $\overline{A} \overline{C} \mapsto \overline{B}$  is not minimal if  $\overline{A} \mapsto \overline{B}$  is valid.  $\overline{A} \overline{B} \mapsto \overline{C}$  is not minimal if  $A \rightarrow B$  holds, since  $\overline{A} \overline{B}$  is not minimal.

In the definition of minimal AODs, rule (1) specifies the requirement of minimal attribute lists, and rules (2), (3) follow the principle of logical implication. All AODs satisfying

the threshold are considered valid, and minimality check is conducted on valid AODs to return only minimal ones.

*Discovery of AODs.* Given a relational instance  $r$ , an error measure function  $g$  and a threshold  $e$ , AOD discovery is to find the complete set of minimal valid AODs on  $r$ .

## 5 ERROR MEASURES FOR LODS

In this section, we adapt two error measures to LODs. We show they satisfy the four criteria stated in Section 4, by providing theoretical results and efficient algorithms. We investigate the lower and upper bounds for these measures, to enable early pruning in AOD discoveries. We also study the relationship between the two measures.

### 5.1 The Percentage of Violating Tuple Pairs

The most common error measure, referred to as  $g_1$ , is introduced for FDs [15], [16], and extended to e.g., DCs [4], [21], [23]. The computation of  $g_1$  is closely related to violations of a dependency. We review violations of LODs first.

*LOD violations* [30], [32]. Violations of a LOD  $\gamma = X \mapsto Y$  are categorized into two types: split and swap.

(1) A tuple pair  $(t, s)$  incurs a split, if  $t =_X s, t \neq_Y s$ .

(2) A tuple pair  $(t, s)$  incurs a swap, if  $t \prec_X s, s \prec_Y t$ .

$X \mapsto Y$  has an “embedded” FD  $\mathcal{X} \rightarrow \mathcal{Y}$ , and  $(t, s)$  incurs a split *w.r.t.*  $X \mapsto Y$  iff  $t, s$  violate  $\mathcal{X} \rightarrow \mathcal{Y}$ . In contrast, a swap is caused by a *swapped* tuple pair  $(t, s)$ , i.e.,  $t$  is before  $s$  if sorted by  $X$ , but  $s$  is before  $t$  if sorted by  $Y$ . This is formalized by *order compatibility dependencies* (OCDs) [5].

**Example 5:** Recall  $\text{Salary} \mapsto \text{LevelYear}$  on the dirty instance  $r$  in Example 2.  $(t_3, t_4)$  is a *swapped* tuple pair, and hence incurs a swap.  $(t_5, t_6)$  incurs a split:  $t_5, t_6$  violate FD  $\text{Salary} \rightarrow \text{Level, Year}$ , since they have the same value on *Salary* but different values on *Level* and *Year*.  $\square$

*Error measure  $g_1$ .* The  $g_1$  is the ratio of the number of violating tuple pairs to the number of all tuple pairs [15]. By considering split and swap, we extend  $g_1$  to LODs.

$$g_{\text{split}}(X \mapsto Y, r) = \frac{|\{(t, s) \in r^2 \mid t =_X s \wedge t \neq_Y s\}|}{|r|^2 - |r|}$$

$$g_{\text{swap}}(X \mapsto Y, r) = \frac{|\{(t, s) \in r^2 \mid t \prec_X s \wedge s \prec_Y t\}|}{|r|^2 - |r|}$$

$$g_1(X \mapsto Y, r) = g_{\text{split}}(X \mapsto Y, r) + 2 \times g_{\text{swap}}(X \mapsto Y, r)$$

Herein,  $|r|$  denotes the number of tuples in  $r$ . We scale up  $g_{\text{swap}}$  to balance out the fact that split is *symmetric*, but swap is *asymmetric*:  $(t, s)$  causes a split iff  $(s, t)$  causes a split, but  $(t, s)$  *does not* cause a swap if  $(s, t)$  causes a swap.

It can be verified that  $g_1(\gamma, r)$  ranges over  $[0, 1]$ , and  $g_1(\gamma, r) = 0$  iff  $\gamma$  holds on  $r$ . The following proposition shows that  $g_1$  satisfies criteria (2), (3) stated in Section 4.

**Proposition 2.** (1)  $g_1(XA \mapsto Y, r) \leq g_1(X \mapsto Y, r)$ , and (2)  $g_1(X \mapsto Y, r) \leq g_1(X \mapsto YY', r)$ .

**Proof.** (1) We prove  $g_1(XA \mapsto Y, r) \leq g_1(X \mapsto Y, r)$  as follows.

(a) Suppose  $(t, s)$  causes a split *w.r.t.*  $XA \mapsto Y$ , i.e.,  $t =_{XA} s, t \neq_Y s$ . It is easy to see that  $(t, s)$  also causes a split *w.r.t.*  $X \mapsto Y$ .

Authorized licensed use limited to: BEIHANG UNIVERSITY. Downloaded on February 06, 2025 at 08:20:50 UTC from IEEE Xplore. Restrictions apply.

TABLE 2  
Relation Instance  $r$

	A	B	C	D	E	F	G	H
$t_1$	1	2	1	1	1	1	1	2
$t_2$	1	2	1	1	1	1	2	3
$t_3$	1	4	3	1	3	2	1	3
$t_4$	1	1	3	2	2	3	1	4
$t_5$	2	5	5	2	4	4	1	5
$t_6$	2	6	5	3	5	5	1	1

$\mapsto Y$ . (b) Suppose  $(t, s)$  causes a swap *w.r.t.*  $XA \mapsto Y$ , i.e.,  $t \prec_{XA} s, s \prec_Y t$ . We see that both  $(t, s)$  and  $(s, t)$  lead to a split *w.r.t.*  $X \mapsto Y$  if  $t =_X s$ , or that  $(t, s)$  causes a swap *w.r.t.*  $X \mapsto Y$  if  $t \prec_X s$ .

(2) We then prove  $g_1(X \mapsto Y, r) \leq g_1(X \mapsto YY', r)$ . We can see that  $(t, s)$  always causes a split (resp. swap) *w.r.t.*  $X \mapsto YY'$  if it causes a split (resp. swap) *w.r.t.*  $X \mapsto Y$ .  $\square$

We show  $g_1$  can be efficiently computed by developing such algorithms. We first review an auxiliary data structure.

*Sorted Partition.* The data structure, referred to as *sorted partition*, is employed in exact LOD discoveries [11], [17]. Given an attribute list  $X$ , the sorted partition  $\tau_X$  on an instance  $r$  is a sorted list of *equivalence classes* (sets). Specifically, tuples with the same value on  $X$  are put into the same equivalence class, and for tuples  $t, s$  with different values on  $X$ , the equivalence class of  $t$  is before that of  $s$  if  $t \prec_X s$ . It is known that a sorted partition is built in  $O(|r| \log(|r|))$  on  $r$ .

We use  $|\tau_X|$  to denote the number (count) of equivalence classes in  $\tau_X$ , and define the *rank* of a tuple  $t$  in  $\tau_X$  as the sequence number of the equivalence class that  $t$  belongs to, denoted by  $I_X[t]$ . Intuitively,  $I_X[t]$  denotes the order of  $t$  on  $X$  in a compact way.

**Example 6:** Consider Table 2.  $\tau_{\vec{A}} = [\{t_1, t_2, t_3, t_4\}, \{t_5, t_6\}]$ ,  $\tau_{\vec{B}} = [\{t_4\}, \{t_1, t_2\}, \{t_3\}, \{t_5\}, \{t_6\}]$ .  $I_{\vec{A}}[t_3] = 1$ ;  $t_3$  is in the first equivalence class of  $\tau_{\vec{A}}$ . Similarly,  $I_{\vec{B}}[t_3] = 3$ .

Observe that a split is always incurred by two tuples in the same equivalence class, while two tuples in different equivalence classes can only lead to a swap. We present two algorithms for computing  $g_{\text{split}}$  and  $g_{\text{swap}}$  respectively.

#### Algorithm 1. Computation of $g_{\text{split}}(X \mapsto Y, r)$

**Input:** sorted partitions  $\tau_X, \tau_Y$

**Output:**  $g_{\text{split}}(X \mapsto Y, r)$

```

1:  $spl \leftarrow 0$ ;
2: foreach equivalence class  $ec$  in  $\tau_X$  do
3:    $map \leftarrow$  an empty hash table;
4:   foreach tuple  $t$  in  $ec$  do
5:      $map[I_Y[t]] \leftarrow map[I_Y[t]] + 1$ ;
6:   foreach  $v$  in the value set of  $map$  do
7:      $spl \leftarrow spl + v \times (|ec| - v)$ ;
8: return  $g_{\text{split}} = \frac{spl}{|r|^2 - |r|}$ ;

```

*Algorithm.* Algorithm 1 is provided for computing  $g_{\text{split}}$ . A tuple pair  $(t, s)$  incurs a split *w.r.t.*  $X \mapsto Y$ , if  $t, s$  have the same values on  $X$  but different values on  $Y$ . Hence, in each equivalence class  $ec$  of  $\tau_X$ , we count the number of tuples




Fig. 1. Example 7 for Algorithm 1 and Example 9 for Algorithm 2.

for each distinct value on  $Y$  using a hash table with  $I_Y[t]$  as the key (lines 3-5). A value  $v$  in the hash table implies  $v$  tuples having the same value on  $Y$ , and each of these tuples forms a split with any tuple from the other  $|ec| - v$  tuples, i.e., tuples in the same equivalence class of  $\tau_X$  but in different equivalence classes of  $\tau_Y$  (lines 6-7). Herein,  $|ec|$  denotes the number of tuples in the equivalence class  $ec$ .

**Example 7:** Consider  $\vec{A} \mapsto \vec{B}$  on Table 2. We show the related split violations in Fig. 1a. (1) In the first equivalence class  $\{t_1, t_2, t_3, t_4\}$  of  $\tau_{\vec{A}}$ ,  $I_{\vec{B}}[t_1] = I_{\vec{B}}[t_2] = 2$ ,  $I_{\vec{B}}[t_3] = 3$ , and  $I_{\vec{B}}[t_4] = 1$ . On the hash table, we have  $map[2] = 2$ ,  $map[3] = map[1] = 1$ . The number of violating tuple pairs is computed as  $2 \times (4-2) + 1 \times (4-1) + 1 \times (4-1) = 10$ . (2) Similarly, we then deal with the second equivalence class  $\{t_5, t_6\}$  of  $\tau_{\vec{A}}$ .  $\square$

**Time Complexity.** Assuming a constant cost for the hash table, Algorithm 1 has a complexity of  $O(|r|)$  on  $\tau_X$  and  $\tau_Y$ .

**Lower and upper bounds of  $g_1$ .** As will be illustrated in Section 6.1, when traversing the search space of candidate AODs, new candidates following  $X \mapsto Y$  are in the form of  $XU \mapsto YV$  ( $U$  or  $V$  can be empty). It is beneficial if we can obtain a lower bound and an upper bound of  $g_1$  for  $XU \mapsto YV$ . This is because (1)  $XU \mapsto YV$  cannot be a valid AOD and can be pruned, if its lower bound of  $g_1$  is larger than the error threshold  $e$ ; and (2)  $XU \mapsto YV$  is always a valid AOD if its upper bound of  $g_1$  is not larger than  $e$ .

The bounds can be efficiently computed together with  $g_{\text{swap}}$  based on the following observations.

**Proposition 3.** [17], [32] If  $(t, s)$  incurs a swap w.r.t.  $X \mapsto Y$ , then  $(t, s)$  also incurs a swap w.r.t.  $XU \mapsto YV$ .

**Proposition 4.** If  $t \prec_X s$  and  $t \prec_Y s$ , then  $(t, s)$  never incurs any LOD violations w.r.t.  $XU \mapsto YV$ .

**Proof.** We can see that  $(t, s)$  does not incur a split w.r.t.  $XU \mapsto YV$  since  $t \neq_{XU} s$ , and that  $(t, s)$  does not incur a swap w.r.t.  $XU \mapsto YV$ , since  $t \prec_{XU} s$  and  $t \prec_{YV} s$ .  $\square$

We refer to a tuple pair  $(t, s)$  as an *ordered pair* if  $t \prec_X s$  and  $t \prec_Y s$ , and refer to  $(t, s)$  as a *swapped pair* if  $t \prec_X s$  and  $s \prec_Y t$ . Intuitively, when appending more attributes to the LHS and RHS attribute lists, Proposition 3 states that swap violations can never be resolved, while Proposition 4 tells us that ordered pairs still never lead to violations. We hence compute the bounds as follows

$$LB_{g_1}(XU \mapsto YV, r) = 2 \times g_{\text{swap}}(X \mapsto Y, r)$$

$$UB_{g_1}(XU \mapsto YV, r) = 1 - 2 \times \frac{|\{(t, s) \in r^2 \mid t \prec_X s \wedge t \prec_Y s\}|}{|r|^2 - |r|}$$

The example below shows that both bounds are tight.

**Example 8:** Recall Table 2. For  $\vec{F} \mapsto \vec{H}$ , it can be seen that there are 5 swapped pairs and 8 ordered pairs. We have the

lower and upper bounds are  $10/30$  and  $14/30$  respectively. It can be verified that  $\vec{F}\vec{G} \mapsto \vec{H}$  has  $g_1$  of  $10/30$  and  $\vec{F} \mapsto \vec{H}\vec{G}$  has  $g_1$  of  $14/30$ .  $\square$

We present one auxiliary structure to facilitate our algorithm for computing  $g_{\text{swap}}$  and the number of ordered pairs.

**Segment Tree.** We employ a simple yet effective data structure, known as segment tree [2]. A segment tree supports various range queries, e.g., range sum/max/min queries, and it takes  $O(n)$  to build and  $O(\log(n))$  to update and query a segment tree built on range  $[1, n]$  (integer values). A segment tree also has a space complexity of  $O(n)$ . As will be seen shortly, the ranks of tuples in a sorted partition are well suited for transforming operations on the sorted partition to the corresponding segment tree.

**Algorithm 2.** Computation of  $g_{\text{swap}}$  and the number of ordered pairs

**Input:** sorted partitions  $\tau_X$  and  $\tau_Y$

**Output:**  $g_{\text{swap}}$  and the number of ordered pairs for  $X \mapsto Y$

- 1:  $\text{swap} \leftarrow 0, \text{ordered} \leftarrow 0$ ;
- 2:  $\text{seg} \leftarrow$  an empty segment tree on range  $[1, |\tau_Y|]$ ;
- 3: **foreach** equivalence class  $ec$  in  $\tau_X$  **do**
- 4:   **foreach** tuple  $t$  in  $ec$  **do**
- 5:      $\text{swap} \leftarrow \text{swap} + \text{seg.query}([I_Y[t] + 1, |\tau_Y|])$ ;
- 6:      $\text{ordered} \leftarrow \text{ordered} + \text{seg.query}([1, I_Y[t] - 1])$ ;
- 7:   **foreach** tuple  $t$  in  $ec$  **do**
- 8:      $\text{seg.insert}(I_Y[t])$ ;
- 9: **return**  $g_{\text{swap}} = \frac{\text{swap}}{|r|^2 - |r|}, \text{ordered}$ ;

**Algorithm.** Algorithm 2 aims to compute  $g_{\text{swap}}$  and the number of ordered pairs simultaneously. The application of a segment tree built on range  $[1, |\tau_Y|]$  is at the core of this algorithm. On this tree,  $\text{insert}(\text{key})$  increases the value associated with  $\text{key}$  by 1, for counting the number of tuples in the same equivalence class of  $\tau_Y$  (line 8), and  $\text{query}([a, b])$  performs a range sum query on  $[a, b]$ , for the total number of tuples in several equivalence classes (lines 5-6).

Equivalence classes  $ec$  in  $\tau_X$  are processed one by one in order (line 3), and updates of the segment tree with tuples in  $ec$  (line 8) are conducted after the queries concerning these tuples (lines 5-6). Therefore, the number of swap violations w.r.t. a tuple  $t$  is the number of tuples whose rank in  $\tau_Y$  is larger than that of  $t$  (line 5), and the number of ordered pairs w.r.t.  $t$  is the number of tuples whose rank in  $\tau_Y$  is smaller than that of  $t$  (line 6). Recall that a larger (resp. smaller) rank in  $\tau_Y$  implies a larger (resp. smaller) value on  $Y$ .

**Example 9:** Consider  $\vec{C} \mapsto \vec{B}$  on Table 2 (shown in Fig. 1b). (1) The first equivalence class  $\{t_1, t_2\}$  of  $\tau_{\vec{C}}$  does not lead to swapped or ordered pairs. After  $t_1, t_2$  are inserted into the segment tree, there are two tuples whose rank in  $\tau_{\vec{B}}$  is 2. (2) In the second equivalence class of  $\tau_{\vec{C}}$ ,  $t_3$  leads to 2 ordered pairs (solid lines) since the rank of  $t_1, t_2$  is smaller than  $I_{\vec{B}}[t_3] = 3$ , while  $t_4$  incurs 2 swapped pairs (dashed lines) since the rank of  $t_1, t_2$  is larger than  $I_{\vec{B}}[t_4] = 1$ . The segment tree is employed to facilitate an efficient range sum query. We then update the tree with  $t_3, t_4$ . (3)  $\{t_5, t_6\}$  of  $\tau_{\vec{C}}$  is processed similarly, which leads to more ordered pairs.

**Time Complexity.** The segment tree has a range of  $[1, |\tau_Y|]$ , and  $|\tau_Y|$  equals  $|r|$  in the worst case. It hence takes at most  $O(|r|)$  to build and  $O(\log(|r|))$  to update and query the tree.



TABLE 3  
Relation Instance  $r$

	$A$	$B$	$C$	$D$
$t_1$	1	2	1	3
$t_2$	1	1	1	4
$t_3$	1	3	2	4
$t_4$	1	4	2	3
$t_5$	2	2	1	1
$t_6$	2	1	1	2
$t_7$	2	3	2	2
$t_8$	2	4	2	1

The update and query are conducted for each tuple once. Algorithm 2 has a worst-case complexity of  $O(|r|\log(|r|))$ .

## 5.2 The Minimum Number of Removed Tuples

Another error measure, referred to as  $g_3$  in literature, is also originally introduced for FDs [15]. It is extended to CFDs [24], set-based canonical LODs [12], [29], comparable dependencies [26] and DCs [21], among others. The computation of  $g_3$  can be very expensive, e.g., *NP-Complete* for comparable dependencies [26] and DCs [21].

*Error measure  $g_3$  for LODs.* The  $g_3$  measures the minimum number of tuples that must be removed from the given instance such that the dependency is satisfied. Specifically,

$$g_3(\gamma, r) = \frac{|r| - \max\{|r'| \mid r' \subseteq r, r' \text{ satisfies } \gamma\}}{|r|}$$

Obviously, we have  $g_3(\gamma, r) \in [0, 1]$ , and  $g_3(\gamma, r) = 0$  iff  $\gamma$  holds on  $r$ . We then show how to compute  $g_3$ , for illustrating the satisfaction of the criteria in Section 4. We present one more definition to facilitate our approach.

*OD sequence.* For a LOD  $\gamma = X \mapsto Y$ , an *OD sequence* on an instance  $r$  is a list of tuples  $t_{a_1}, t_{a_2}, \dots, t_{a_k}$  from  $r$ , such that for any  $1 \leq i < j \leq k$ , (a)  $t_{a_i} \preceq_X t_{a_j}$ , (b)  $t_{a_i} \preceq_Y t_{a_j}$ , and (c)  $t_{a_i} =_Y t_{a_j}$  if  $t_{a_i} =_X t_{a_j}$ . It is easy to see that any two tuples in this sequence *cannot* form a LOD violation. We say an OD sequence is the *longest OD sequence*, denoted by  $\text{LOS}(\gamma, r)$ , whose  $k$  is the maximum among all OD sequences (choose an arbitrary one if several ones have the same  $k$ ).

**Example 10:** Consider Table 3. For  $\vec{A} \mapsto \vec{C}$ , we have a LOS  $[t_1, t_2, t_5, t_6]$ . There can be more than one LOS, e.g.,  $[t_3, t_4, t_7, t_8]$  or  $[t_1, t_2, t_7, t_8]$  is also a LOS.  $\square$

We show  $g_3$  can be readily computed from LOS.

**Proposition 5.**  $g_3(\gamma, r) = 1 - \frac{|\text{LOS}(\gamma, r)|}{|r|}$ .

**Proof.** Suppose that  $g_3$  is computed based on a maximum subset of  $r$  that satisfies  $\gamma$ , say  $r'$ . Let  $\gamma = X \mapsto Y$ . We can order all tuples in  $r'$  by  $X$ , to form a sequence. It can be verified that this is an OD sequence, since any two tuples in  $r'$  do not incur a violation. The sequence is also the longest. If we have any longer OD sequence, then all tuples in it form a subset larger than  $r'$ , which contradicts the assumption.  $\square$

*Remark.* LOS differs from the *longest increasing sequence* in e.g., [6], [9], [20]. This is because LOD violations consist of both swap and split. In an OD sequence, we must have (a)

$t_{a_i} \preceq_Y t_{a_j}$  if  $t_{a_i} \preceq_X t_{a_j}$ , similar to the longest increasing sequence, and (b)  $t_{a_i} =_Y t_{a_j}$  if  $t_{a_i} =_X t_{a_j}$ , which is unique.

We show the monotonicity of  $g_3$ , based on Proposition 5.

**Proposition 6.** (1)  $g_3(XA \mapsto Y, r) \leq g_3(X \mapsto Y, r)$ , and (2)  $g_3(X \mapsto Y, r) \leq g_3(X \mapsto YY', r)$ .

**Proof.** (1) We prove  $g_3(XA \mapsto Y, r) \leq g_3(X \mapsto Y, r)$  by showing that  $\text{LOS}(X \mapsto Y, r)$  can always be transformed into an OD sequence (not necessarily the longest) for  $XA \mapsto Y$ . Let  $\text{LOS}(X \mapsto Y, r) = t_{a_1}, t_{a_2}, \dots, t_{a_k}$ , where for any  $1 \leq i < j \leq k$ , (a)  $t_{a_i} \preceq_X t_{a_j}$ , (b)  $t_{a_i} \preceq_Y t_{a_j}$  and (c)  $t_{a_i} =_Y t_{a_j}$  if  $t_{a_i} =_X t_{a_j}$ . The transformation is done by a linear scan of  $\text{LOS}(X \mapsto Y, r)$ , and we exchange the position of  $t_{a_i}$  and  $t_{a_{i+1}}$ , if  $t_{a_i} =_X t_{a_{i+1}}$  (hence  $t_{a_i} =_Y t_{a_{i+1}}$ ) and  $t_{a_{i+1}} \prec_{XA} t_{a_i}$ . It can be verified that the new sequence must be an OD sequence for  $XA \mapsto Y$ .

(2) We show  $\text{LOS}(X \mapsto YY', r)$  is also an OD sequence (not necessarily the longest) for  $X \mapsto Y$ . Suppose  $t_{a_1}, t_{a_2}, \dots, t_{a_k}$  is a LOS for  $X \mapsto YY'$ . For  $t_{a_i}$  and  $t_{a_j}$  where  $1 \leq i < j \leq k$ : (a) If  $t_{a_i} =_X t_{a_j}$ , then we have  $t_{a_i} =_{Y'} t_{a_j}$  by the definition of LOS, and hence  $t_{a_i} =_Y t_{a_j}$ . (b) If  $t_{a_i} \prec_X t_{a_j}$ , then we have  $t_{a_i} \preceq_{Y'} t_{a_j}$  by the definition of LOS. We must have  $t_{a_i} \preceq_Y t_{a_j}$ , otherwise  $t_{a_i} \succ_{YB} t_{a_j}$  if  $t_{a_i} \succ_Y t_{a_j}$ , which leads to a contradiction. Putting (a) and (b) together, we know  $t_{a_1}, t_{a_2}, \dots, t_{a_k}$  is an OD sequence for  $X \mapsto Y$ .

Recall that LOS is the longest OD sequence and  $g_3 = 1 - \frac{|\text{LOS}|}{|r|}$ . We have  $g_3(XA \mapsto Y, r) \leq g_3(X \mapsto Y, r)$ , and  $g_3(X \mapsto Y, r) \leq g_3(X \mapsto YY', r)$ .  $\square$

**Example 11:** Recall Table 3. (1)  $[t_1, t_2, t_5, t_6]$  is a LOS for  $\vec{A} \mapsto \vec{C}$ . After being transformed into  $[t_2, t_1, t_6, t_5]$ , it is an OD sequence (but not LOS) for  $\vec{A} \vec{B} \mapsto \vec{C}$ . We see one LOS for  $\vec{A} \vec{B} \mapsto \vec{C}$  is  $[t_2, t_1, t_6, t_5, t_7, t_8]$ . (2)  $[t_1, t_8]$  is a LOS for  $\vec{A} \mapsto \vec{C} \vec{D}$ , and is also an OD sequence (but not LOS) for  $\vec{A} \mapsto \vec{C}$ .

Similar to  $g_1$ , we aim for the lower and upper bounds of  $g_3$  for  $XU \mapsto YV$ . We find that they can be computed with two other sequences that slightly differ from LOS.

*Strict increasing sequence and non-decreasing sequence.* For  $\gamma = X \mapsto Y$ , (1) a *strict increasing sequence* on  $r$  is a list of tuples  $t_{a_1}, \dots, t_{a_k}$ , such that for any  $1 \leq i < j \leq k$ , (a)  $t_{a_i} \prec_X t_{a_j}$ , and (b)  $t_{a_i} \prec_Y t_{a_j}$ . The *longest strict increasing sequence*, denoted by  $\text{LSIS}(\gamma, r)$ , is a strict increasing sequence with the maximum  $k$ .

(2) A *non-decreasing sequence* is a list of tuples  $t_{a_1}, \dots, t_{a_k}$ , such that for any  $1 \leq i < j \leq k$ , (a)  $t_{a_i} \preceq_X t_{a_j}$ , and (b)  $t_{a_i} \preceq_Y t_{a_j}$ . The *longest non-decreasing sequence*, denoted by  $\text{LNDS}(\gamma, r)$ , is a non-decreasing sequence with the maximum  $k$ .

We have the following results.

**Proposition 7.** For a LOD  $\gamma$  and a relation instance  $r$ , a *strict increasing sequence* must be an OD sequence, and an OD sequence must be a *non-decreasing sequence*.

**Proof.** These can be directly seen by the definitions.  $\square$

**Proposition 8.** Any two tuples in  $\text{LSIS}(X \mapsto Y, r)$  do not form a violation (split or swap) w.r.t.  $XU \mapsto YV$ .

**Proof.** For  $t_{a_i}, t_{a_j} \in \text{LSIS}(X \mapsto Y, r)$ , we know (a)  $t_{a_i} \neq_{XU} t_{a_j}$ ; and (b)  $t_{a_i} \prec_{YV} t_{a_j}$  if  $t_{a_i} \prec_{XU} t_{a_j}$ . This guarantees no split or swap violations w.r.t.  $XU \mapsto YV$ .  $\square$

**Proposition 9.** *Tuples in  $\text{LNDS}(X \mapsto Y, r)$  form a subset  $r'$  of  $r$  such that (a) any two tuples in  $r'$  do not incur a swap w.r.t.  $X \mapsto Y$ ; and (b)  $r'$  is maximum among all subsets of  $r$  that satisfy (a).*

**Proof.** Obviously any two tuples in  $\text{LNDS}(X \mapsto Y, r)$  do not incur a swap w.r.t.  $X \mapsto Y$ . Let  $r'$  be the set of all tuples in  $\text{LNDS}(X \mapsto Y, r)$ . If we can find a tuple  $s \in r$  such that (a)  $s \notin r'$ ; and (b)  $s$  does not incur a swap w.r.t.  $X \mapsto Y$  against any tuple in  $r'$ , then  $s$  can be inserted into  $\text{LNDS}(X \mapsto Y, r)$ , contradicting the assumption that  $\text{LNDS}$  is the longest.  $\square$

In summary, (1)  $\text{LSIS}(X \mapsto Y, r)$  must be an OD sequence for  $XU \mapsto YV$  on  $r$  (not necessarily the longest), and hence can serve as a lower bound of  $\text{LOS}(XU \mapsto YV, r)$ . (2)  $\text{LNDS}(X \mapsto Y, r)$  corresponds to a maximum subset of  $r$  that is free of swap violations w.r.t.  $X \mapsto Y$ . Recall that  $XU \mapsto YV$  can never resolve swap w.r.t.  $X \mapsto Y$ . Therefore,  $\text{LNDS}(X \mapsto Y, r)$  is an upper bound of  $\text{LOS}(XU \mapsto YV, r)$ .

The following example shows that both bounds are tight.

**Example 12:** Recall Table 3. For  $A \mapsto C$ , we have a  $\text{LSIS}$   $[t_1, t_8]$  and a  $\text{LNDS}$   $[t_1, t_2, t_5, t_6, t_7, t_8]$ , which results in a lower bound of 2 and an upper bound of 6 for  $\text{LOS}(AU \mapsto CV, r)$ . It can be seen that  $A \mapsto CD$  has a  $\text{LOS}$   $[t_1, t_8]$ , and  $AB \mapsto C$  has a  $\text{LOS}$   $[t_2, t_1, t_6, t_5, t_7, t_8]$ .

Based on the observation, we are ready to define the lower and upper bounds of  $g_3$  for  $XU \mapsto YV$ .

$$LB_{g_3}(XU \mapsto YV, r) = 1 - \frac{|\text{LNDS}(X \mapsto Y, r)|}{|r|}$$

$$UB_{g_3}(XU \mapsto YV, r) = 1 - \frac{|\text{LSIS}(X \mapsto Y, r)|}{|r|}$$

### Algorithm 3. Computations of LOS, LSIS and LNDS

---

**Input:** sorted partitions  $\tau_X, \tau_Y$   
**Output:**  $\text{LOS}(X \mapsto Y, r)$ ,  $\text{LSIS}(X \mapsto Y, r)$ ,  $\text{LNDS}(X \mapsto Y, r)$

- 1:  $los \leftarrow$  an empty segment tree on range  $[1, |\tau_Y|]$ ;
- 2:  $lsis \leftarrow$  an empty segment tree on range  $[1, |\tau_Y|]$ ;
- 3:  $lnds \leftarrow$  an empty segment tree on range  $[1, |\tau_Y|]$ ;
- 4:  $\tau_{XY} \leftarrow \tau_X.\text{expand}(\tau_Y)$ ;
- 5: **foreach** equivalence class  $ecx$  in  $\tau_X$  **do**
- 6:   **foreach** equivalence class  $ecxy$  in  $\tau_{XY}$  from  $ecx$  **do**
- 7:      $iy \leftarrow I_Y[ecxy[1]]$ ;
- 8:      $ecxy.\text{los} \leftarrow |ecxy| + los.\text{query}([1, iy])$ ;
- 9:      $ecxy.\text{lsis} \leftarrow 1 + lsis.\text{query}([1, iy - 1])$ ;
- 10:     $ecxy.\text{lnds} \leftarrow |ecxy| + lnds.\text{query}([1, iy])$ ;
- 11:     $lnds.\text{insert}(iy, ecxy.\text{lnds})$ ;
- 12:   **foreach** equivalence class  $ecxy$  in  $\tau_{XY}$  from  $ecx$  **do**
- 13:      $iy \leftarrow I_Y[ecxy[1]]$ ;
- 14:      $los.\text{insert}(iy, ecxy.\text{los})$ ;
- 15:      $lsis.\text{insert}(iy, ecxy.\text{lsis})$ ;
- 16: **return**  $los.\text{query}([1, |\tau_Y|])$ ,  $lsis.\text{query}([1, |\tau_Y|])$ , and  $lnds.\text{query}([1, |\tau_Y|])$ ;

---

**Algorithm.** Algorithm 3 is a three-in-one approach to computing LOS, LSIS and LNDS, and we use a segment tree for each of them (lines 1-3). The segment trees here have different operation semantics from those in Algorithm 2, but the complexity of each operation remains unchanged. Specifically,  $\text{insert}(x, y)$  updates the value associated with key  $x$  to  $y$ , and  $\text{query}([a, b])$  returns the  $\max$  value associated

with keys in the range of  $[a, b]$ . We use segment trees to facilitate our computation in the dynamic programming fashion. The tuple rank in  $\tau_Y$  is used as the key for querying trees, and the value is the length of the longest sequence (LOS, LSIS or LNDS) that ends with that key (rank).

We first compute  $\tau_{XY}$  with  $\tau_X$  and  $\tau_Y$  (line 4). This is a basic operation on sorted partitions [17]. One equivalence class in  $\tau_X$  may be divided into several equivalence classes in  $\tau_{XY}$ , such that tuples in the same equivalence class of  $\tau_{XY}$  have the same values on both  $X$  and  $Y$ .

The outer loop (line 5) enumerates equivalence classes  $ecx$  in  $\tau_X$ , ordered by  $X$ 's value. The first inner loop (lines 6-11) enumerates equivalence classes  $ecxy$  in  $\tau_{XY}$  that are obtained from  $ecx$  (with the same  $X$ 's value), in the order of  $Y$ 's value. We then identify the rank related to  $ecxy$  in  $\tau_Y$  (all tuples in  $ecxy$  have the same value on  $Y$ ), denoted by  $iy$  in the algorithm (line 7). We use  $iy$  as the key for querying and updating segment trees. Note that the  $Y$ 's values related to keys (equivalence classes in  $\tau_Y$ ) in the range of  $[1, iy - 1]$  are less than the  $Y$ 's value related to  $iy$ .

Consider the computation of LOS. The LOS that ends with  $ecxy$  is obtained by appending all tuples in  $ecxy$  to the longest LOS that ends with an equivalence class (a) already inserted into the tree (a smaller value on  $X$ ) and (b) having a  $Y$ 's value not larger than that of  $ecxy$  (line 8). We save the length of the LOS for  $ecxy$  (line 8), and update the tree with it in the second inner loop (line 14). This is necessary since the same value on  $Y$  is required for the same value on  $X$  in LOS; a different equivalence class  $ecxy'$  from the same  $ecx$  cannot be combined with  $ecxy$ . In contrast, we update the tree for LNDS immediately (line 11). This is because  $ecxy'$  can contribute to the LNDS of  $ecxy$  as long as the  $Y$ 's value of  $ecxy'$  is not larger than that of  $ecxy$ . The computation of LSIS differs in the following. Only one (arbitrary) tuple in  $ecxy$  can be appended to LSIS and the range query is conducted on  $[1, iy - 1]$  (line 9), since an equal value on  $X$  or  $Y$  is not allowed in LSIS.

Finally, we get LOS, LNDS and LSIS, by querying  $\max$  values from their corresponding trees (line 16).

**Example 13:** Consider  $\bar{D} \mapsto \bar{E}$  on Table 2. In Fig. 2, we illustrate the process of Algorithm 3, by showing the (length of) LNDS, LOS or LSIS that ends with the rank  $iy$ . We have  $\tau_{\bar{D}} = [\{t_1, t_2, t_3\}, \{t_4, t_5\}, \{t_6\}]$  and  $\tau_{\bar{D} \mapsto \bar{E}} = [\{t_1, t_2\}, \{t_3\}, \{t_4\}, \{t_5\}, \{t_6\}]$ . In the outer loop, we enumerate equivalence classes in  $\tau_{\bar{D}}$ .

(1) We deal with  $\{t_1, t_2, t_3\}$ . In the inner loop, we enumerate equivalence classes in  $\tau_{\bar{D} \mapsto \bar{E}}$  that are from  $\{t_1, t_2, t_3\}$ , i.e.,  $\{t_1, t_2\}$  and  $\{t_3\}$ . (a)  $t_1, t_2$  are in the same equivalence class of  $\tau_{\bar{D} \mapsto \bar{E}}$ ; they hence exist (or do not exist) simultaneously in any LOS (resp. LNDS), but only one of them can exist in any LSIS. (b)  $t_3$  can be appended to  $\{t_1, t_2\}$  in LNDS, but not in LOS or LSIS.

(2) We handle  $\{t_4, t_5\}$  in  $\tau_{\bar{D}}$ , which is divided into  $\{t_4\}$  and  $\{t_5\}$  in  $\tau_{\bar{D} \mapsto \bar{E}}$ . (a)  $\{t_4\}$  can be appended to  $\{t_1, t_2\}$  in LOS, LNDS and LSIS. (b) For LNDS,  $\{t_5\}$  can be appended to  $\{t_1, t_2\}$ ,  $\{t_3\}$  or  $\{t_4\}$ ; in Fig. 2 we choose  $\{t_4\}$  for the longest sequence (the same for  $\{t_3\}$ ). For LOS and LSIS,  $\{t_5\}$  can only be appended to  $\{t_1, t_2\}$  or  $\{t_3\}$ . We must choose  $\{t_1, t_2\}$  for LOS, but it is the same to choose  $\{t_1, t_2\}$  or  $\{t_3\}$  for LSIS, since only one tuple in  $\{t_1, t_2\}$  can exist in LSIS.

(3) We deal with  $\{t_6\}$  in  $\tau_{\bar{D}}$  (and  $\tau_{\bar{D} \mapsto \bar{E}}$ ). We must choose  $\{t_5\}$  for LNDS, but it is the same to choose  $\{t_4\}$  or  $\{t_5\}$  for LOS and LSIS.

**Time Complexity.** Each segment tree has a range of  $[1, |\tau_Y|]$ .  $|\tau_Y| = |r|$  in the worst case; it takes at most  $O(|r|)$  to




Fig. 2. Example 13 for Algorithm 3.

build and  $O(\log(|r|))$  to update and query segment trees. The two inner loops are linear in the size of equivalence classes in  $\tau_{XY}$ . To sum up, Algorithm 3 has a worst-case complexity of  $O(|r|\log(|r|))$ .

### 5.3 Relationship Between Measures $g_1$ and $g_3$

Both  $g_1$  and  $g_3$  of a LOD  $\gamma$  on a relation instance  $r$  measure the error rate of  $\gamma$  in  $r$ . However, the values of  $g_1$  and  $g_3$  can differ significantly, as illustrated below.

**Example 14:** Consider an instance  $r = \{t^1, t^2, \dots, t^k\}$  of schema  $R(A, B)$  and an AOD  $\vec{A} \mapsto \vec{B}$ . Suppose  $t^i_A = t^i_B = i$  for all  $i \in [1, k-1]$ , and  $t^k_A = k$ . We consider two cases of  $t^k_B$ : (a)  $t^k_B = k - 1.5$ , or (b)  $t^k_B = 0$ . In both cases, we can remove  $t^k$  from  $r$  such that  $\vec{A} \mapsto \vec{B}$  is satisfied, and hence,  $g_3 = \frac{1}{k}$ . In contrast, we have  $g_1^a = \frac{2}{k(k-1)}$  in case (a), but  $g_1^b = \frac{2}{k}$  in case (b). The ratios of  $g_3$  to  $g_1^a$  and  $g_1^b$  to  $g_1^a$  can be arbitrarily large, as the number  $k$  of tuples increases.  $\square$

Formally, we have the following result.

**Proposition 10.**  $\frac{2g_3(\gamma, r)}{|r|-1} \leq g_1(\gamma, r) \leq 2g_3(\gamma, r)$ .

**Proof.** By the definition of  $g_3$ , we can find a maximum subset  $r'$  of  $r$  with the size of  $(1 - g_3)|r|$ , where any two tuples in  $r'$  do not form a violation. For each tuple from the  $g_3|r|$  tuples that are removed from  $r$ , it violates  $\gamma$  against (a) at most  $|r| - 1$  tuples, and (b) at least one tuple from  $r'$ . Recall that (1) both  $(t, s)$  and  $(s, t)$  are counted in  $g_{\text{split}}$ , and (2)  $g_{\text{swap}}$  is scaled by 2 when computing  $g_1$ . We have

$$\frac{2g_3|r|}{|r|(|r|-1)} \leq g_1 \leq \frac{2g_3|r|(|r|-1)}{|r|(|r|-1)}$$

$$\frac{2g_3}{|r|-1} \leq g_1 \leq 2g_3$$

$\square$

## 6 DISCOVERY OF AODS

In this section, we first present an AOD discovery algorithm that is suitable to both  $g_1$  and  $g_3$ . We then study several optimizations and ranking functions.

### 6.1 Algorithms for AOD Discovery

**Algorithm.** DisAOD (Algorithm 4) discovers the complete set  $\Sigma$  of minimal and valid AODs on a given instance  $r$ , with a given error measure function  $g$  ( $g_1$  or  $g_3$ ) and a threshold  $e$ . DisAOD traverses the search space of AODs by following a depth-first-search (DFS) strategy implemented by recursion. The AOD traversal is organized in a *forest*. Each tree is rooted at an AOD of the form  $\vec{A} \mapsto \vec{B}$ , i.e.,  $\vec{A} \mapsto \vec{B}$  or  $\vec{A} \mapsto \vec{B}$ , where  $B \in R$  and  $A \in R \setminus B$  (lines 2-3). Recall that it suffices to consider AODs with asc on the leftmost attribute on the RHS due to *symmetry* (Section 3).

For each candidate AOD  $X \mapsto Y$ , its error measure  $Vg$  and lower/upper bound  $LBg/UBg$  are computed with the given

function  $g$  (line 8). If  $X \mapsto Y$  is valid, then we add it into  $\Sigma$  (lines 9-10). If the  $UBg$  is larger than the threshold  $e$ , then we further test  $X \mapsto Y\vec{C}$  (both  $X \mapsto Y\vec{C}$  and  $X \mapsto Y\vec{C}$ ) for all  $C \in R \setminus XY$ , by recursively calling function *Search* (lines 11-14). As a prerequisite, we check whether  $Y\vec{C}$  is a minimal attribute list by calling function *MinimalAttributelist* (line 13); candidates with non-minimal attribute lists are directly discarded. By definition (Section 4), it is to check whether  $Y \rightarrow C$ . This is done by checking whether appending  $\vec{C}$  to  $Y$  incurs any changes to  $\tau_Y$ , i.e., whether  $\tau_{Y\vec{C}} = \tau_Y$ . It is easy to prove that  $\tau_{Y\vec{C}} = \tau_Y$  iff  $Y \rightarrow C$ . If  $Y\vec{C}$  is found to be a minimal attribute list, then the checking incurs no extra cost since the computation of  $\tau_{Y\vec{C}}$  is originally required. We will further develop optimization techniques for this in Section 6.2.

If the upper bound  $UBg$  is not larger than the threshold  $e$ , then AODs of the form  $X \mapsto YV$  are all valid. We generate minimal ones among them by calling function *Extend* (line 16). Specifically, this is done by enumerating  $X \mapsto YW$ , where  $W$  is a list on *all* attributes in  $R \setminus XY$  (the others cannot be minimal) and  $YW$  is a minimal attribute list.

If  $X \mapsto Y$  is invalid and  $LBg \leq e$ , then we further consider  $X\vec{C} \mapsto Y$  (both  $X\vec{C} \mapsto Y$  and  $X\vec{C} \mapsto Y$ ) for all  $C \in R \setminus XY$ , if  $X\vec{C}$  is a minimal attribute list (lines 17-21).

As the final step, we remove non-minimal AODs  $X \mapsto Y$  if there exists valid AOD  $X \mapsto YU$  ( $U$  is not empty) in  $\Sigma$ , by calling function *MinimalAOD* (line 4). This is necessary by the definition of minimal AODs (Section 4).

### Algorithm 4. DisAOD

**Input:** a relation  $r$  of schema  $R$ , an error measure function  $g$  and a threshold  $e$

**Output:** the complete set  $\Sigma$  of minimal and valid AODs

```

1:  $\Sigma \leftarrow \emptyset$ ;
2: foreach  $B \in R, A \in R \setminus B$  do
3:   Search( $\vec{A} \mapsto \vec{B}$ );
4:  $\Sigma \leftarrow \text{MinimalAOD}(\Sigma)$ ;
5: return  $\Sigma$ ;
6:
7: Function Search(AOD candidate  $X \mapsto Y$ )
8:  $Vg, LBg, UBg \leftarrow \text{Compute}(X \mapsto Y, g, r)$ ;
9: if  $Vg \leq e$  then
10:   $\Sigma \leftarrow \Sigma \cup \{X \mapsto Y\}$ ;
11:  if  $UBg > e$  then
12:    foreach  $C \in R \setminus XY$  do
13:      if MinimalAttributelist( $Y\vec{C}$ ) then
14:        Search( $X \mapsto Y\vec{C}$ );
15:  else
16:     $\Sigma \leftarrow \Sigma \cup \text{Extend}(X \mapsto Y)$ ;
17: else
18:  if  $LBg \leq e$  then
19:    foreach  $C \in R \setminus XY$  do
20:      if MinimalAttributelist( $X\vec{C}$ ) then
21:        Search( $X\vec{C} \mapsto Y$ );

```



**Example 15:** Consider  $R(A, B, C, D)$ . In the tree rooted at  $\vec{A} \mapsto \vec{B}$ , we consider candidates  $\vec{A} \mapsto \vec{B}\vec{C}$ ,  $\vec{A} \mapsto \vec{B}\vec{D}$  (resp.  $\vec{A}\vec{C} \mapsto \vec{B}$ ,  $\vec{A}\vec{D} \mapsto \vec{B}$ ), if  $\vec{A} \mapsto \vec{B}$  is valid (resp. invalid) and the pruning by lower/upper bounds does not apply to  $\vec{A} \mapsto \vec{B}$ . Before computing the error measure of a candidate, say  $\vec{A} \mapsto \vec{B}\vec{C}$ , we check whether  $\vec{B}\vec{C}$  is a minimal attribute list by calling function *MinimalAttributelist*. It is to check whether  $B \rightarrow C$ . If (a)  $\vec{B}\vec{C}$  is a minimal attribute list, and (b)  $\vec{A} \mapsto \vec{B}\vec{C}$  is valid and the upper bound of error measure is not larger than the threshold  $e$ , then AODs obtained by appending more attributes to  $\vec{B}\vec{C}$  are all valid. Function *Extend* is called to generate minimal ones among them. Specifically, it enumerates candidates by appending all remaining attributes to the RHS, e.g.,  $\vec{A} \mapsto \vec{B}\vec{C}\vec{D}$ , and further checks whether  $\vec{B}\vec{C}\vec{D}$  is a minimal attribute list. If  $\vec{B}\vec{C}\vec{D}$  is minimal, then  $\vec{A} \mapsto \vec{B}\vec{C}\vec{D}$  is a minimal valid AOD.  $\square$

**Complexity.** AODs concern lists of attributes. In the worst-case, the number of all candidate AODs is  $O(|R|!)$ , and it takes  $O(|r|\log(|r|))$  to check each candidate AOD, where  $|R|$  (resp.  $|r|$ ) is the number of attributes (resp. tuples). The actual complexity is usually much smaller (Section 8).

## 6.2 Optimizations

In this subsection, we further develop several optimization techniques for DisAOD.

### Algorithm 5. Incremental Computation of $g_{\text{swap}}$ and the Number of Ordered Tuple Pairs

---

**Input:** sorted partitions  $\tau_X, \tau_Y, \tau_{X\bar{A}}$ , the numbers of swapped and ordered tuple pairs for  $X \mapsto Y$

**Output:**  $g_{\text{swap}}$  and the number of ordered tuple pairs for  $X\bar{A} \mapsto Y$

---

```

1: foreach  $i$  in  $[1, |\tau_X|]$  do
2:    $\text{rank}[i] \leftarrow 0$ ;  $\text{newrank}[i] \leftarrow 0$ ;
3: foreach equivalence class  $ecy$  in  $\tau_Y$  do
4:   foreach tuple  $t$  in  $ecy$  do
5:     if  $I_Y(t) > \text{rank}[I_X(t)]$  then
6:        $\text{newrank}[I_X(t)] \leftarrow \text{newrank}[I_X(t)] + 1$ ;
7:        $\text{rank}[I_X(t)] \leftarrow I_Y(t)$ ;
8:        $m[t] \leftarrow \text{newrank}[I_X(t)]$ ;
9:  $\text{swap}, \text{ordered} \leftarrow$  the numbers of swapped and ordered pairs for  $X \mapsto Y$ ;
10: foreach equivalence class  $ecx$  in  $\tau_X$  do
11:    $\text{range} \leftarrow 0$ ;
12:   foreach tuple  $t$  in  $ecx$  do
13:      $\text{range} \leftarrow \max(\text{range}, m[t])$ ;
14:    $\text{seg} \leftarrow$  an empty segment tree on range  $[1, \text{range}]$ ;
15:   foreach equivalence class  $ecxa$  in  $\tau_{X\bar{A}}$  from  $ecx$  do
16:     foreach tuple  $t$  in  $ecxa$  do
17:        $\text{swap} \leftarrow \text{swap} + \text{seg.query}([m[t] + 1, \text{range}])$ ;
18:        $\text{ordered} \leftarrow \text{ordered} + \text{seg.query}([1, m[t] - 1])$ ;
19:     foreach tuple  $t$  in  $ecxa$  do
20:        $\text{seg.insert}(m[t])$ ;
21: return  $g_{\text{swap}} = \frac{\text{swap}}{|r|^2 - |r|}$ ,  $\text{ordered}$ ;

```

---

**Incremental computations.** Following  $X \mapsto Y$ , we consider new candidate  $X\bar{A} \mapsto Y$  or  $X \mapsto Y\bar{A}$ . Leveraging results of  $X \mapsto Y$ , incremental computations not only apply to sorted partitions, but also to error measure functions. We present an “incremental” version of Algorithm 2 with better efficiency.

Authorized licensed use limited to: BEIHANG UNIVERSITY. Downloaded on February 06, 2025 at 08:20:50 UTC from IEEE Xplore. Restrictions apply.

Recall that a swapped (resp. an ordered) tuple pair *w.r.t.*  $X \mapsto Y$  is still a swapped (resp. an ordered) pair *w.r.t.*  $X\bar{A} \mapsto Y$  or  $X \mapsto Y\bar{A}$ . Hence, the number of swapped (resp. ordered) pairs monotonically increases. Without loss of generality, we consider  $X\bar{A} \mapsto Y$ . We see the following.

**Proposition 11.** If  $t \prec_{X\bar{A}} s$  and  $t \not\prec_X s$ , then  $t =_X s$ .

Proposition 11 tells us that if tuples  $t, s$  form a new swapped or an ordered pair *w.r.t.*  $X\bar{A} \mapsto Y$ , then  $t, s$  must have the same value on  $X$ , i.e., in the same equivalence class of  $\tau_X$ . Therefore, we can leverage  $\tau_X$  that is already computed for  $X \mapsto Y$ , and cope with each equivalence class in  $\tau_X$  separately, when computing the incremental swapped and ordered tuple pairs for  $X\bar{A} \mapsto Y$ .

**Algorithm.** Algorithm 5 incrementally computes  $g_{\text{swap}}$  and the number of ordered tuple pairs for  $X\bar{A} \mapsto Y$ , based on the known numbers of swapped and ordered tuple pairs for  $X \mapsto Y$ . As stated earlier, it handles each equivalence class of  $\tau_X$  one by one (lines 10-21), in contrast to Algorithm 2 that deals with all equivalence classes of  $\tau_X$  as a whole.

A segment tree is leveraged for each equivalence class *ecx*. Recall that the complexity of a segment tree is closely related to the range; it takes  $O(n)$  to build and  $O(\log(n))$  to update and query a segment tree on  $[1, n]$ . We use the idea of *state compaction* to build each segment tree on a compact range. We obtain the *local* rank  $m[t]$  of a tuple  $t$  in *ecx* based on the original rank  $I_Y[t]$ , which preserves the order. As an example, suppose there are 5 tuples in *ecx* whose  $I_Y[t]$  values are 3, 5, 5, 1, 10 respectively, their ranks in *ecx* are 2, 3, 3, 1, 4, respectively.

More specifically, we convert the rank  $I_Y[t]$  to  $m[t]$  as follows (lines 1-8). We use two auxiliary arrays *rank* and *newrank* that are indexed by  $I_X(t)$ , i.e., the  $\tau_X$  rank of tuple  $t$ . For the  $i$ th equivalence class of  $\tau_X$ , denoted by  $\tau_X[i]$ , we use  $\text{rank}[i]$  to store the  $\tau_Y$  rank of the last visited tuple from  $\tau_X[i]$ , and  $\text{newrank}[i]$  to store the *local*  $\tau_Y$  rank of the last visited tuple from  $\tau_X[i]$ . All tuples  $t$  are processed according to their orders in  $\tau_Y$  (lines 3-4), and hence  $I_Y(t)$  can never decrease during the processing. For  $t$  belonging to  $\tau_X[i]$  ( $I_X(t) = i$ ),  $\text{rank}[I_X(t)]$  and  $\text{newrank}[I_X(t)]$  are updated if  $t$  has a larger  $\tau_Y$  rank than the last visited tuple (Line 5-7), and the *local* rank  $m[t]$  is  $\text{newrank}[I_X(t)]$  (line 8).

The range of the segment tree for the equivalence class *ecx* is set to be the max value of the new ranks, i.e., the number of distinct  $I_Y(t)$  values for all  $t$  from *ecx* (lines 11-14), which reduces from  $|\tau_Y|$  to (at most)  $|ecx|$  by the state compaction. Note that the sum of the ranges of all segment trees used in Algorithm 5 is at most the number  $|r|$  of tuples.

**Example 16:** On Table 2,  $\vec{A} \mapsto \vec{B}$  incurs no swapped but 8 ordered pairs (Fig. 3a). Now consider  $\vec{A}\vec{C} \mapsto \vec{B}$ . The equivalence class  $\{t_1, t_2, t_3, t_4\}$  in  $\tau_{\vec{A}\vec{C}}$  is divided into  $\{t_1, t_2\}$  and  $\{t_3, t_4\}$  in  $\tau_{\vec{A}\vec{C}}$ , which incurs two new ordered pairs (solid lines) and two new swapped pairs (dashed lines), shown in Fig. 3b.  $\square$

**Time Complexity.** The state compaction for segment trees (lines 1-8 and 11-13) is done in  $O(|r|)$ , so is the initialization of all segment trees (line 14). Algorithm 5 has the same worst-case complexity as Algorithm 2, but is experimentally verified to be much more efficient in practice.

**Index for checking the attribute list minimality.** The computation of  $\tau_{Y\vec{C}}$  is required for checking whether appending  $\vec{C}$  to  $Y$  leads to a non-minimal attribute list. To avoid some unnecessary computations, we employ an indexing




Fig. 3. Example 16 for Algorithm 5.

structure on  $C$ , for fetching all  $\mathcal{X}$  if we find  $\mathcal{X} \rightarrow C$  in DisAOD. Specifically, we do the following when  $\bar{C}$  is appended to  $Y$ .

(1) If  $\mathcal{Y}$  is a *superset* of any  $\mathcal{X}$  related to  $C$  in the index, then we know  $Y\bar{C}$  is not a minimal attribute list.

(2) Otherwise, we compute  $\tau_{Y\bar{C}}$ . (a) If  $\tau_{Y\bar{C}} = \tau_Y$ , then  $Y\bar{C}$  is not a minimal attribute list since  $\mathcal{Y} \rightarrow C$ . We update the index with  $\mathcal{Y}$ , and also remove any  $\mathcal{X}$  related to the key  $C$  if  $\mathcal{X}$  is a superset of  $\mathcal{Y}$ . (b) If  $\tau_{Y\bar{C}} \neq \tau_Y$ , then  $Y\bar{C}$  is a minimal attribute list. We continue to the next step of DisAOD with the computed sorted partition  $\tau_{Y\bar{C}}$ .

**Sorted Partition Cache.** Sorted partitions are heavily used in AOD discovery. In the traversal, the same attribute list may occur multiple times (possibly) on different sides. For example, we may generate  $\bar{A}\bar{C} \mapsto \bar{B}$  from  $\bar{A} \mapsto \bar{B}$ ,  $\bar{A}\bar{C} \mapsto \bar{D}$  from  $\bar{A} \mapsto \bar{D}$ , and  $\bar{B} \mapsto \bar{A}\bar{C}$  from  $\bar{B} \mapsto \bar{A}$ , all with the list  $\bar{A}\bar{C}$ , and hence, the same sorted partition  $\tau_{\bar{A}\bar{C}}$ . DisAOD adopts a DFS traversal with a small memory footprint, which enables us to maintain a cache for the created sorted partitions. In addition to the sorted partitions necessary for the DFS traversal, we also use free memory to preserve more sorted partitions for possible reuse. We use a simple LRU (least recently used) strategy when the memory is used up.

### 6.3 Ranking Functions

The number of minimal valid AODs can be large on some instances. In this subsection, we present two functions to rank AODs, for helping users quickly select a small set of more relevant AODs. Specifically, (a) *ability to sort* is a new measure proposed for AODs, and (b) *succinctness* is a common measure that is adapted to AODs.

**Ability to sort.** The ability of a list  $X$  in terms of sorting can be judged by the sorted partition  $\tau_X$ . We aim for  $\tau_X$  that is of a large size (many distinct values on  $X$ ) and doesn't have very large equivalence classes (tuple values on  $X$  are evenly distributed). Specifically, we measure the ability of  $X$  as

$$\text{sort}(X) = \frac{\sum_{ec \in \tau_X} |ec|^2}{|r|^2}$$

In the equation,  $|ec|$  denotes the number of tuples in the equivalence class  $ec$  of  $\tau_X$ . It is easy to see that a small value is obtained if tuples are evenly distributed among many equivalence classes. For a LOD  $X \mapsto Y$ , we define its ability to sort as the average of the abilities of its LHS and RHS.

$$\text{sort}(X \mapsto Y) = \text{avg}(\text{sort}(X), \text{sort}(Y))$$

**Succinctness.** Based on the minimum description length principle, succinctness is a common criterion for ranking discovered dependencies [4], [23]. We define the succinctness of  $X \mapsto Y$  as follows and again prefer a small value.

$$\text{succ}(X \mapsto Y) = \frac{|X| + |Y|}{|r|}$$

Note that appending attributes to the LHS (RHS) of an AOD harms the *succinctness*, but favors the *ability to sort*. As an add-on feature of AOD discovery, DisAOD can additionally sort all discovered AODs according to the two functions, for *top-k* AODs. We will experimentally study the effectiveness of ranking functions in Section 8.

## 7 ERROR MEASURE ESTIMATION BY SAMPLING

In this section, we first study a method to estimate error measure  $g_1$  by sampling. Based on the method, we then propose an alternative approach to AOD discovery on a sample of the full instance.

It can be very costly to compute error measures on large instances. As explained in [15], [21] for approximate FD and DC discoveries, a good estimation of error measures suffices in many applications. In this paper we focus on the estimation of  $g_1$ .

**Estimation of  $g_1$ .** Consider a LOD  $\gamma = X \mapsto Y$  on an instance  $r$ . We reorder tuples in  $r$  such that for two tuples  $t_i, t_j \in r$ ,  $t_i \preceq_X t_j$  if  $i < j$ . Recall that a tuple pair  $(t_i, t_j)$  may incur (a) a split only when  $t_i =_X t_j$ , or (b) a swap only when  $t_i \prec_X t_j$ .

Suppose the  $g_1$  value of  $\gamma$  on  $r$  is known. When we randomly sample  $k$  tuple pairs  $(t_i, t_j)$  ( $i < j$ ) from  $r$ , it is to conduct Bernoulli experiments  $k$  times with a probability of  $g_1$  each time, in terms of the number of violations (split or swap). Hence, we know that the number  $|vio|$  of violations in the sampling follows the binomial distribution

$$|vio| \sim B(k, g_1)$$

The  $g_1$  can be easily estimated

$$\hat{g}_1 = \frac{|vio|}{k}$$

In practice we can use normal approximation to the binomial distribution, when  $k$  is big enough

$$|vio| \sim N(kg_1, kg_1(1 - g_1))$$

Our aim is to determine the size  $k$  of the sampled tuple pairs, for a  $\hat{g}_1$  close enough to its true value. Specifically, for a given error threshold  $e$  and two more parameters  $\delta, \alpha$ , we hope the probability that the distance between  $\hat{g}_1$  and the true value  $g_1$  is less than  $\delta e$  is at least  $1 - \alpha$ , i.e.,

$$\Pr(|\hat{g}_1 - g_1| < \delta e) \geq 1 - \alpha$$

By normal approximation, we have the following

$$\begin{aligned} & \Pr(|\hat{g}_1 - g_1| < \delta e) \\ &= \Pr\left(\frac{|\hat{g}_1 - g_1|}{\sqrt{g_1(1 - g_1)/k}} < \frac{\delta e}{\sqrt{g_1(1 - g_1)/k}}\right) \\ &= 2\Phi\left(\frac{\delta e}{\sqrt{g_1(1 - g_1)/k}}\right) - 1 \geq 1 - \alpha \end{aligned}$$

$$\Phi\left(\frac{\delta e}{\sqrt{g_1(1-g_1)/k}}\right) \geq 1 - \frac{\alpha}{2}$$

$$\frac{\delta e}{\sqrt{g_1(1-g_1)/k}} \geq \mu_{1-\frac{\alpha}{2}}$$

$$k \geq \frac{\mu_{1-\frac{\alpha}{2}}^2 g_1(1-g_1)}{\delta^2 e^2}$$

We mainly care about  $g_1$  that is not greater than the given error threshold  $e$  and  $e$  is typically much smaller than 1. By replacing  $(1 - g_1)$  with 1 and  $g_1$  with  $e$  respectively, we have

$$k \geq \frac{\mu_{1-\frac{\alpha}{2}}^2}{\delta^2 e} \quad (1)$$

The value of  $\mu$  for a specific  $\alpha$  in the equation can be obtained by referring to the standard normal distribution table. The theoretical result shows an estimation of high accuracy can be obtained with a  $k$  independent of the size  $|r|$  of the instance  $r$ . For example, using  $k = 54,103$ , we have a probability of at least 98% that the gap between  $\hat{g}_1$  and  $g_1$  is less than 0.001, i.e., in the setting of  $e = 0.01$ ,  $\delta = 0.1$  and  $\alpha = 0.02^1$ . Note that a very small instance of 1,000 tuples already has 999,000/2 tuple pairs available. Better, the sampling strategy is widely applicable, since it is a uniform random sampling without any assumption on data distributions.

**AOD discovery by sampling.** We propose a lightweight approach to AOD discovery with  $g_1$ . Given an instance  $r$ , (1) we randomly choose  $k$  tuple pairs ( $k$  is determined according to Equation 1), and use all tuples from the pairs to form a sample  $r'$ . (2) We run DisAOD with  $g_1$  on  $r'$ .

The goal is to perform AOD discovery on  $r'$  as a substitution for AOD discovery on  $r$ , with high accuracy and far less cost. Our experimental evaluations (Section 8) demonstrate that we can indeed achieve this goal.

## 8 EXPERIMENTAL EVALUATIONS

In this section, we present an in-depth experimental study. We conduct experiments to (1) demonstrate the efficiency of AOD discoveries and optimization techniques, and to (2) verify the effectiveness of AOD discoveries, the ranking functions and the estimation of  $g_1$  by sampling.

### 8.1 Experimental Settings

**Datasets.** We use a set of real-life and synthetic data from <http://metanome.de> that are evaluated in OD discoveries [5], [11], [17], [28], [29]. (1) NCV, FLI, Hepa and Atom are real-life data, concerning voters, flights, hepatitis disease and atom sites, respectively. (2) DB, Letter and FDR are synthetic datasets with complicated attribute relationships. We summarize the properties of datasets in Table 4.

**Algorithms.** We implement all the algorithms in Java. (1) AOD<sub>1</sub> and AOD<sub>3</sub>, different versions of DisAOD for measures  $g_1$  and  $g_3$  respectively. (2) Some variants of DisAOD, for testing the effectiveness of optimization techniques and ranking functions (details are provided later). (3) FastOD, the

TABLE 4  
Datasets, Execution Statistics of AOD<sub>1</sub> and AOD<sub>3</sub>

Dataset Properties			AOD <sub>1</sub> ( $e = 0.001$ )		AOD <sub>3</sub> ( $e = 0.01$ )		Exact LOD	
DataSet	$ r $	$ R $	Time(s)	$ AOD $	Time(s)	$ AOD $	Time(s)	$ LOD $
NCV	930K	15	2,036	184	3,374	102	26	20
FLI	500K	17	60,684	76,288	4,142	386	59,857 <sup>+</sup>	526,906 <sup>+</sup>
Hepa	155	20	0.136	0	0.096	0	0.163	0
Atom	33K	10	2	48	5	126	0.485	108
DB	250K	13	21	10	57	2	3	2
Letter	20K	17	0.833	0	0.191	0	0.148	0
FDR	250K	30	110	10	1,200	40	47	1,917

algorithm for discovering approximate set-based canonical ODs with measure  $g_3$  [12]. (4) SampleAOD<sub>1</sub>, the method that performs AOD discovery with  $g_1$  by sampling (Section 7).

**Parameter settings.** In addition to  $|r|$  and  $|R|$ , we use one more parameter: the error threshold  $e$ . We use random sampling (resp. projection) to vary  $|r|$  (resp.  $|R|$ ) when required.

**Running environment.** We run all experiments on a PC server with an Intel Xeon E-2224 3.4G CPU, 64GB of memory and CentOS 7, and report the average results of 5 runs. For space limitation, we provide more experiments in the full version at <https://github.com/chenjixuan20/AOD>, along with all the tested datasets and code.

### 8.2 Efficiency of AOD Discovery

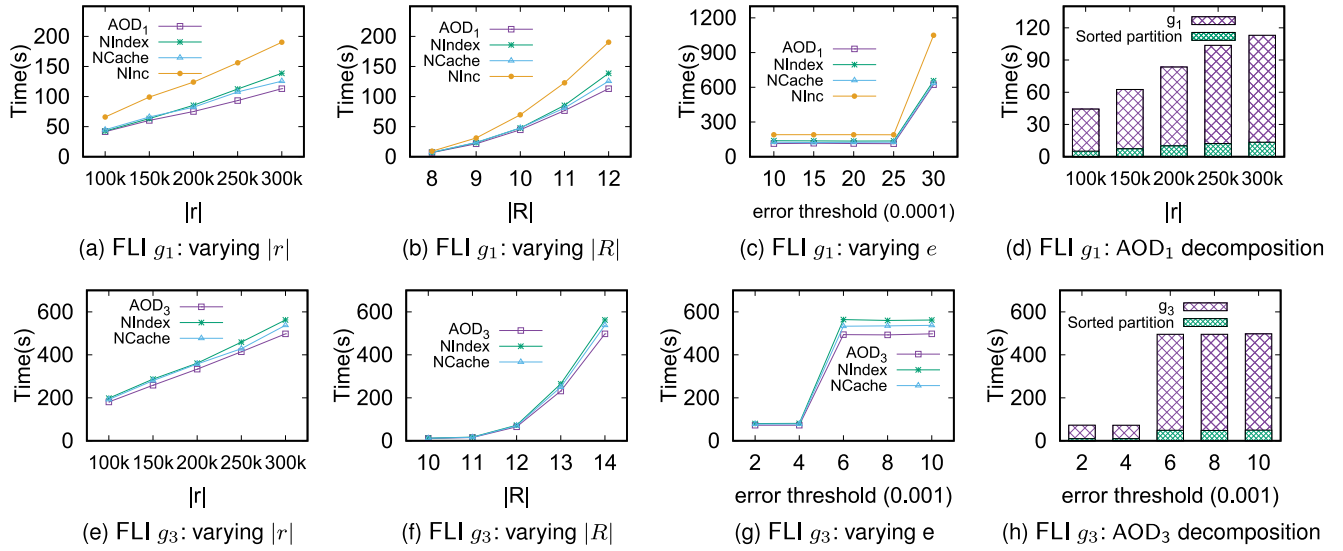
**Exp-1: AOD<sub>1</sub> and AOD<sub>3</sub> on all datasets.** We summarize results of AOD<sub>1</sub> and AOD<sub>3</sub> on all tested data in Table 4, with running times (in seconds) and the number  $|AOD|$  of discovered AODs. We use different threshold  $e$  for AOD<sub>1</sub> and AOD<sub>3</sub>. Intuitively, an erroneous tuple incurs a  $g_3$  value of  $\frac{1}{|r|}$ , but a single violation incurs a  $g_1$  value of  $\frac{1}{|r|^2 - |r|}$ . Note that the times of AOD<sub>1</sub> and AOD<sub>3</sub> cannot be compared, since the sets of discovered AODs by them are quite different.

We are not aware of any existing works on approximate LOD discovery. For reference, we show the experimental results of the state-of-the-art exact LOD discovery algorithm [11] (obtained from the authors). We see approximate discoveries are usually slower than the exact counterpart. This is reasonable since exact discovery is a special case of approximate one, and the method of [11] fully leverages the properties of exact LODs for an efficient pruning. A valid exact LOD is always a valid AOD, but it may be a non-minimal AOD. Hence, the result of approximate discovery does not necessarily contain that of exact discovery. Indeed, the number of exact LODs can be much larger than that of AODs. For example, the exact discovery already finds more than 500K exact LODs on FLI when we manually terminate it. Intuitively, instead of discovering a proper AOD, exact discovery may find many specializations of it to resolve violations that are actually incurred by errors in the data.

In the sequel we further compare AOD<sub>1</sub> (resp. AOD<sub>3</sub>) against its variants, by varying parameters.

**Exp-2: AOD<sub>1</sub> against Variants.** We implement some variants by disabling an optimization each time (Section 6.2). (1) NInc disables the incremental computation of *swapped* and *ordered* pairs; (2) NIndex disables the index for minimality check; and (3) NCache disables the sorted partition cache. We also test a variant without leveraging lower/upper bounds; it is always

<sup>1</sup>  $\mu_{0.98} = 2.326$  in the standard normal distribution table.

Fig. 4. AOD<sub>1</sub>, AOD<sub>3</sub> against Variants.

orders of magnitude slower (not shown). The results show that pruning rules with bounds are crucial to the efficiency of AOD discovery.

We set  $|r| = 300K$ ,  $|R| = 12$  and  $e = 0.001$  by default on FLI, and vary  $|r|$  from 100K to 300K in Fig. 4a,  $|R|$  from 8 to 12 in Fig. 4b, and  $e$  from 0.001 to 0.003 in Fig. 4c.

We see the following. (1) AOD<sub>1</sub> scales well with  $|r|$ , consistent with the complexity analysis. As  $|r|$  increases from 100K to 300K, the time increases from 41s to 113s. (2)  $|R|$  significantly affects the efficiency, as expected. We find the number of discovered AODs increases from 2 to 56 (not shown), as  $|R|$  increases from 8 to 12. However, we see the performance is much better than the worst-case complexity that is factorial in  $|R|$ . Indeed, the running time is mostly affected by the size of the actual search space, i.e., the number of AODs that are required to be checked (not pruned). Recall that function *Compute* (line 8 of DisAOD) is called to compute the error measures and lower/upper bounds of candidate AODs, and hence, we use the number of calling this function as the size of the actual search space. We experimentally find that only 232 (resp. 4,668) AODs are checked for  $|R| = 8$  (resp. 12). (3) The threshold  $e$  can significantly affect the results of AOD<sub>1</sub>. Intuitively, a large  $e$  value enables AODs with more attributes on the RHS and in turn more attributes on the LHS, while a small  $e$  value usually leads to fewer attributes on the RHS. We find the number of discovered AODs almost remains unchanged when  $e$  is in the range of  $[0.001, 0.0025]$ , but sharply increases by more than 50 times as  $e$  increases to 0.003. The results also demonstrate the benefit of the upper bound. A relatively large  $e$ , e.g., 0.003, helps AOD<sub>1</sub> quickly generate valid AODs. Hence, the time increases by less than 6 times as the number of valid AODs increases by more than 50 times.

In terms of the optimizations, we see the following. (1) AOD<sub>1</sub> is up to 1.7 times and on average 61% faster than NInc. The cost of AOD<sub>1</sub> mainly consists of the times for creating sorted partitions and for computing  $g_1$  values. Along the same setting as Fig. 4a, we show the two times respectively in Fig. 4d. We find the latter governs the overall time, and hence the incremental computation of  $g_1$  significantly improves the efficiency. (2) NIndex and NCache concern computations of

sorted partitions. AOD<sub>1</sub> is on average 14% and 10% and up to 23% and 15% faster than NIndex and NCache, respectively.

*Exp-3: AOD<sub>3</sub> against Variants.* We compare AOD<sub>3</sub> against variants (excluding NInc). The usage of lower/upper bounds is again experimentally found to be crucial (not shown).

We set  $|r| = 300K$ ,  $|R| = 14$  and  $e = 0.01$  by default on FLI, and vary  $|r|$  from 100K to 300K in Fig. 4e,  $|R|$  from 10 to 14 in Fig. 4f, and  $e$  from 0.002 to 0.01 in Fig. 4g. We see the following. (1) AOD<sub>3</sub> scales well with  $|r|$ . (2) The efficiency of AOD<sub>3</sub> is sensitive to  $|R|$ . We find the number of discovered AODs increases from 12 to 142, and the size of the actual search space increases from 574 to 9,654 (not shown), as  $|R|$  increases from 10 to 14. The actual search space is much smaller than the theoretical worst-case one. (3) The running time increases from 73s to 495s on FLI when  $e$  increases from 0.004 to 0.006. This is because the number of discovered AODs increases from 20 to 142. (4) AOD<sub>3</sub> is on average 12% and 7% and up to 18% and 10% faster than NIndex and NCache, respectively. (5) The computation of  $g_3$  on average takes more than 89% of the total time, as shown in Fig. 4h (along the same setting as Fig. 4g).

*Exp-4: AOD<sub>3</sub> against FastOD.* We compare AOD<sub>3</sub> against FastOD [12]. FastOD discovers approximate set-based canonical ODs with measure  $g_3$ . Recall that set-based ODs generalize LODs (Section 2), but this does not apply to AODs. Intuitively, a valid AOD does not guarantee that ODs from the corresponding set of canonical ODs are all valid approximate ones. Hence, the connection between approximate LODs and set-based canonical ODs is broken.

The results in Table 5 show AOD<sub>3</sub> is usually much faster than FastOD. We find this is mainly because the number of

TABLE 5  
AOD<sub>3</sub> Against FastOD on Various Datasets

Dataset Properties			AOD <sub>3</sub> ( $e = 0.01$ )		FastOD ( $e = 0.01$ )	
DataSet	$ r $	$ R $	Time(s)	$ AOD $	Time(s)	$ OD $
NCV	30K	15	109	138	650	245
FLI	50K	12	42	56	78	82
Hepa	155	20	0.096	0	198	6,862
DB	250K	13	57	2	333	785
FDR	250k	15	96	6	482	1,208

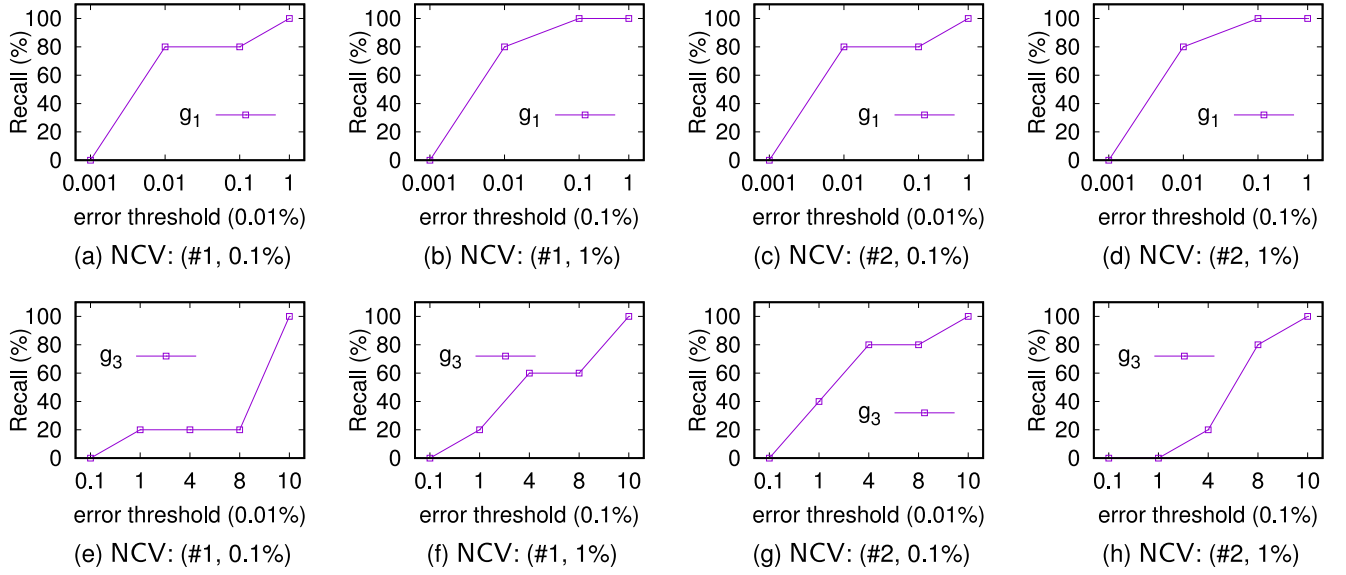


Fig. 5. Effectiveness of AOD discovery on NCV.

discovered approximate canonical ODs is typically much larger than that of AODs. Similar results are reported in the exact canonical OD and LOD discoveries [11], [29].

### 8.3 Effectiveness of AOD Discovery and Ranking

*Exp-5: Recall of AOD discovery.* In this set of experiments, we verify the effectiveness of AOD discovery by recalling LODs hidden in dirty data. We first manually identify some “golden” LODs verified by domain experts on dataset NCV. We then use 10K tuples of NCV and introduce noises to them; values close to the original correct ones are used as noises, which is common in practice. The introduction of noise is controlled by noise ratio  $\theta$  and two strategies [21]. #1: on each attribute, each value has a probability of  $\theta$  to be assigned a new value from the active domain. #2: each tuple has a probability of  $\theta$  to be selected, and new values from the active domain are assigned to all values of selected tuples. Intuitively, #2 is a setting that favors  $g_3$ , since noises are on fewer tuples in #2 than #1.

We run DisAOD on the dirty dataset, and compute the *recall* as the ratio of the number of discovered golden LODs to the total number of golden LODs. From Fig. 5a to 5h, we vary the error threshold  $e$  and test various settings ( $\#i, \theta$ ). We see the following. (a) In contrast to exact LOD discoveries with a recall of 0 in all settings (not shown), both  $AOD_1$  and  $AOD_3$  have a recall of 100% when  $e$  is above a threshold. We denote this threshold by  $e_o$ . (b) In  $AOD_1$ ,  $e_o$  is much smaller than the noise ratio  $\theta$ , by up to orders of magnitude. We find NCV has sparse value distributions on some attributes, and hence the

introduced new values lead to very few violations. (c) In  $AOD_3$ ,  $e_o$  is very close to the noise ratio  $\theta$ . This is expected in #2 (Figures 5g and 5h);  $g_3$  concerns the number of violating tuples. We find  $e_o$  is also close to  $\theta$  in #1 (Figures 5e and 5f). This is because there are very few violations in NCV and the number of violating tuples in #1 is similar to that in #2.

*Exp-6: Ranking functions.* Using FLI, we verify the effectiveness of ranking functions in identifying meaningful AODs. Besides the two ranking functions proposed in Section 6.3, we also rank AODs in terms of their error rates. All discovered AODs have error rates no larger than the given threshold, but their error rates are usually different. A discovered AOD is an exact LOD only when its error rate is 0.

We measure ranking functions in terms of *precision*, along the same line as [23]. Specifically, we use 10K tuples of FLI, perform AOD discovery with  $g_1$ , obtain top-10 AODs for each ranking function, and ask humans to judge whether the AODs are meaningful or not. The situation is practical, since it is typically much easier for humans to verify discovered dependencies than to manually design dependencies. The precision of each ranking function is the ratio of the number of meaningful AODs (judged by humans) to 10.

The results in Table 6 show that all the three ranking functions can effectively help identify meaningful AODs. Most meaningful AODs on FLI have short LHS and RHS attribute lists (some examples are given in Table 7 with their actual error rates). Note that only AOD discovery but not exact LOD

TABLE 6  
Precisions of Ranking Functions on FLI

error threshold $e$	ability to sort	succinctness	error rate
0.01	1	1	1
0.001	1	1	1
0.0001	0.8	0.8	0.6

TABLE 7  
Example AODs on FLI

AOD	error rate
OriginAirportSeqID $\mapsto$ OriginAirportID	0.000128
DestAirportSeqID $\mapsto$ DestAirportID	0.0
Origin $\mapsto$ OriginAirportSeqID	0.000024



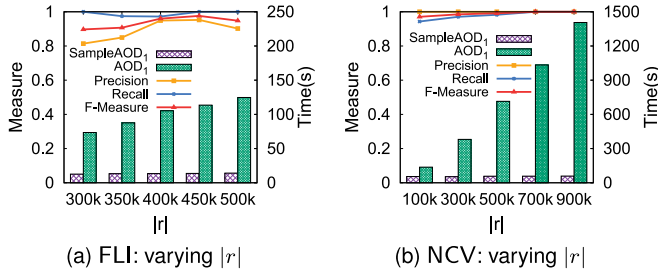


Fig. 6. SampleAOD<sub>1</sub> against AOD<sub>1</sub>.

discovery can find meaningful AODs whose error rates are not 0. We test different error thresholds. When 0.0001 is used, we find some too “specialized” AODs are discovered instead of meaningful AODs that do not satisfy the threshold. For example, instead of  $\text{OriginAirportSeqID} \mapsto \text{OriginAirportID}$ ,  $\text{OriginAirportSeqIDFlightDate} \mapsto \text{OriginAirportID}$  is discovered when  $\epsilon = 0.0001$ , but the latter one is rejected by human verification. This hinders the precisions of ranking functions, since rankings are applied to the discovered AODs. This suggests that a small set of AODs can be inspected by users to help select a proper threshold.

#### 8.4 Effectiveness of Estimation by Sampling

*Exp-7: AOD discovery by sampling.* We experimentally study SampleAOD<sub>1</sub> on FLI ( $|R| = 11$  and  $\epsilon = 0.01$ ) and NCV ( $|R| = 14$  and  $\epsilon = 0.01$ ) in Figs. 6a and 6b, respectively. We also give the running time of AOD<sub>1</sub> for comparison. To avoid the bias of random sampling, we report the average of 10 runs. In this set of experiments, the *precision* ( $P$ ) is the proportion of AODs discovered by SampleAOD<sub>1</sub> that are also discovered by AOD<sub>1</sub>, the *recall* ( $R$ ) is the proportion of AODs discovered by AOD<sub>1</sub> that are also discovered by SampleAOD<sub>1</sub>, and the *F-measure* ( $F$ ) is  $2 \cdot (P \cdot R) / (P + R)$ .

We see the following. (1) SampleAOD<sub>1</sub> always achieves high *precision*, *recall* and *F-measure*, i.e.,  $P \geq 0.81$ ,  $R \geq 0.97$  and  $F \geq 0.90$  on FLI and  $P \geq 0.99$ ,  $R \geq 0.94$  and  $F \geq 0.97$  on NCV. The values of  $P$ ,  $R$  and  $F$  are not affected by the increase of  $|r|$ , consistent with our theoretical analysis. (2) The running time of SampleAOD<sub>1</sub> almost remains unchanged on each dataset, since the sample size is independent of  $|r|$ . Hence, SampleAOD<sub>1</sub> is much faster than AOD<sub>1</sub>, up to 8.8 times on FLI and 24.7 times on NCV.

The results demonstrate that SampleAOD<sub>1</sub> can deliver results of high accuracy in far less time.

## 9 CONCLUSION

We have formalized the AOD discovery problem, developed efficient algorithms and optimizations for measures  $g_1$ ,  $g_3$ , related lower/upper bounds and AOD discoveries. We have studied techniques for estimating  $g_1$  by sampling. We have also experimentally verified the benefits of our methods.

We intend to further study the workflow of AOD discovery. There are some technical issues that need to be addressed, e.g., the method to help users determine and adjust the error threshold  $\epsilon$  when required. We also intend to further study the complexity in the framework of *fixed-parameter tractable*, along the same lines as [3].

## REFERENCES

- [1] Z. Abedjan, L. Golab, and F. Naumann, “Data profiling: A tutorial,” in *Proc. ACM Int. Conf. Manage. Data*, 2017, pp. 1747–1751.
- [2] J. L. Bentley and J. H. Friedman, “Data structures for range searching,” *ACM Comput. Surv.*, vol. 11, no. 4, pp. 397–409, 1979.
- [3] T. Bläsius, T. Friedrich, and M. Schirneck, “The parameterized complexity of dependency detection in relational databases,” in *Proc. Int. Symp. Parameterized Exact Comput.*, 2016, vol. 63, pp. 6:1–6:13.
- [4] X. Chu, I. F. Ilyas, and P. Papotti, “Discovering denial constraints,” *Proc. VLDB Endowment*, vol. 6, no. 1, pp. 1498–1509, 2013.
- [5] C. Consonni, P. Sottovia, A. Montresor, and Y. Velegrakis, “Discovering order dependencies through order compatibility,” in *Proc. Int. Conf. Extending Database Technol.*, 2019, pp. 409–420.
- [6] M. L. Fredman, “On computing the length of longest increasing subsequences,” *Discret. Math.*, vol. 11, no. 1, pp. 29–35, 1975.
- [7] S. Ginsburg and R. Hull, “Order dependency in the relational model,” *Theor. Comput. Sci.*, vol. 26, pp. 149–195, 1983.
- [8] S. Ginsburg and R. Hull, “Sort sets in the relational model,” *J. ACM*, vol. 33, no. 3, pp. 465–488, 1986.
- [9] L. Golab, H. J. Karloff, F. Korn, A. Saha, and D. Srivastava, “Sequential dependencies,” *Proc. VLDB Endowment*, vol. 2, no. 1, pp. 574–585, 2009.
- [10] Y. Jin, Z. Tan, W. Zeng, and S. Ma, “Approximate order dependency discovery,” in *Proc. IEEE 37th Int. Conf. Data Eng.*, 2021, pp. 25–36.
- [11] Y. Jin, L. Zhu, and Z. Tan, “Efficient bidirectional order dependency discovery,” in *Proc. IEEE 36th Int. Conf. Data Eng.*, 2020, pp. 61–72.
- [12] R. Karegar, P. Godfrey, L. Golab, M. Kargar, D. Srivastava, and J. Szlichta, “Efficient discovery of approximate order dependencies,” in *Proc. Int. Conf. Extending Database Technol.*, 2021, pp. 427–432.
- [13] B. Kenig, P. Munda, G. Prasad, B. Salimi, and D. Suciu, “Mining approximate acyclic schemes from relations,” in *Proc. ACM SIGMOD Int. Conf. Manage. Data*, 2020, pp. 297–312.
- [14] B. Kenig and D. Suciu, “Integrity constraints revisited: From exact to approximate implication,” in *Proc. Int. Conf. Database Theory*, 2020, pp. 18:1–18:20.
- [15] J. Kivinen and H. Mannila, “Approximate dependency inference from relations,” in *Proc. Int. Conf. Database Theory*, 1992, pp. 86–98.
- [16] S. Kruse and F. Naumann, “Efficient discovery of approximate dependencies,” *Proc. VLDB Endowment*, vol. 11, no. 7, pp. 759–772, 2018.
- [17] P. Langer and F. Naumann, “Efficient order dependency detection,” *Int. J. Very Large Data Bases J.*, vol. 25, no. 2, pp. 223–241, 2016.
- [18] P. Li, M. H. Böhlen, J. Szlichta, and D. Srivastava, “Discovery of band order dependencies,” *CoRR*, vol. abs/1905.11948, 2019. [Online]. Available: <http://arxiv.org/abs/1905.11948>
- [19] P. Li, J. Szlichta, M. H. Böhlen, and D. Srivastava, “Discovering band order dependencies,” in *Proc. IEEE 36th Int. Conf. Data Eng.*, 2020, pp. 1878–1881.
- [20] D. Liben-Nowell, E. Vee, and A. Zhu, “Finding longest increasing and common subsequences in streaming data,” *J. Comb. Optim.*, vol. 11, no. 2, pp. 155–175, 2006.
- [21] E. Livshits, A. Heidari, I. F. Ilyas, and B. Kimelfeld, “Approximate denial constraints,” *Proc. VLDB Endowment*, vol. 13, no. 10, pp. 1682–1695, 2020.
- [22] T. Papenbrock and F. Naumann, “A hybrid approach to functional dependency discovery,” in *Proc. Int. Conf. Manage. Data*, 2016, pp. 821–833.
- [23] E. H. M. Pena, E. C. de Almeida, and F. Naumann, “Discovery of approximate (and exact) denial constraints,” *Proc. VLDB Endowment*, vol. 13, no. 3, pp. 266–278, 2019.
- [24] J. Rammelaere and F. Geerts, “Revisiting conditional functional dependency discovery: Splitting the “C” from the “FD”,” in *Proc. Eur. Conf. Mach. Learn. Princ. Pract. Knowl. Discov. Databases*, 2018, pp. 552–568.
- [25] D. E. Simmen, E. J. Shekita, and T. Malkemus, “Fundamental techniques for order optimization,” in *Proc. ACM SIGMOD Int. Conf. Manage. Data*, 1996, pp. 57–67.
- [26] S. Song, L. Chen, and P. S. Yu, “Comparable dependencies over heterogeneous data,” *VLDB J.*, vol. 22, no. 2, pp. 253–274, 2013.
- [27] S. Song, F. Gao, R. Huang, and C. Wang, “Data dependencies extended for variety and veracity: A family tree,” *IEEE Trans. Knowl. Data Eng.*, early access, Dec. 25, 2020, doi: [10.1109/TKDE.2020.3046443](https://doi.org/10.1109/TKDE.2020.3046443).

- [28] J. Szlichta, P. Godfrey, L. Golab, M. Kargar, and D. Srivastava, "Effective and complete discovery of order dependencies via set-based axiomatization," *Proc. VLDB Endowment*, vol. 10, no. 7, pp. 721–732, 2017.
- [29] J. Szlichta, P. Godfrey, L. Golab, M. Kargar, and D. Srivastava, "Effective and complete discovery of bidirectional order dependencies via set-based axioms," *VLDB J.*, vol. 27, no. 4, pp. 573–591, 2018.
- [30] J. Szlichta, P. Godfrey, and J. Gryz, "Fundamentals of order dependencies," *Proc. VLDB Endowment*, vol. 5, no. 11, pp. 1220–1231, 2012.
- [31] J. Szlichta, P. Godfrey, J. Gryz, W. Ma, W. Qiu, and C. Zuzarte, "Business-intelligence queries with order dependencies in DB2," in *Proc. Extending Database Technol.*, 2014, pp. 750–761.
- [32] J. Szlichta, P. Godfrey, J. Gryz, and C. Zuzarte, "Expressiveness and complexity of order dependencies," *Proc. VLDB Endowment*, vol. 6, no. 14, pp. 1858–1869, 2013.
- [33] Z. Tan, A. Ran, S. Ma, and S. Qin, "Fast incremental discovery of pointwise order dependencies," *Proc. VLDB Endowment*, vol. 13, no. 10, pp. 1669–1681, 2020.



**Yifeng Jin** received the master's degree in computer science from Fudan University in 2021. He is currently a software engineer with Alibaba Group. His research interests include databases on modern hardware and distributed computations.



**Zijong Tan** received the PhD degree from Fudan University. He was a visiting researcher with the University of Edinburgh. He is currently an associate professor with the School of Computer Science, Fudan University, China. His research interests include data profiling and data quality management.



**Jixuan Chen** received the BS degree in software engineering from the Huazhong University of Science and Technology in 2020. He is currently working toward the master's degree with the School of Computer Science, Fudan University, China. His research interests include data profiling and distributed computations.



**Shuai Ma** received the PhD degree from the University of Edinburgh in 2011 and the PhD degree from Peking University in 2004. He is currently a professor with the School of Computer Science and Engineering, Beihang University, China. His research interests include database theory and systems and big data. He was the recipient of the Best Paper Award for VLDB 2010. He has been an associate editor for *VLDB Journal* since 2017 and the *IEEE Transactions On Big Data and Knowledge and Information Systems* since 2020.

► For more information on this or any other computing topic, please visit our Digital Library at [www.computer.org/csdl](http://www.computer.org/csdl).