

Certain Answers meet Zero-One Laws

Leonid Libkin
[University of Edinburgh](#)

A company database: orders, customers, payments

Orders

ORDER_ID	TITLE	PRICE
Ord1	"Big Data"	30
Ord2	"SQL"	35
Ord3	"Logic"	50

Pay

CUST_ID	ORDER
c1	Ord1
c2	Ord2

Customer

CUST_ID	NAME
c1	John
c2	Mary

A company database: orders, customers, payments

Orders

ORDER_ID	TITLE	PRICE
Ord1	"Big Data"	30
Ord2	"SQL"	35
Ord3	"Logic"	50

Pay

CUST_ID	ORDER
c1	Ord1
c2	Ord2

Customer

CUST_ID	NAME
c1	John
c2	Mary

Typical queries, as we teach students to write them:

A company database: orders, customers, payments

Orders

ORDER_ID	TITLE	PRICE
Ord1	"Big Data"	30
Ord2	"SQL"	35
Ord3	"Logic"	50

Pay

CUST_ID	ORDER
c1	Ord1
c2	Ord2

Customer

CUST_ID	NAME
c1	John
c2	Mary

Typical queries, as we teach students to write them:

Unpaid orders:

```
select O.order_id  
from Orders O  
where O.order_id not in  
      (select order from Pay P)
```

A company database: orders, customers, payments

Orders

ORDER_ID	TITLE	PRICE
Ord1	"Big Data"	30
Ord2	"SQL"	35
Ord3	"Logic"	50

Pay

CUST_ID	ORDER
c1	Ord1
c2	Ord2

Customer

CUST_ID	NAME
c1	John
c2	Mary

Typical queries, as we teach students to write them:

Unpaid orders:

```
select O.order_id  
from Orders O  
where O.order_id not in  
      (select order from Pay P)
```

Customers without an order:

```
select C.cust_id from Customer C  
where not exists  
      (select * from Orders O, Pay P  
       where C.cust_id=P.cust_id  
       and P.order=O.order_id)
```

A company database: orders, customers, payments

Orders

ORDER_ID	TITLE	PRICE
Ord1	"Big Data"	30
Ord2	"SQL"	35
Ord3	"Logic"	50

Pay

CUST_ID	ORDER
c1	Ord1
c2	Ord2

Customer

CUST_ID	NAME
c1	John
c2	Mary

Typical queries, as we teach students to write them:

Unpaid orders:

```
select O.order_id  
from Orders O  
where O.order_id not in  
      (select order from Pay P)
```

Answer: Ord3.

Customers without an order:

```
select C.cust_id from Customer C  
where not exists  
      (select * from Orders O, Pay P  
       where C.cust_id=P.cust_id  
       and P.order=O.order_id)
```

Answer: none.

A company database: orders, customers, payments

Orders

ORDER_ID	TITLE	PRICE
Ord1	"Big Data"	30
Ord2	"SQL"	35
Ord3	"Logic"	50

Pay

CUST_ID	ORDER
c1	Ord1
c2	Ord2

Customer

CUST_ID	NAME
c1	John
c2	Mary

Typical queries, as we teach students to write them:

Unpaid orders:

```
select O.order_id  
from Orders O  
where O.order_id not in  
      (select order from Pay P)
```

Customers without an order:

```
select C.cust_id from Customer C  
where not exists  
      (select * from Orders O, Pay P  
       where C.cust_id=P.cust_id  
         and P.order=O.order_id)
```

A company database: orders, customers, payments

Orders

ORDER_ID	TITLE	PRICE
Ord1	"Big Data"	30
Ord2	"SQL"	35
Ord3	"Logic"	50

Pay

CUST_ID	ORDER
c1	Ord1
c2	Ord2

Customer

CUST_ID	NAME
c1	John
c2	Mary

Unpaid orders:

```
select O.order_id  
from Orders O  
where O.order_id not in  
      (select order from Pay P)
```

Customers without an order:

```
select C.cust_id from Customer C  
where not exists  
      (select * from Orders O, Pay P  
       where C.cust_id=P.cust_id  
       and P.order=O.order_id)
```


A company database: orders, customers, payments

Orders

ORDER_ID	TITLE	PRICE
Ord1	"Big Data"	30
Ord2	"SQL"	35
Ord3	"Logic"	50

Pay

CUST_ID	ORDER
c1	Ord1
c2	Ord2

Customer

CUST_ID	NAME
c1	John
c2	Mary

In the real world, information is often missing

Unpaid orders:

```
select O.order_id  
from Orders O  
where O.order_id not in  
      (select order from Pay P)
```

Customers without an order:

```
select C.cust_id from Customer C  
where not exists  
      (select * from Orders O, Pay P  
       where C.cust_id=P.cust_id  
       and P.order=O.order_id)
```

A company database: orders, customers, payments

Orders

ORDER_ID	TITLE	PRICE
Ord1	"Big Data"	30
Ord2	"SQL"	35
Ord3	"Logic"	50

Pay

Customer

CUST_ID	NAME
c1	John
c2	Mary

In the real world, information is often missing

Unpaid orders:

```
select O.order_id  
from Orders O  
where O.order_id not in  
      (select order from Pay P)
```

Customers without an order:

```
select C.cust_id from Customer C  
where not exists  
      (select * from Orders O, Pay P  
       where C.cust_id=P.cust_id  
       and P.order=O.order_id)
```

A company database: orders, customers, payments

Orders

ORDER_ID	TITLE	PRICE
Ord1	"Big Data"	30
Ord2	"SQL"	35
Ord3	"Logic"	50

Pay

CUST_ID	ORDER
c1	Ord1
c2	--

Customer

CUST_ID	NAME
c1	John
c2	Mary

In the real world, information is often missing

Unpaid orders:

```
select O.order_id
from Orders O
where O.order_id not in
      (select order from Pay P)
```

Customers without an order:

```
select C.cust_id from Customer C
where not exists
      (select * from Orders O, Pay P
       where C.cust_id=P.cust_id
       and P.order=O.order_id)
```

A company database: orders, customers, payments

Orders

ORDER_ID	TITLE	PRICE
Ord1	"Big Data"	30
Ord2	"SQL"	35
Ord3	"Logic"	50

Pay

CUST_ID	ORDER
c1	Ord1
c2	--

Customer

CUST_ID	NAME
c1	John
c2	Mary

In the real world, information is often missing

Unpaid orders:

```
select O.order_id  
from Orders O  
where O.order_id not in  
      (select order from Pay P)
```

Customers without an order:

```
select C.cust_id from Customer C  
where not exists  
      (select * from Orders O, Pay P  
       where C.cust_id=P.cust_id  
       and P.order=O.order_id)
```

Old Answer: Ord3 New: NONE!

Old answer: none New: c2!

Problems

- There is a good understanding of what correctness is: **certain answers**
 - true in all completions (restricted **validity**)
- Computationally hard: **coNP**-hard for basic SQL
 - Hence DBMSs **sacrifice correctness to ensure efficiency**
 - SQL DBMSs use special rules based on 3-valued logic to get query answers
 - and these answers can be very wrong...

How to solve it

For many years, the community adopted this approach



Recently, a new idea emerged: approximations of certain answers

Can be found efficiently for all relational algebra queries

**Behave well in theory (L., ACM TODS 2016)
and practice (Guagliardo, L. PODS'16 + followups)**

Naive Evaluation

- Treat nulls as **new constants**
- Evaluate query using standard techniques
- Heavily used: data integration/exchange, OBDA etc

Orders

ORDER_ID	TITLE	PRICE
Ord1	“Big Data”	30
Ord2	“SQL”	35
Ord3	“Logic”	50

Pay

Customer

CUST_ID	NAME
c1	John
c2	Mary

Naive Evaluation

- Treat nulls as **new constants**
- Evaluate query using standard techniques
- Heavily used: data integration/exchange, OBDA etc

Orders

ORDER_ID	TITLE	PRICE
Ord1	"Big Data"	30
Ord2	"SQL"	35
Ord3	"Logic"	50

Pay

Customer

CUST_ID	NAME
c1	John
c2	Mary

Unpaid orders:

```
select O.order_id
from Orders O
where O.order_id not in
      (select order from Pay P)
```


Naive Evaluation

- Treat nulls as **new constants**
- Evaluate query using standard techniques
- Heavily used: data integration/exchange, OBDA etc

Orders

ORDER_ID	TITLE	PRICE
Ord1	"Big Data"	30
Ord2	"SQL"	35
Ord3	"Logic"	50

Pay

Customer

CUST_ID	NAME
c1	John
c2	Mary

Unpaid orders:

```
select O.order_id
from Orders O
where O.order_id not in
      (select order from Pay P)
```

Customers without an order:

```
select C.cust_id from Customer C
where not exists
      (select * from Orders O, Pay P
       where C.cust_id=P.cust_id
       and P.order=O.order_id)
```

Naive Evaluation

- Treat nulls as **new constants**
- Evaluate query using standard techniques
- Heavily used: data integration/exchange, OBDA etc

Orders

ORDER_ID	TITLE	PRICE
Ord1	"Big Data"	30
Ord2	"SQL"	35
Ord3	"Logic"	50

Pay

Customer

CUST_ID	NAME
c1	John
c2	Mary

Unpaid orders:

```
select O.order_id
from Orders O
where O.order_id not in
      (select order from Pay P)
```

Answer: Ord2, Ord3.

Customers without an order:

```
select C.cust_id from Customer C
where not exists
      (select * from Orders O, Pay P
       where C.cust_id=P.cust_id
       and P.order=O.order_id)
```

Answer: c2.

Naive Evaluation

- Treat nulls as **new constants**
- Evaluate query using standard techniques
- Heavily used: data integration/exchange, OBDA etc

Orders

ORDER_ID	TITLE	PRICE
Ord1	"Big Data"	30
Ord2	"SQL"	35
Ord3	"Logic"	50

Pay

Customer

CUST_ID	NAME
c1	John
c2	Mary

Unpaid orders:

```
select O.order_id
from Orders O
where O.order_id not in
      (select order from Pay P)
```

Customers without an order:

```
select C.cust_id from Customer C
where not exists
      (select * from Orders O, Pay P
       where C.cust_id=P.cust_id
       and P.order=O.order_id)
```

Naive Evaluation

- Treat nulls as **new constants**
- Evaluate query using standard techniques
- Heavily used: data integration/exchange, OBDA etc

Orders

ORDER_ID	TITLE	PRICE
Ord1	"Big Data"	30
Ord2	"SQL"	35
Ord3	"Logic"	50

Pay

CUST_ID	ORDER
c1	Ord1
c2	\perp

Customer

CUST_ID	NAME
c1	John
c2	Mary

Unpaid orders:

```
select O.order_id  
from Orders O  
where O.order_id not in  
  (select order from Pay P)
```

Customers without an order:

```
select C.cust_id from Customer C  
where not exists  
  (select * from Orders O, Pay P  
   where C.cust_id=P.cust_id  
     and P.order=O.order_id)
```


How bad are bad answers?

- What if \perp is a constant that is different from Ord1, Ord2, Ord3?
- Then naive evaluation actually produces **correct** answers!
- If we know nothing about it is not such an unreasonable assumption: there could be many orders?
- But what if we know $\perp \in \{\text{Ord1}, \text{Ord2}, \text{Ord3}\}$?
- Then answer to the first query is Ord2 with 50% chance and Ord3 with 50% chance. Answer to the second query is empty.

Questions

- Is naive evaluation always good without constraints on nulls, or we just got lucky?
 - Yes, it always is
- Can we get the second type of answers, with constraints?
 - Yes, but with more work
- Now revisit certain answers, and connect with a well know subject in logic and probability

Incomplete data and certain answers




Incomplete database D represents many complete databases D_1, D_2, \dots

This is done by interpreting incompleteness

For example, by assigning values to every null that occurs in D

Incomplete data and certain answers



Tuple a is **certain answer** to query Q in D
 $\Leftrightarrow a$ is an answer to Q in every D_i

Certainty is **hard** computationally:
coNP-hard for relational algebra queries

Zero-One Laws

A formula α over graphs; **green** = true; **red** = false



α is **almost surely valid**: true in almost all graphs

Examples:

- $\mu(\text{has isolated node}) = 0$
- $\mu(\text{is a tree}) = 0$
- $\mu(\text{connected}) = 1$
- $\mu(\text{has diameter at most 2}) = 1$

Zero-One Laws

A formula α over graphs; **green** = true; **red** = false



α is **almost surely valid**: true in almost all graphs

- pick a graph G at random
- calculate the probability $\mu(\alpha)$ that α is true in G
- $\mu(\alpha) = 1 \Leftrightarrow \alpha$ is almost surely valid

Examples:

- $\mu(\text{has isolated node}) = 0$
- $\mu(\text{is a tree}) = 0$
- $\mu(\text{connected}) = 1$
- $\mu(\text{has diameter at most 2}) = 1$

Zero-One Laws

**Extended to many other logics: Fixed-point, Infinitary logics,
Fragments of second-order logic; Other distributions too**

A very active subject in logic/combinatorics

Zero-One Laws

Fagin 1976:
if α is first-order, then $\mu(\alpha)$ is 0 or 1

**Extended to many other logics: Fixed-point, Infinitary logics,
Fragments of second-order logic; Other distributions too**

A very active subject in logic/combinatorics

Zero-One Laws


Fagin 1976:
if α is first-order, then $\mu(\alpha)$ is 0 or 1

α is valid (true in all graphs) - **undecidable**.
 α is almost surely valid ($\mu(\alpha) = 1$) - **easy to decide**.

**Extended to many other logics: Fixed-point, Infinitary logics,
Fragments of second-order logic; Other distributions too**

A very active subject in logic/combinatorics

Certainty and Zero-One Laws




For query **Q**:

- pick a complete database **D_i** at random
- $\mu(Q, D, a)$: probability that $a \in Q(D_i)$

$$\mu(Q, D, a) = 1 \Rightarrow$$

a = almost certainly true answer to **Q** in **D**

Certainty and Zero-One Laws



For query Q :

- pick a complete database D_i at random
- $\mu(Q, D, a)$: probability that $a \in Q(D_i)$

$$\mu(Q, D, a) = 1 \Rightarrow$$

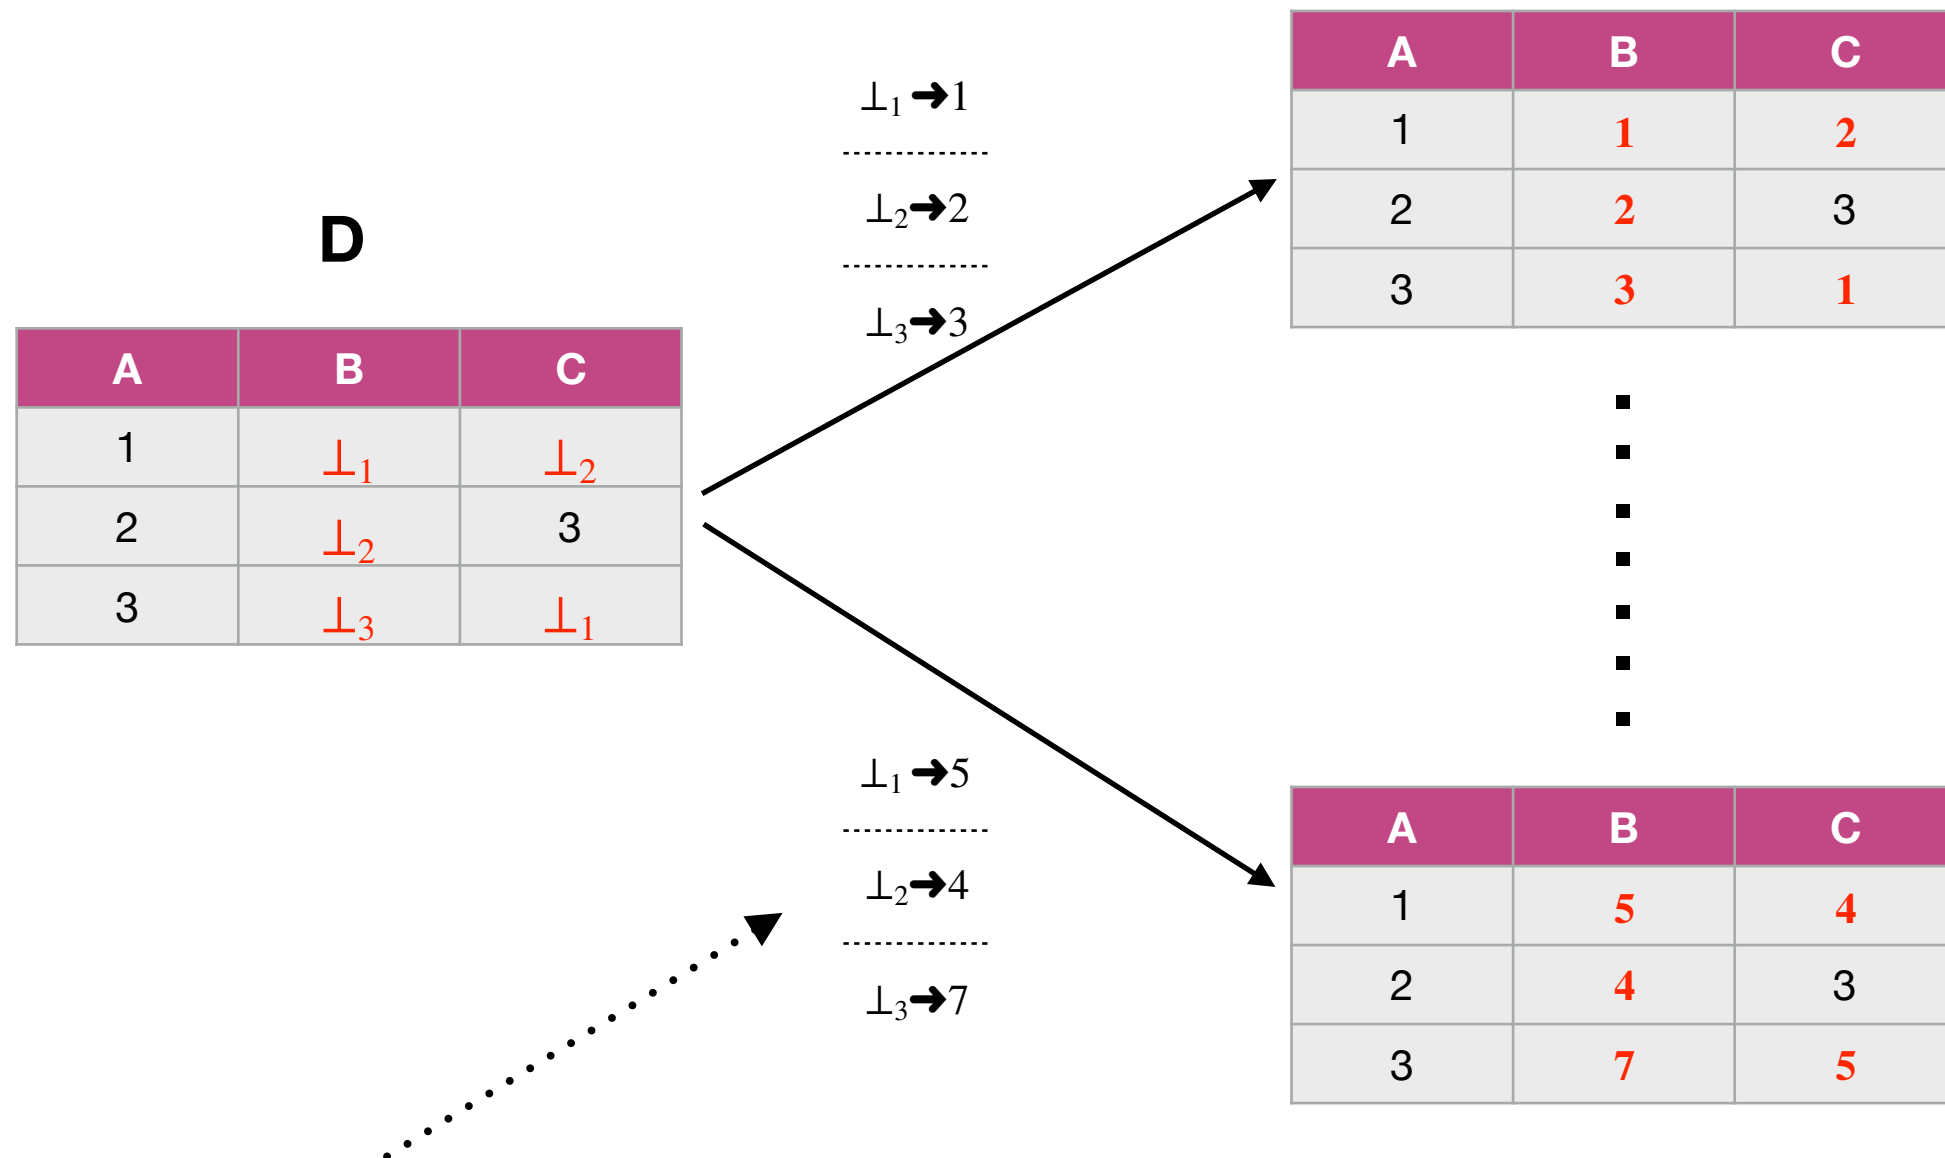
a = almost certainly true answer to Q in D

Questions

1. When is $\mu(Q, D, a) = 1$?
2. How easy is it to compute?
3. Can an answer be 50% true?
4. Is one tuple a better answer than another?

The model

Marked nulls - common in data integration, exchange, OBDA, generalize SQL nulls



Valuations **v**: Nulls \rightarrow Constants

Certain Answers

A tuple of constants **c** is a certain answer:

$c \in Q(v(D))$ for each valuation **v**

An arbitrary tuple **a** is a certain answer:

$v(a) \in Q(v(D))$ for each valuation **v**

Definition of certain answers from Lipski 1984;
unfortunately forgotten for years in favour of the constants-only definition

Certain Answers

A tuple of constants **c** is a certain answer:

$$\mathbf{c} \in Q(\mathbf{v}(\mathbf{D})) \text{ for each valuation } \mathbf{v}$$

An arbitrary tuple **a** is a certain answer:

$$\mathbf{v}(\mathbf{a}) \in Q(\mathbf{v}(\mathbf{D})) \text{ for each valuation } \mathbf{v}$$

Definition of certain answers from Lipski 1984;
unfortunately forgotten for years in favour of the constants-only definition

Support of **a**:

$$\text{Supp}(\mathbf{Q}, \mathbf{D}, \mathbf{a}) = \{ \text{valuations } \mathbf{v} \mid \mathbf{v}(\mathbf{a}) \in Q(\mathbf{v}(\mathbf{D})) \}$$

Certain Answers

Certain Answers

Support of **a**: $\text{Supp}(Q,D,a) = \{\text{valuations } v \mid v(a) \in Q(v(D))\}$

Answer **a** is **certain** \Leftrightarrow every valuation **v** is in $\text{Supp}(Q,D,a)$

Certain Answers

Support of **a**: $\text{Supp}(Q,D,a) = \{\text{valuations } v \mid v(a) \in Q(v(D))\}$

Answer **a** is **certain** \Leftrightarrow every valuation **v** is in $\text{Supp}(Q,D,a)$

Idea: answer **a** is **almost certainly true**

\Leftrightarrow a randomly chosen valuation **v** is in $\text{Supp}(Q,D,a)$

Certain Answers

Support of **a**: $\text{Supp}(Q,D,a) = \{\text{valuations } v \mid v(a) \in Q(v(D))\}$

Answer **a** is **certain** \Leftrightarrow every valuation **v** is in $\text{Supp}(Q,D,a)$

Idea: answer **a** is **almost certainly true**

\Leftrightarrow a randomly chosen valuation **v** is in $\text{Supp}(Q,D,a)$

A small problem: there are infinitely many valuations.
But techniques from zero-one laws help: look at finite approximations.

Measuring Certainty

Constants (non-nulls) = $\{c_1, c_2, c_3, \dots\}$

$Valuation_k$ = finite set of valuations with range $\subseteq \{c_1, \dots, c_k\}$

$$Supp_k(Q, D, a) = Supp(Q, D, a) \cap Valuation_k$$

Measuring Certainty

Constants (non-nulls) = $\{c_1, c_2, c_3, \dots\}$

$Valuation_k$ = finite set of valuations with range $\subseteq \{c_1, \dots, c_k\}$

$$Supp_k(Q, D, a) = Supp(Q, D, a) \cap Valuation_k$$

$$\mu_k(Q, D, a) = \frac{|Supp_k(Q, D, a)|}{|Valuation_k|} \quad (\text{a number in } [0,1])$$

Interpretation: Probability that a randomly chosen valuation with range in $\{c_1, \dots, c_k\}$ witnesses that a is an answer to Q

Measuring Certainty

$$\mu(Q, D, a) = \lim_{k \rightarrow \infty} \mu_k(Q, D, a)$$

Interpretation: Probability that a randomly chosen valuation witnesses that a is an answer to Q

Observation: the value $\mu(Q, D, a)$ does not depend on a particular enumeration of $\{c_1, c_2, c_3, \dots\}$

Zero-One Law

Zero-One Law

- **Q**: any reasonable query
 - definable in a query language such as relational algebra, datalog, second-order logic etc - formally, generic

Zero-One Law and Naive Evaluation

Zero-One Law and Naive Evaluation

- $\mu(Q,D,a) = 1 \Leftrightarrow a$ is returned by the naive evaluation of Q
 - thus almost certainly true answers are much easier to compute than certain answers
 - and naive evaluation is justified as being very close to certainty

Naive evaluation: treat nulls as values

A	B
1	\perp_1
2	\perp_1
2	\perp_2

⊖

A	B
1	\perp_2
2	\perp_1

=

A	B
1	\perp_1
2	\perp_2

Naive evaluation: treat nulls as values

A	B
1	\perp_1
2	\perp_1
2	\perp_2

 $\mathbf{-}$

A	B
1	\perp_2
2	\perp_1

 $\mathbf{=}$

A	B
1	\perp_1
2	\perp_2

Certain answer is empty because of valuations $\perp_1, \perp_2 \rightarrow c$

If the range of nulls is infinite, such valuations are **unlikely**.
Returned tuples are **almost** certainly true answers - but not certain.

Naive evaluation: treat nulls as values

A	B
1	\perp_1
2	\perp_1
2	\perp_2

 \neq

A	B
1	\perp_2
2	\perp_1

 $=$

A	B
1	\perp_1
2	\perp_2

Certain answer is empty because of valuations $\perp_1, \perp_2 \rightarrow c$

If the range of nulls is infinite, such valuations are **unlikely**.
Returned tuples are **almost** certainly true answers - but not certain.

In general, **naive evaluation \neq certain answers**. Exceptions:

- unions of conjunctive queries
- their extension with **$Q \div R$** where **R** is a relation

Naive evaluation: treat nulls as values

A	B
1	\perp_1
2	\perp_1
2	\perp_2

 \neq

A	B
1	\perp_2
2	\perp_1

 $=$

A	B
1	\perp_1
2	\perp_2

What if:

1. **A** is a key of the second relation (which forces $\perp_1 = \perp_2$), or
2. there is a restriction on the range of **B**?

The reasoning that valuations $\perp_1, \perp_2 \rightarrow c$ are unlikely no longer works

This is due to the presence of **constraints**.

Certainty with constraints

- Only interested in databases satisfying integrity constraints Σ - for example, keys or foreign keys
- Standard approach: find certain answers to $\Sigma \rightarrow Q$
- Not very informative (because $\Sigma \rightarrow Q$ is $\neg\Sigma \vee Q$)
 - if $\mu(\Sigma, D) = 0$, then $\mu(\Sigma \rightarrow Q, D, a) = 1$
 - if $\mu(\Sigma, D) = 1$, then $\mu(\Sigma \rightarrow Q, D, a) = \mu(Q, D, a)$

Certainty with constraints

- A better idea: use **conditional probability** $\mu(Q \mid \Sigma, D, a)$
 - probability that a randomly chosen valuation that satisfies Σ also witnesses that a is an answer to Q
- Still defined as a limit since there are infinitely many valuations

Measuring certainty with constraints

$$Supp_k(Q, D, a) = \{ \text{valuations } v \in Valuation_k \mid v(a) \in Q(v(D)) \}$$

$$\mu_k(Q \mid \Sigma, D, a) = \frac{|Supp_k(Q \wedge \Sigma, D, a)|}{|Supp_k(\Sigma, D, a)|}$$

Interpretation: Probability that a randomly chosen valuation with range in $\{c_1, \dots, c_k\}$ that witnesses constraints Σ also witnesses that a is an answer to Q

Measuring certainty with constraints

$$\mu(Q \mid \Sigma, D, a) = \lim_{k \rightarrow \infty} \mu_k(Q \mid \Sigma, D, a)$$

Interpretation: Probability that a randomly chosen valuation that witnesses constraints Σ also witnesses that a is an answer to Q

Observation: the value $\mu(Q \mid \Sigma, D, a)$ does not depend on a particular enumeration of $\{c_1, c_2, c_3, \dots\}$

Zero-One Law fails with constraints

- Database **D**: $R = \{\perp\}$, $S = \{1\}$, $U = \{1,2\}$
- Constraint: $R \subseteq U$
- Query **Q**: is $R \subseteq S$?
- $\mu(Q \mid \Sigma, D, a) = 0.5$

What if zero-one fails?

- The best next thing: **convergence**
- Consider, for example, **ordered** graphs.
- Zero-one law fails: $\mu(\text{edge between the smallest and the largest element}) = 0.5$
- But $\mu(\alpha)$ exists for every first-order α
 - and is a rational of the form $n/2^m$ (Lynch 1980)

Convergence with constraints

Convergence with constraints

- **Q**: any reasonable query, **Σ** : any reasonable constraints (both **generic**)

Convergence with constraints

- **Q**: any reasonable query, **Σ** : any reasonable constraints (both **generic**)
- **Theorem**: $\mu(Q \mid \Sigma, D, a)$ always exists
 - $\mu(Q \mid \Sigma, D, a)$ is a rational number between 0 and 1

Convergence with constraints

- **Q**: any reasonable query, **Σ** : any reasonable constraints (both **generic**)
- **Theorem**: $\mu(Q \mid \Sigma, D, a)$ always exists
 - $\mu(Q \mid \Sigma, D, a)$ is a rational number between 0 and 1
- Every rational number in **[0,1]** can appear as $\mu(Q \mid \Sigma, D, a)$ for a conjunctive query **Q** and inclusion constraints **Σ**

Convergence with constraints

Convergence with constraints

- Computing $\mu(Q \mid \Sigma, D, a)$

Convergence with constraints

- Computing $\mu(Q \mid \Sigma, D, a)$
- It is a rational number, so need a function complexity class

Convergence with constraints

- Computing $\mu(Q \mid \Sigma, D, a)$
- It is a rational number, so need a function complexity class
- It can be computed in $FP^{\#P}$
 - functions computable in polynomial time with access to $\#P$ oracle

Convergence with constraints

- Computing $\mu(Q \mid \Sigma, D, a)$
- It is a rational number, so need a function complexity class
- It can be computed in $FP^{\#P}$
 - functions computable in polynomial time with access to $\#P$ oracle
- Computing $\mu(Q \mid \Sigma, D, a)$ could be hard for $FP^{\#P}$
 - under the appropriate definition of hardness for function classes

Constraints and zero-one laws

- Zero-one law still holds for some constraints, e.g., **functional dependencies**
- Σ : a set of functional dependencies.
- Known: if Q is a **conjunctive query**, then **certain answers under $\Sigma = Q(\text{chase}(D, \Sigma))$**
- If Q is an **arbitrary query**, then **almost certainly true answers under $\Sigma = Q(\text{chase}(D, \Sigma))$**
 - $\mu(Q \mid \Sigma, D, a) = \mu(Q, \text{chase}(D, \Sigma), a)$

Qualitative Measures

- We can also use supports $\text{Supp}(Q,D,a)$ to define qualitative measures:

Qualitative Measures

- We can also use supports $\text{Supp}(Q,D,a)$ to define qualitative measures:
 - a is at least as good an answer as b , to query Q if $\text{Supp}(Q,D,b) \subseteq \text{Supp}(Q,D,a)$

Qualitative Measures

- We can also use supports $\text{Supp}(Q,D,a)$ to define qualitative measures:
 - a is at least as good an answer as b , to query Q if $\text{Supp}(Q,D,b) \subseteq \text{Supp}(Q,D,a)$
 - a is a better answer than b , to query Q if $\text{Supp}(Q,D,b) \subsetneq \text{Supp}(Q,D,a)$

Qualitative Measures

- We can also use supports $\text{Supp}(Q,D,a)$ to define qualitative measures:
 - a is at least as good an answer as b , to query Q if $\text{Supp}(Q,D,b) \subseteq \text{Supp}(Q,D,a)$
 - a is a better answer than b , to query Q if $\text{Supp}(Q,D,b) \subsetneq \text{Supp}(Q,D,a)$
 - a is a best answer to Q if there is no better answer

Qualitative measure: example

A	B
1	\perp_1
2	\perp_1
2	\perp_2

—

A	B
1	\perp_2
2	\perp_1

Qualitative measure: example

A	B
1	\perp_1
2	\perp_1
2	\perp_2

—

A	B
1	\perp_2
2	\perp_1

- No **certain** answers
- Naive evaluation gives **(1, \perp_1)** and **(2, \perp_2)**
- **(2, \perp_2)** is a **better answer** than **(1, \perp_1)**
- **Best answer** = **(2, \perp_2)**

Qualitative measure: example

A	B
1	\perp_1
2	\perp_1
2	\perp_2

—

A	B
1	\perp_2
2	\perp_1

- No **certain** answers
- Naive evaluation gives **(1, \perp_1)** and **(2, \perp_2)**
- **(2, \perp_2)** is a **better answer** than **(1, \perp_1)**
- **Best answer** = **(2, \perp_2)**

Unlike certain answers, best answers always exist

Qualitative measures: complexity

- Fix a query Q of relational algebra/calculus
- Input: database D , tuples a and b

Qualitative measures: complexity

- Fix a query **Q** of relational algebra/calculus
- Input: database **D**, tuples **a** and **b**
- Is **a** at least as good as **b**? **coNP-complete**

Qualitative measures: complexity

- Fix a query **Q** of relational algebra/calculus
- Input: database **D**, tuples **a** and **b**
- Is **a** at least as good as **b**? **coNP-complete**
- Is **a** better than **b**? **DP-complete**

Qualitative measures: complexity

- Fix a query **Q** of relational algebra/calculus
- Input: database **D**, tuples **a** and **b**
- Is **a** at least as good as **b**? **coNP-complete**
- Is **a** better than **b**? **DP-complete**
- Identify the set of best answers: **P^{NP}[log n]-complete**

Qualitative measures: complexity

- Fix a query **Q** of relational algebra/calculus
- Input: database **D**, tuples **a** and **b**
- Is **a** at least as good as **b**? **coNP-complete**
- Is **a** better than **b**? **DP-complete**
- Identify the set of best answers: **P^{NP}[log n]-complete**
- For unions of conjunctive queries, all in **PTIME**.
 - Does not go via naive evaluation; the algorithm is of very different nature

Quantitative vs qualitative

- Quantitative: $\mu(Q,D,a) = 0$ or $\mu(Q,D,a) = 1$
- Qualitative: **best** or **not best**
- All 4 combinations are possible, even for first-order queries

Summary

Almost certainly true answers to queries on incomplete databases

- **Very good - though not always perfect**
- **Much better than certain answers computationally**

Summary

Almost certainly true answers to queries on incomplete databases

- **Very good - though not always perfect**
- **Much better than certain answers computationally**

Future

Summary

Almost certainly true answers to queries on incomplete databases

- **Very good - though not always perfect**
- **Much better than certain answers computationally**

Future

- **Other data models**
- **SQL nulls and SQL evaluation**
- **Other distributions**
- **Applications of certain answers (integration, exchange, etc)**