



Http://act.buaa.edu.cn

北京航空航天大学计算机新技术研究所

The Institute of Advanced Computing Technology



Computing Issues for Big Data

Theory, Systems, and Applications

Beihang University
Chunming Hu (hucm@buaa.edu.cn)

Big Data Summit, with CyberC 2013
October 10, 2013. Beijing, China.



TI

Institute of Advanced Computing Technology



Bio of Myself


- Chunming Hu
 - Got my Ph.D in 2006, Beihang University
 - Associate Professor at Institute of Advanced Computing Technology (ACT), School of Computer Science, Beihang University
- Research Interests
 - Service Computing (2001-2008)
 - Grid Computing (2005-2009)
 - Cloud Infrastructure and System Virtualization (2008-)
 - Network System Virtualization
 - Cloud-Client Computing
 - Distributed Systems for Data Processing





Agenda

- Understanding the Big Data
 - Background
 - Computing Issues (4V \rightarrow 3I)
- Big Data Research at Beihang University
 - RCBD, lead by Wenfei Fan
 - 973 Project on Big Data, lead by Beihang University
- Early Experience on Big Data
 - BD-Tractable by Preprocessing
 - Performance Model for Hadoop
 - Distributed Graph Pattern Matching
 - Event Early Detection via Social Data
- Summary





Big Data in Cyberspace

• Large Scale with Rapid Updates

Social Networks

- 4 Micro-blogger Provider in China:
- 800M Users, 200M tweets everyday, 20M+ Photos.

Internet Search

- **Baidu**: 1PB log data per Day. Handling 1000PB
- **Google**: Processing 20PB data everyday


1PB data in DVD:

~25km

1ZB=1PB × 10⁶

IDC Report


- Data doubled every 18 months
- Data in Cyberspace
- IDC Report :
 - 2009: 0.8ZB
 - 2012: 2.7 ZB
 - 2020(E): 35ZB




Airplane
15,000m

Chomolung
ma 8,800m


4V Features in Big Data




VOLUME



VELOCITY



VARIETY



VALUE

Volume



- In PB or EB
- Distributed data

Velocity



- Dynamic Changes
- Updated constantly

Variety



- Heterogeneous
- Semi-structured or unstructured

Value



- Biz opportunity
- Sensitive Data

The Economist



Information has gone from scarce to superabundant. That brings huge new benefits, says Kenneth Cukier, but also his headache

Data Deluge

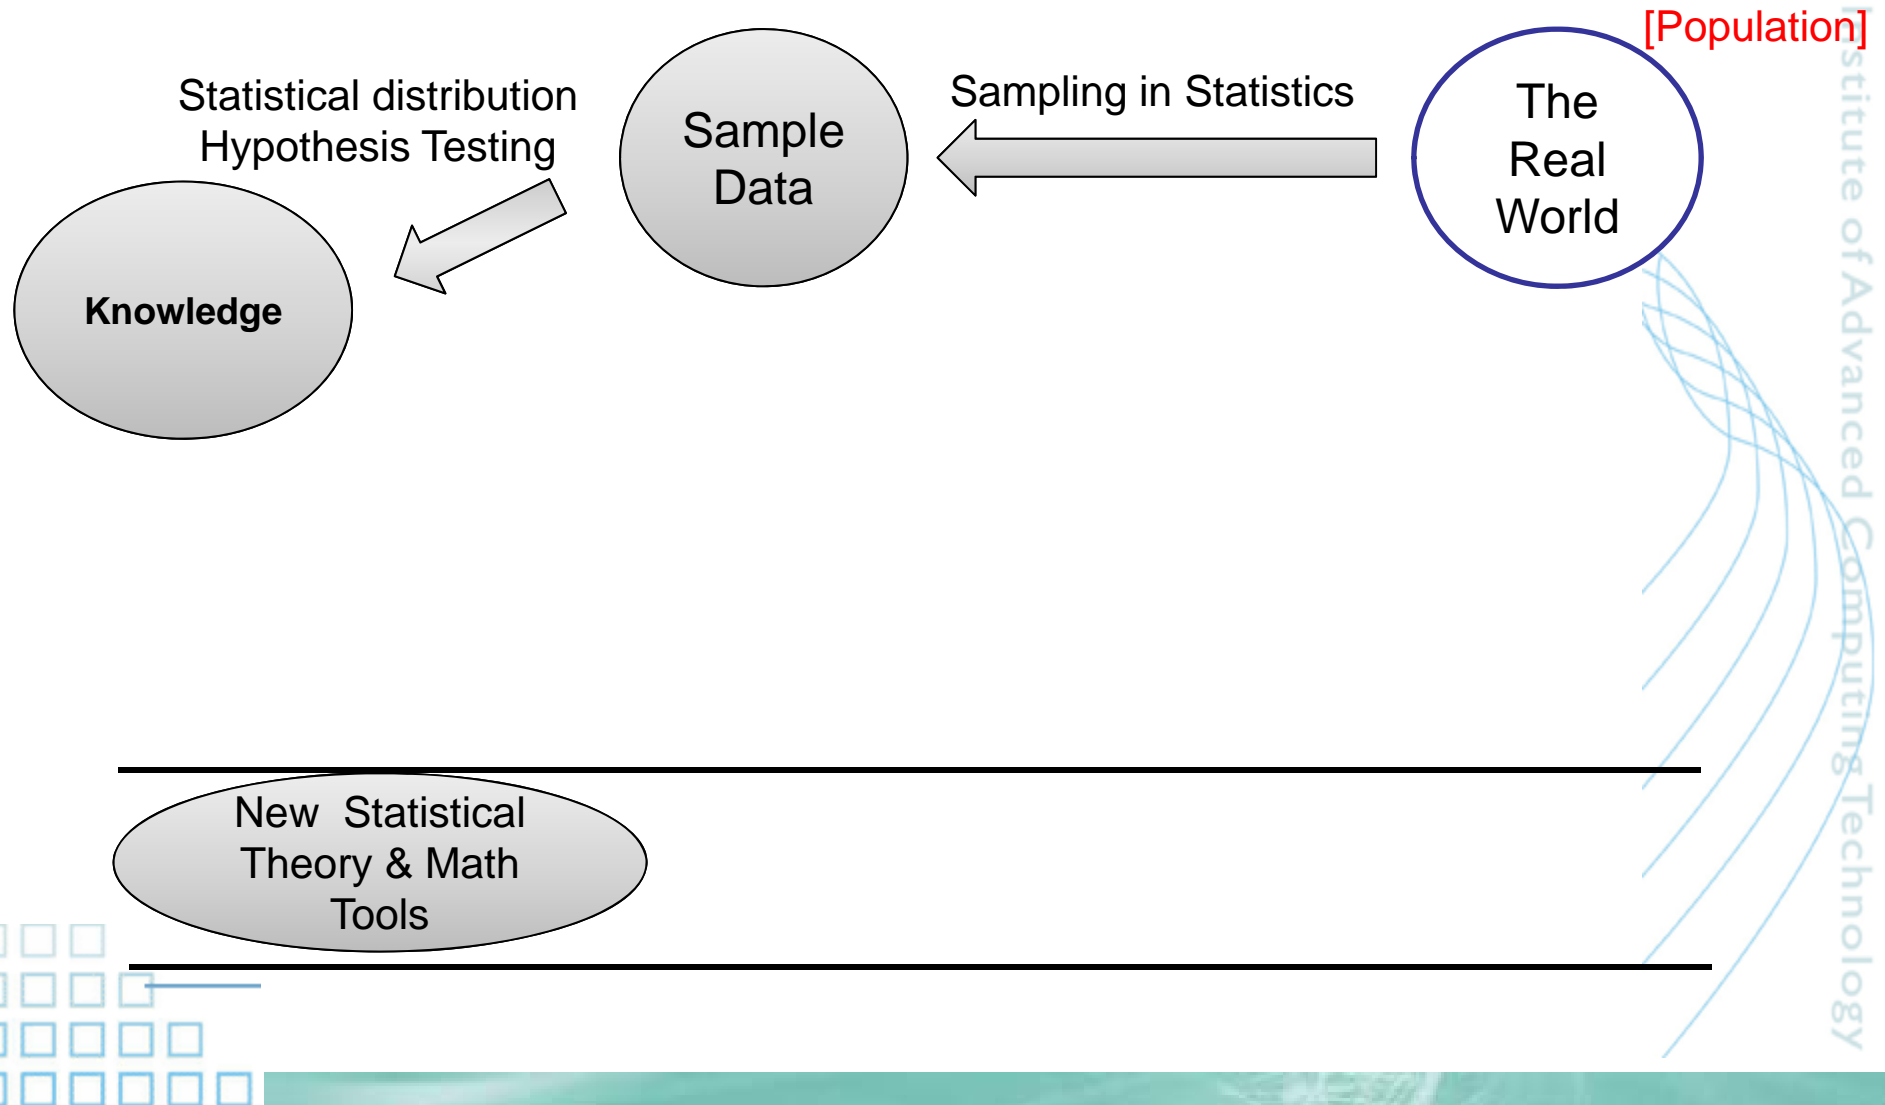


Wikipedia

large and complex datasets, which is quite difficult to process using existing data management tools, and traditional data processing applications


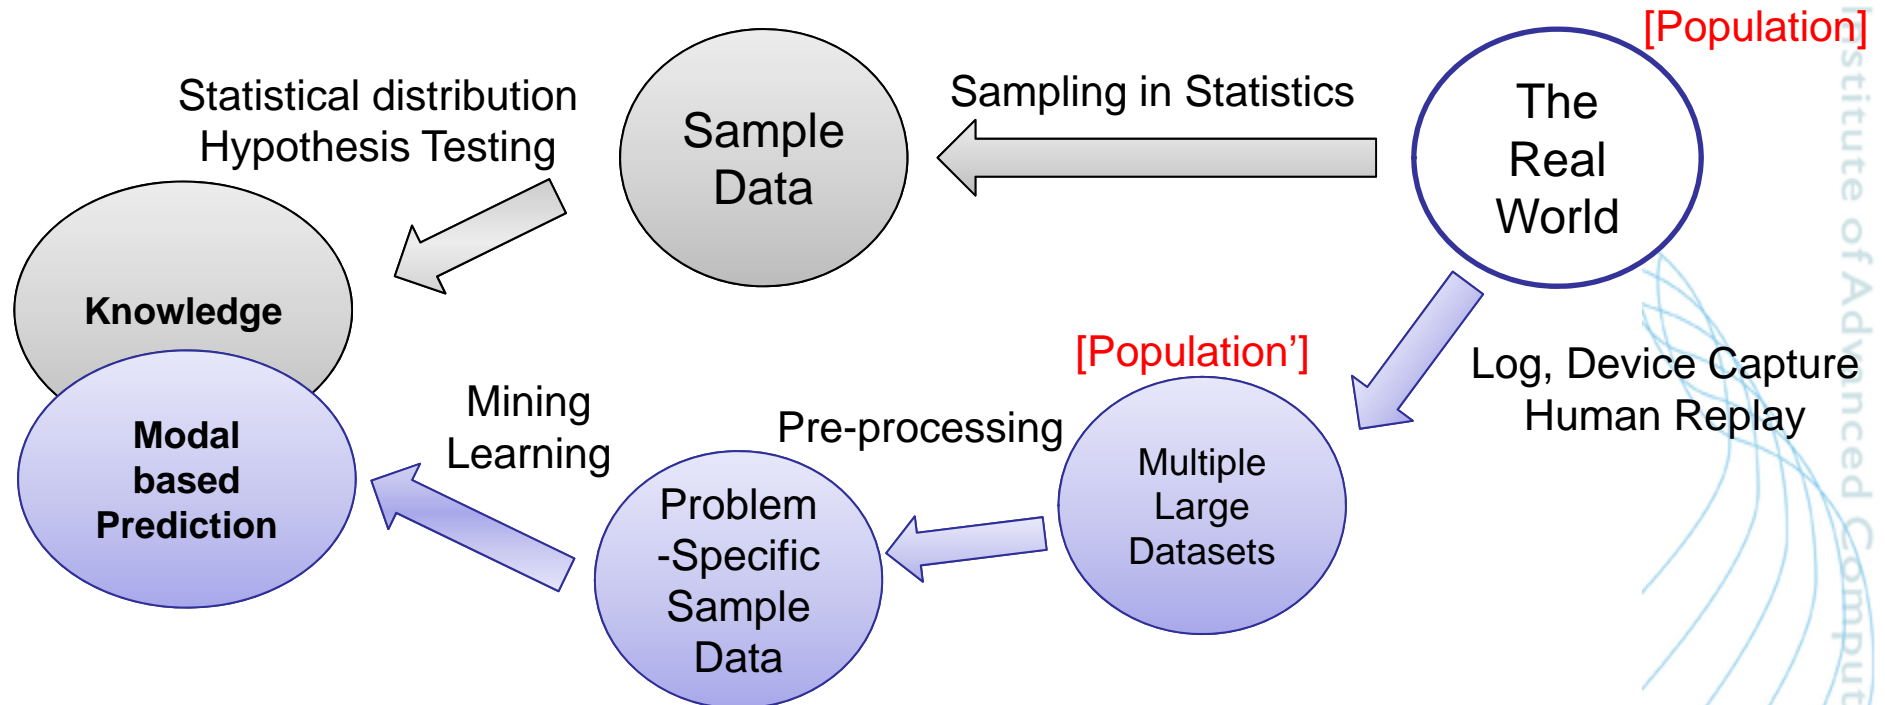


Big Data Changes



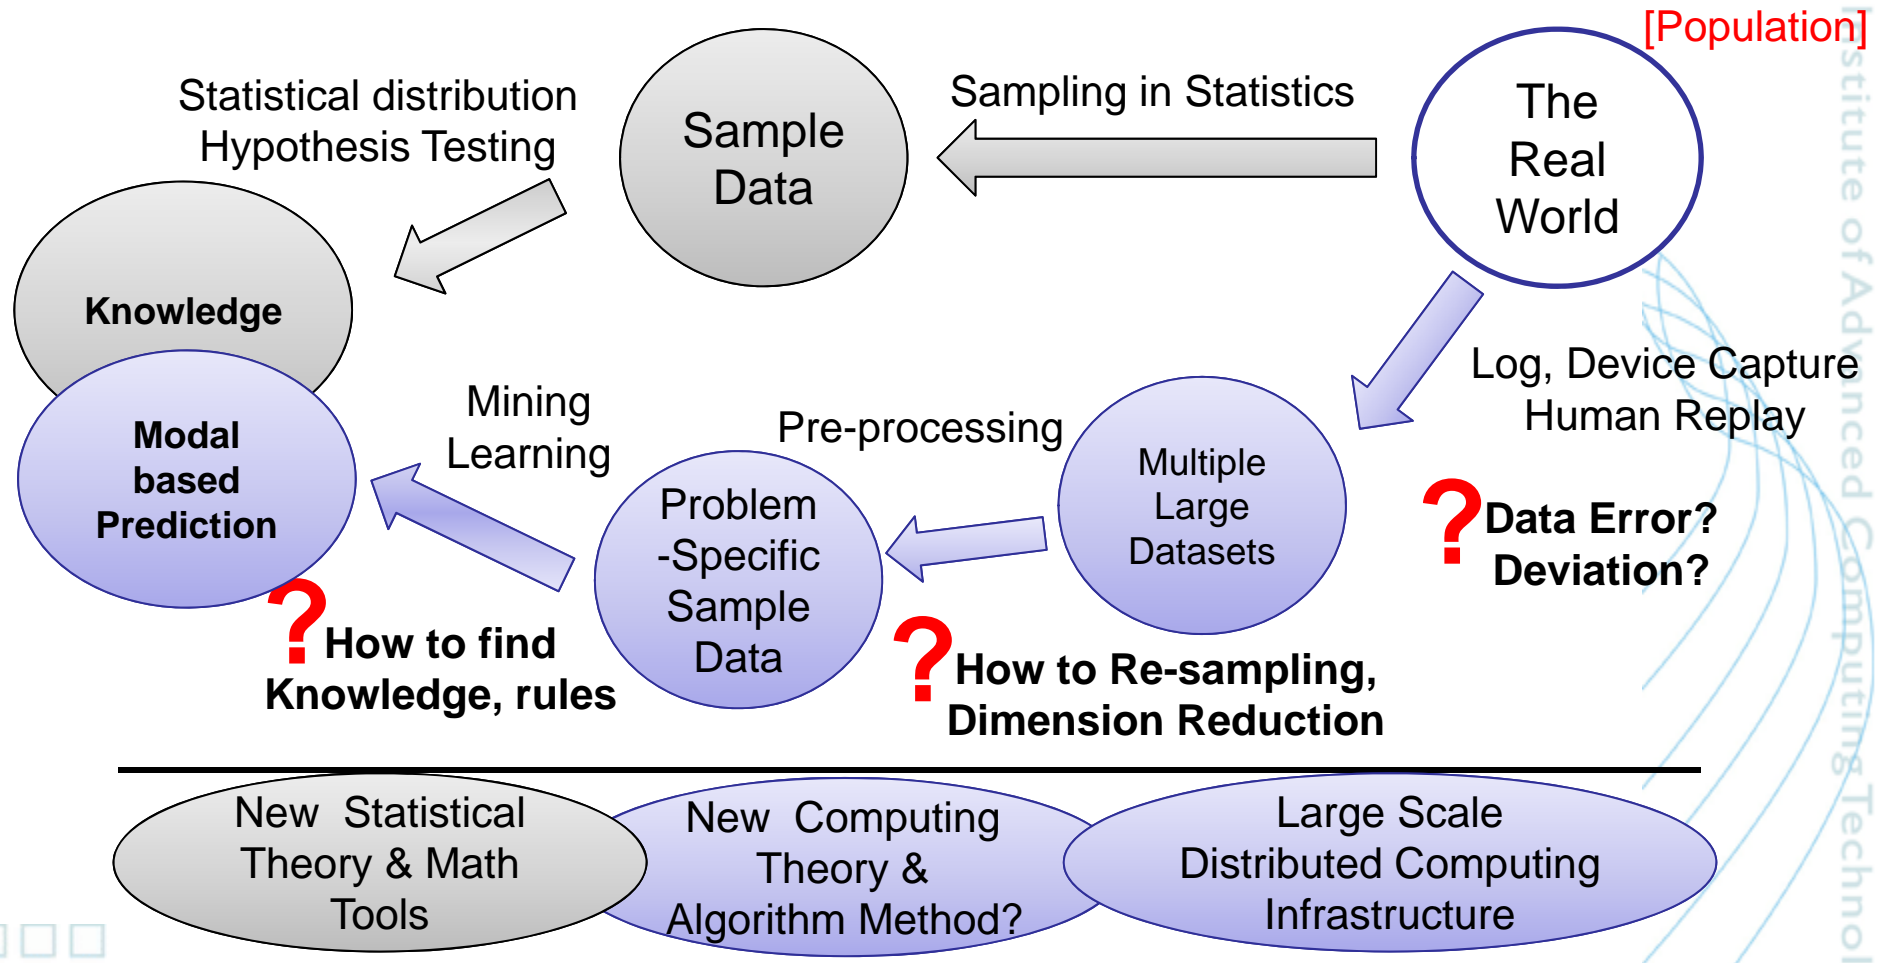


Big Data Changes

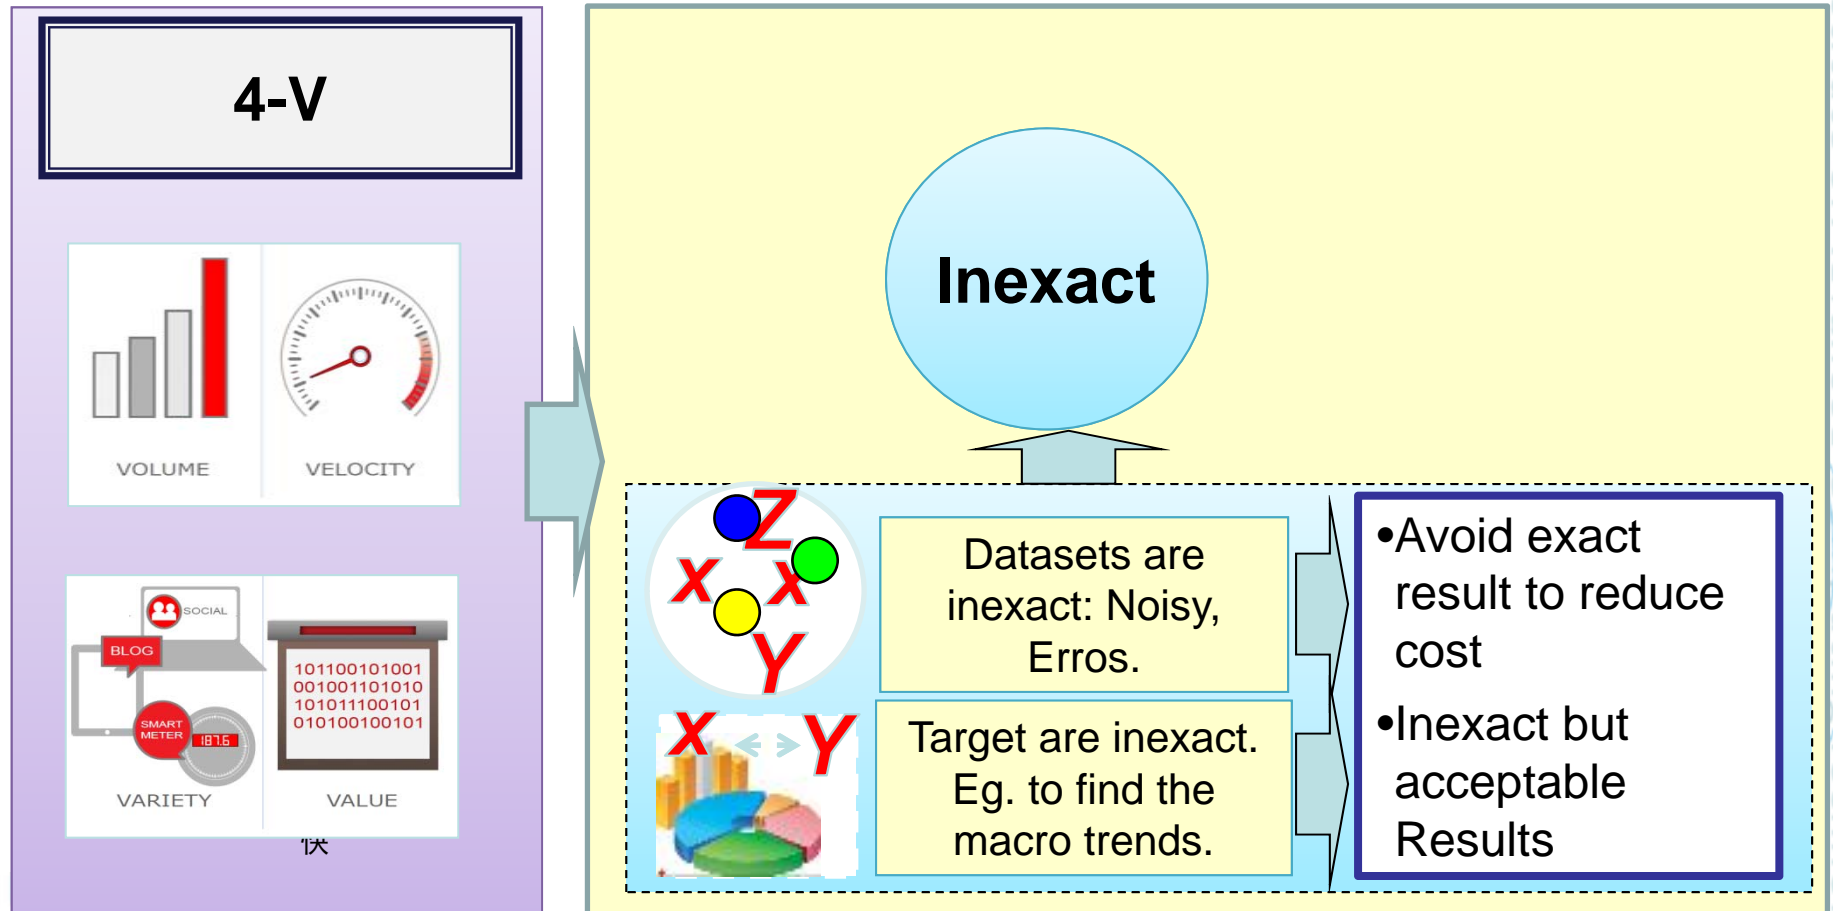




Big Data Changes

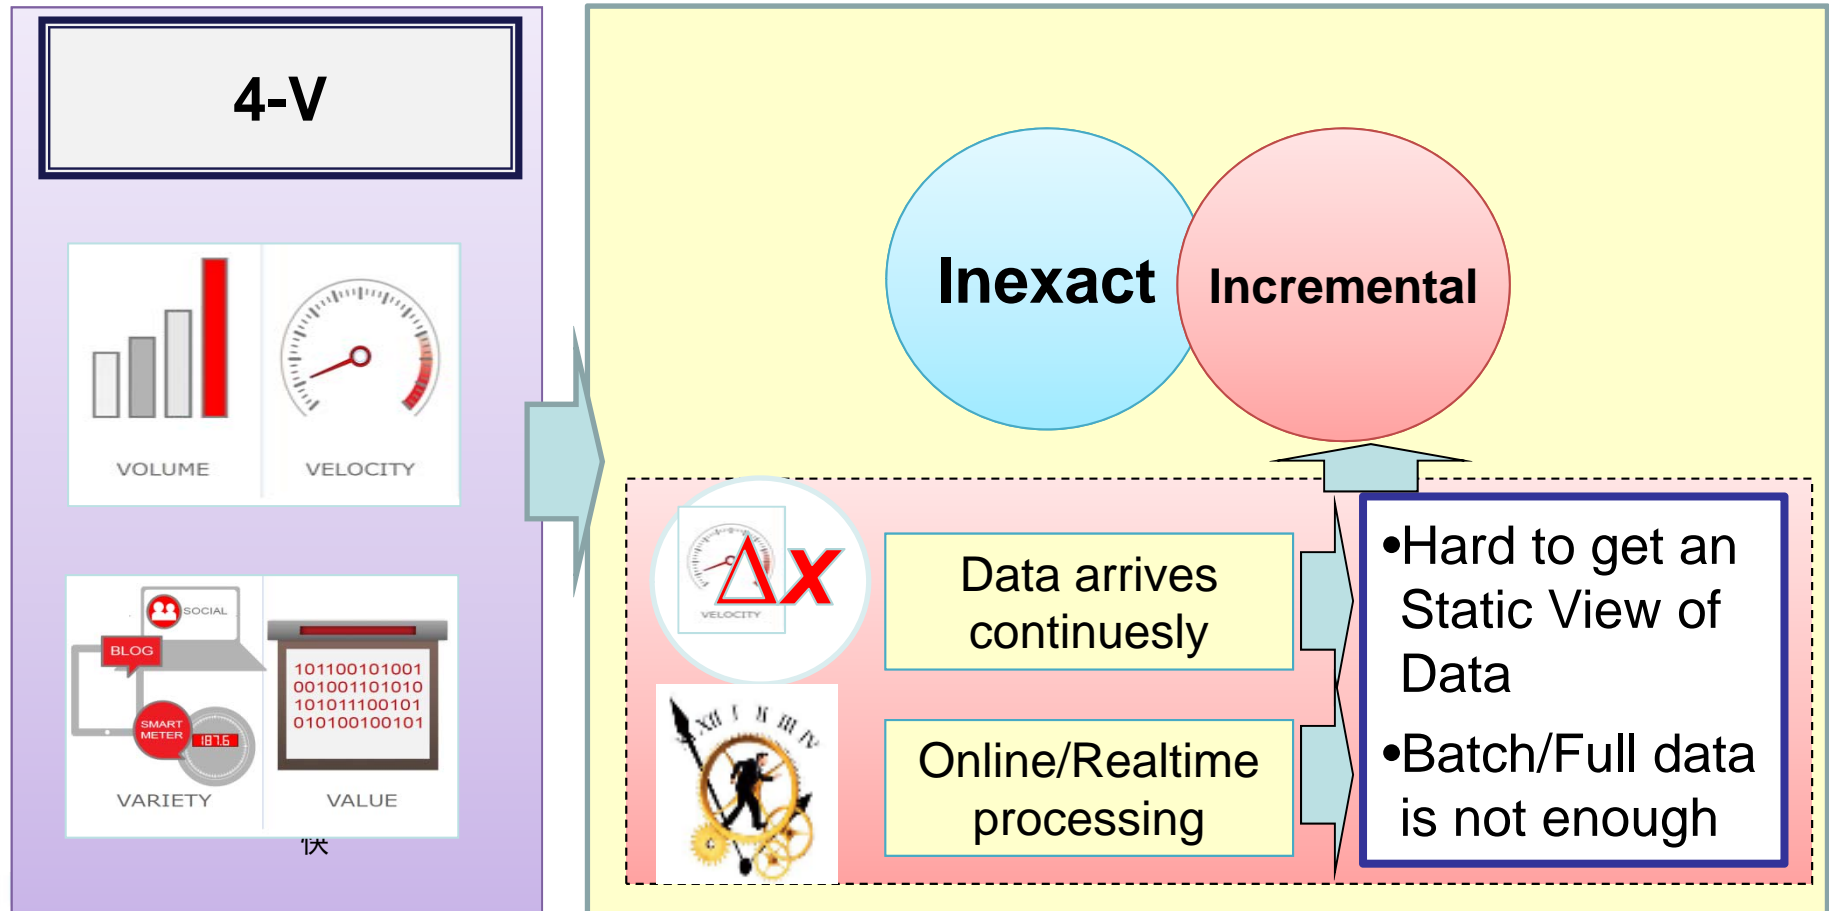


Focusing on the Computing







Focusing on the Computing





Focusing on the Computing

4-V



伏

Features of Big Data Computing

Inexact

Incremental

Inductive


Multi-source
Datasets

References
between
Datasets

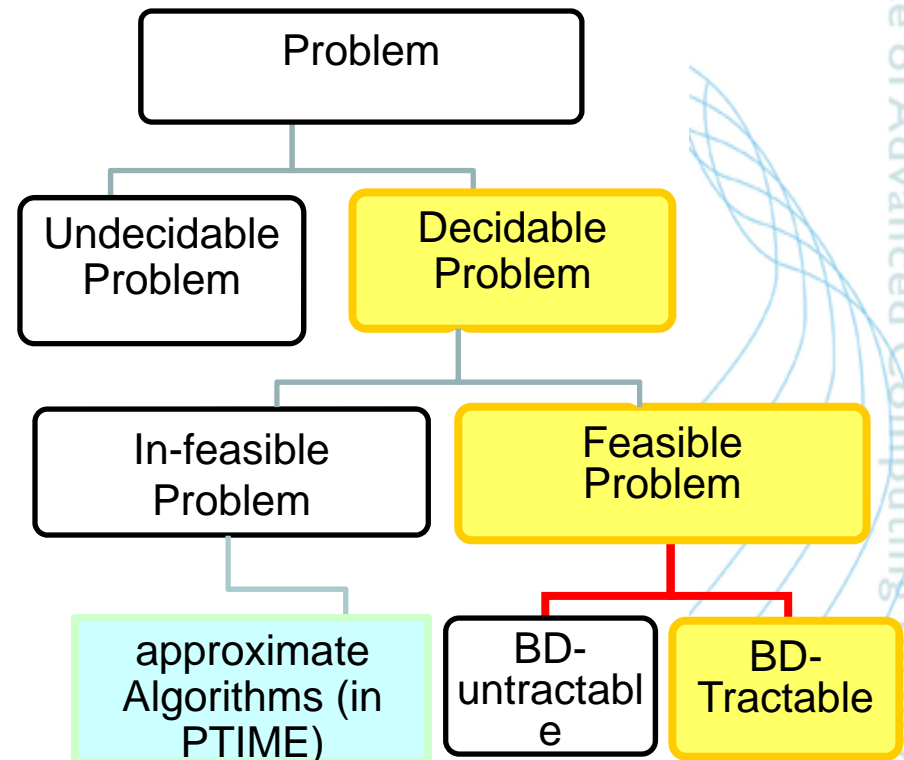
- Use the data correlations to adjust the errors
- Transfer Learning

Questions on Big Data Computing

- Question 1: Is there new Theory for Big Data Computing?



- **Good**: PTIME
- **Bad**: NP-Hard
- **Ugly**: PSPACE-hard, or EXPTIME-hard, undecidable




Questions on Big Data Computing

- Question 2: Is Hadoop a must to data processing?
 - Different Computing requirements, User Scenarios
 - Different Algorithm Design methods


- MapReduce (OSDI 2004)

$x \rightarrow x_1, x_2, x_3 \dots x_n$

- Handling All Data in a Distributed way



Incremental




MR is not the only solution

- Incremental Computing: Percolator by Google (OSDI 2010)



3I



New Algorithm Methods

- Resampling
- Query preserving data compression
- Partial evaluation and distributed processing
- Top-k query and terminating... ..

Questions on Big Data Computing


- Question 3: How to handle the data computing algorithms in an operatable manner?
 - Application specific Features Analysis
 - Data Pattern, Arriving Rate, Query, ...
 - General Purpose vs. Specific Purpose
 - Domain-specific Knowledges
 - Data Mining and ML Methods
 - Distributed system
 - Offline/Online
 - Batch/Incremental/Streaming
 - In-memory Computing





Agenda

- Understanding the Big Data
 - Background
 - Computing Issues (4V \rightarrow 3I)
- **Big Data Research at Beihang University**
 - RCBD, lead by Wenfei Fan
 - 973 Project on Big Data, lead by Beihang University
- Early Experience on Big Data
 - BD-Tractable by Preprocessing
 - Performance Model for Hadoop
 - Distributed Graph Pattern Matching
 - Event Early Detection via Social Data
- Summary





RCBD: International Research Centre on Big Data

International Research Centre on Big Data (Founded in Sept 2012)



Beihang U.



U. Edinburgh



HKUST



U. Pennsylvania



Baidu

广州市香港科大
霍英东研究院



HKUST Fok Ying Tung
Graduate School



NOAH'S ARK LAB

From Big Data to Deep Knowledge





973 Project on Big Data

- Computing Theory on Big Data (2014-2018)
 - a 973 Project Funding from Ministry of Science & Technology, with 8 institutional organization

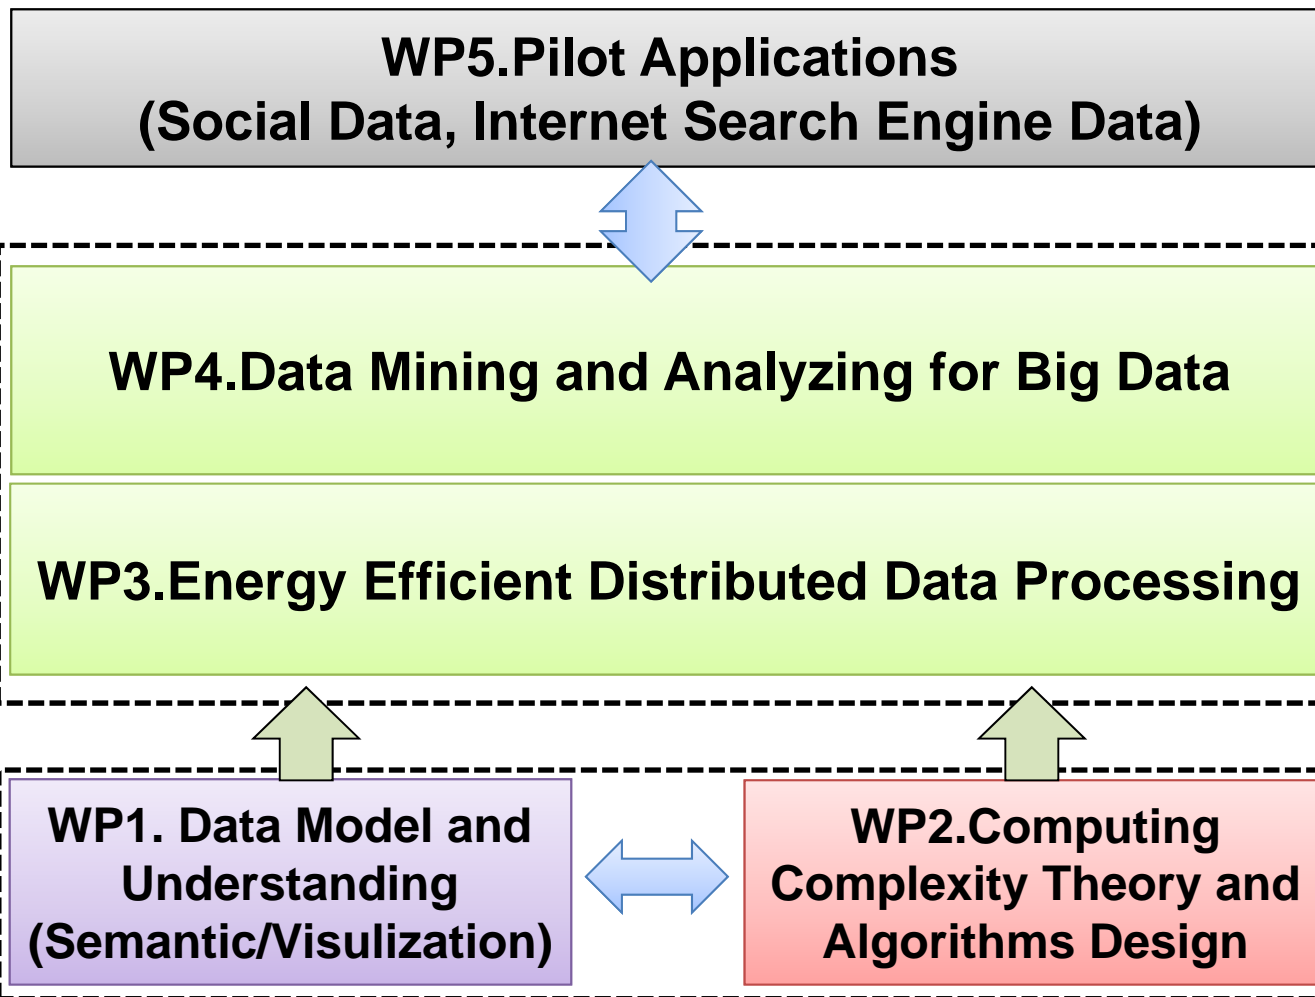


- Focusing on the Computing Issue of Big Data Processing






973 Project on Big Data



Agenda

- Understanding the Big Data
 - Background
 - Computing Issues (4V \rightarrow 3I)
- Big Data Research at Beihang University
 - RCBD, lead by Wenfei Fan
 - 973 Project on Big Data, lead by Beihang University
- **Early Experience on Big Data**
 - BD-Tractable by Preprocessing
 - Performance Model for Hadoop
 - Distributed Graph Pattern Matching
 - Event Early Detection via Social Data
- Summary




Some Early Experiences

- Theory
 - BD-Tractable via Data Preprocessing
- Systems
 - Performance Model for Hadoop
 - Graph Pattern Matching
- Applications
 - Using Mood to Detect Sudden Occurrence (Early Event Detection via Social Data)



BD-Tractable with Preprocessing

- *Polynomial time queries become **intractable** on big data*
- We want to be able to tell what queries are feasible on big data



BD-tractable queries: queries feasible on big data



BD-Tractable with Preprocessing

- How do we dealing with SQL querys on a large DATABASE?
 - Scan through all the records? NO!!
 - Using Index to get better query performance!
 - B-Tree index, from $O(n)$ to $O(\log n)$
 - Query Optimizations!
- Two steps of computing
 - Set up the “index”: preprocessing
 - Doing query on the “index”



BD-Tractable with Preprocessing

A class Q of queries is **BD-tractable** if there exists a **PTIME** preprocessing function Π such that

- ✓ for any database D on which queries of Q are defined,


$$D' = \Pi(D)$$

hence D' is of polynomial size

- ✓ for any query Q defined on D , $Q(D)$ can be computed on D' in **parallel polylog time (NC)**

possible rewriting

parallel $\log^k(|D|, |Q|)$



Does it work? If a linear scan of D could be done in $\log(|D|)$ time:

- ✓ **15 seconds** when D is of 1 PB instead of **1.99 days**
- ✓ **18 seconds** when D is of 1 EB rather than **5.28 years**

BD-tractable queries are feasible on big data

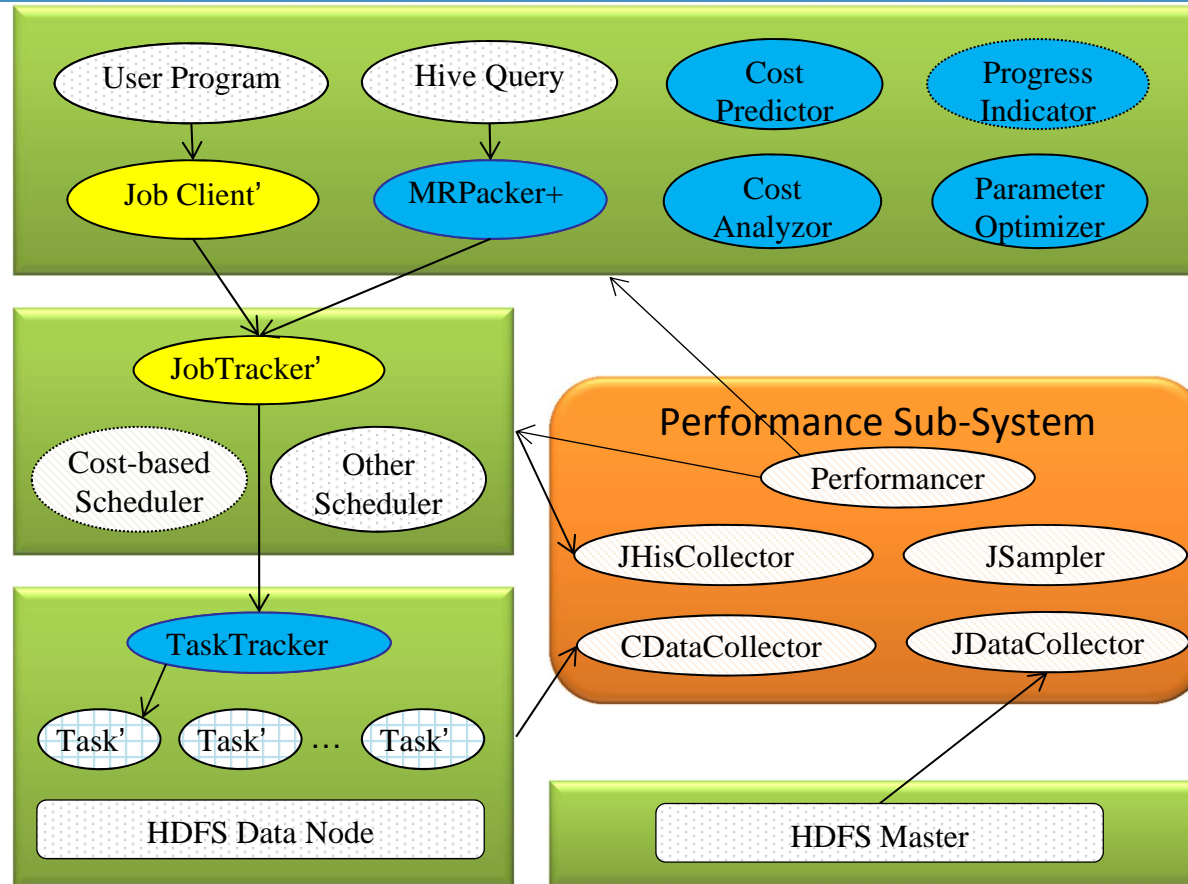
Performance Model for Hadoop

- Accurate performance prediction can help optimize the MR jobs
 - Cost-based scheduling strategies
 - Query Optimization
- Multi-processing steps in MR
- Basic Idea:
 - Benchmarking and Profiling on target machine
 - Find out the parameters for the model






Performance Model for Hadoop



Distributed Graph Pattern Matching

- Graph pattern matching




- Providing evaluation algorithms and optimizations for graph simulation in a distributed setting

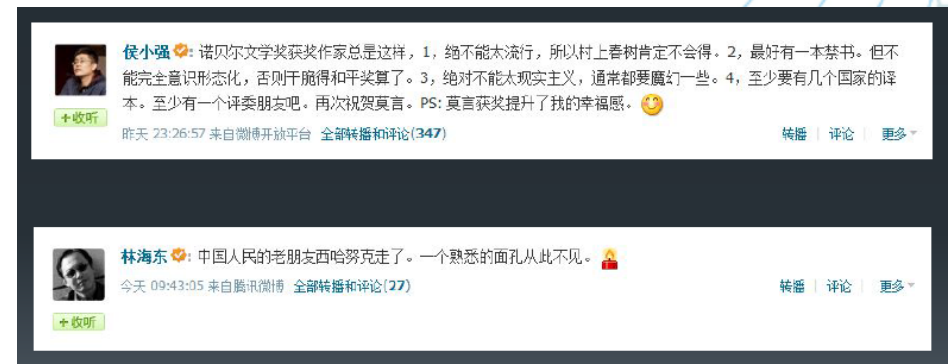


Early Event Detection via Social Data

- Social Data reflect the physical world
 - Event Detection via Social Data




- Motion plays important role in social media
 - How to detect the user motion through the weibo text??



Early Event Detection via Social Data

- Classification
 - 95 motion icons selected from 1000 icons
 - Use the text with motion icons as the training sets

Sentiment	#Emoticons	Typical emoticons
Angry	9	    
Disgusting	14	    
Joyful	50	    
Sad	22	    

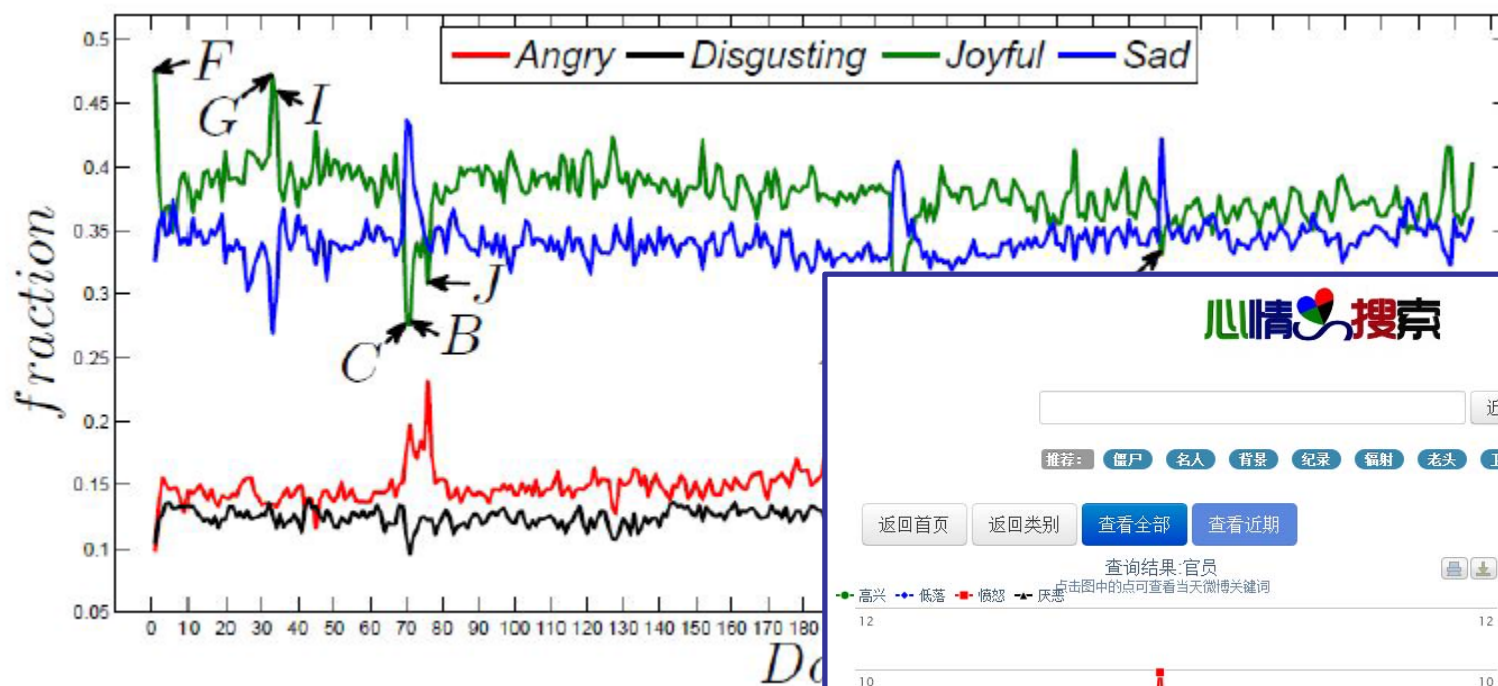


(a) Hourly pattern.



Early Event Detection via Social Data

- Abnormal event detection



- Mood Search

<http://xinqings.nlsde.buaa.edu.cn>





iCOME: a Competition for Big Data

<http://www.icome.org.cn>

中国云·移动互联网
创新大奖赛

设为首页 | 加入收藏

首页 关于竞赛 竞赛题目 通知公告 常见问题 参赛报名 作品提交 作品展示 竞赛平台 论坛

中国云·移动互联网创新大奖赛
Internet Contest for Cloud & Mobile computing

玩转大数据

大赛团队登录
账号
密码
忘记密码? 注册

参赛流程
1 初赛报名 — 2 赛题选定 — 3 作品提交 — 4 专家评审
5 现场答辩 — 6 颁奖典礼

大赛简介
中国云·移动互联网创新大奖赛 (Internet contest for Cloud & Mobile computing, 简称iCOME) 是由云联盟倡议并支持的云计算及移动互联网技术与应用创新竞赛...

通知公告 more

- 参赛选手须知 (重要) 2013/09/26
- 测试结果 2013/09/23
- 大数据时代重新定义隐私 2013/09/27
- 刘经南院士: 大数据时代的泛在... 2013/09/26

 数据集 more
 数据集, 又称为资料集、数据集合或资料集合, 是指一种由数据所组成的集合。

customer center help@icome.org.cn

作品提交.....7月-10月初
颁奖典礼.....10月下旬


【主办单位】 CCA 中国云联盟

【协办单位】 Baidu 百度, China Unicom 中国联通, Longfor 龙湖地产, CYBERPORT 数码港, 用友, CBC 宽带资本, 阿里云, Tencent 腾讯, Lenovo 联想, TCL, 云天基金, 天富大数据, 大数据实验室, 投资

Agenda


- Understanding the Big Data
 - Background
 - Computing Issues (4V \rightarrow 3I)
- Big Data Research at Beihang University
 - RCBD, lead by Wenfei Fan
 - 973 Project on Big Data, lead by Beihang University
- Early Experience on Big Data
 - BD-Tractable by Preprocessing
 - Performance Model for Hadoop
 - Distributed Graph Pattern Matching
 - Event Early Detection via Social Data

- **Summary**



Summary

- Big Data: from 4V to 3I
 - Inexact
 - Incremental
 - Inductive
- Application-Driven Vertical Integration
 - Theory, Algorithm, Distributed Systems, Mining & ML Methods
- Open Data Community
 - Get more data from industry/government





Acknowledgement

- **Part of the slides borrowed from**

- Prof. Wenfei Fan at RCBd,
- Prof. Ke Xu at NLSDE, Beihang University
- Prof. Shuai Ma, Dr. Xuelian Lin at ACT, Beihang University

- **References**

- Wenfei Fan, Floris Geerts, Frank Neven. Making Queries Tractable on Big Data with Preprocessing, VLDB 2013.
- Lin Xuelian, Meng Zide, Xu Chuan and Wang Meng. A Practical Performance Model for Hadoop MapReduce. IEEE International Conference on Cluster Computing Workshops (ClusterW), 2012
- Shuai Ma, Yang Cao, Jinpeng Huai, Tianyu Wo. Distributed Graph Pattern Matching, WWW 2012.
- Shuai Ma, Yang Cao, Wenfei Fan, Jinpeng Huai, Tianyu Wo. Capturing Topology in Graph Pattern Matching, VLDB 2012.
- Jichang Zhao, Li Dong, Junjie Wu, Ke Xu: MoodLens: an emoticon-based sentiment analysis system for chinese tweets. KDD 2012.





Http://act.buaa.edu.cn

北京航空航天大学计算机新技术研究所

The Institute of Advanced Computing Technology



iVIC



Thank You!

School of Computer Science, Beihang University
Chunming Hu (hucm@buaa.edu.cn)

Big Data Summit, with CyberC 2013
October 10, 2013. Beijing, China.



Institute of Advanced Computing Technology