

# Trinity: A Distributed Graph Engine on a Memory Cloud

Bin Shao

Microsoft Research (Beijing, China)

**Why do we need a graph system?**

# Existing Systems


- Mature data processing systems
  - RDBMS
  - Map Reduce Systems, e.g. cosmos
- Systems specialized for certain graph operations:
  - PageRank

# Graph Data is “Special” ...

- Random access (Poor Locality)
  - For a node, its adjacent nodes’ content cannot be accessed without “jumping” no matter how you represent a graph
  - Not cache-friendly, data reuse is hard
- Unstructured nature of graph
  - Difficult to extract parallelism by partitioning data
  - Hard to get an efficient “Divide and Conquer” solution


# Graph in the Jail of Storage

- RDBMS/cosmos, mature but not for graphs
- The commonest graph operation “traversal” incurs excessive amount of table joins

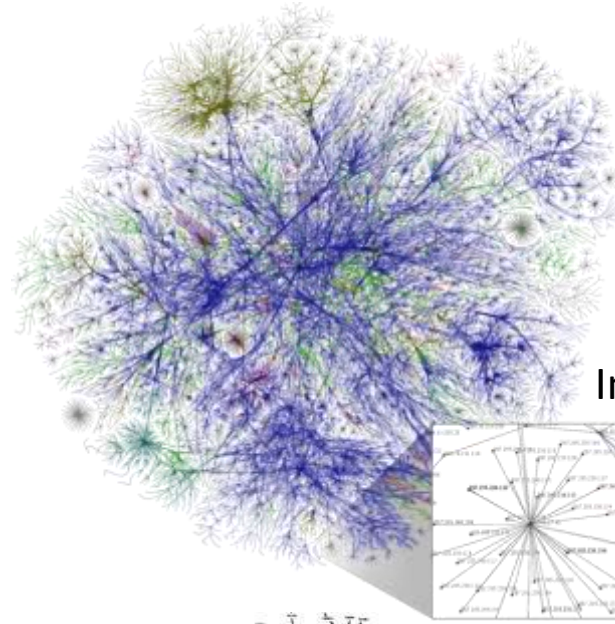


Graph in the  
Jail of the storage

# Challenge I: Diversity of Graphs




Satori Schema Graph



Internet Web Graph



Social Network



Protein Interaction Network

Do we need to design algorithms for each type of graphs?


# Challenge II: Diversity of Computations

- Online query processing
  - Shortest path query
  - Subgraph matching query
  - SPARQL query
  - ...
- Offline graph analytics
  - PageRank
  - Community detection
  - ...
- Other graph operations
  - Graph generation, visualization, interactive exploration, etc.

Do we need to implement systems for each graph operation?

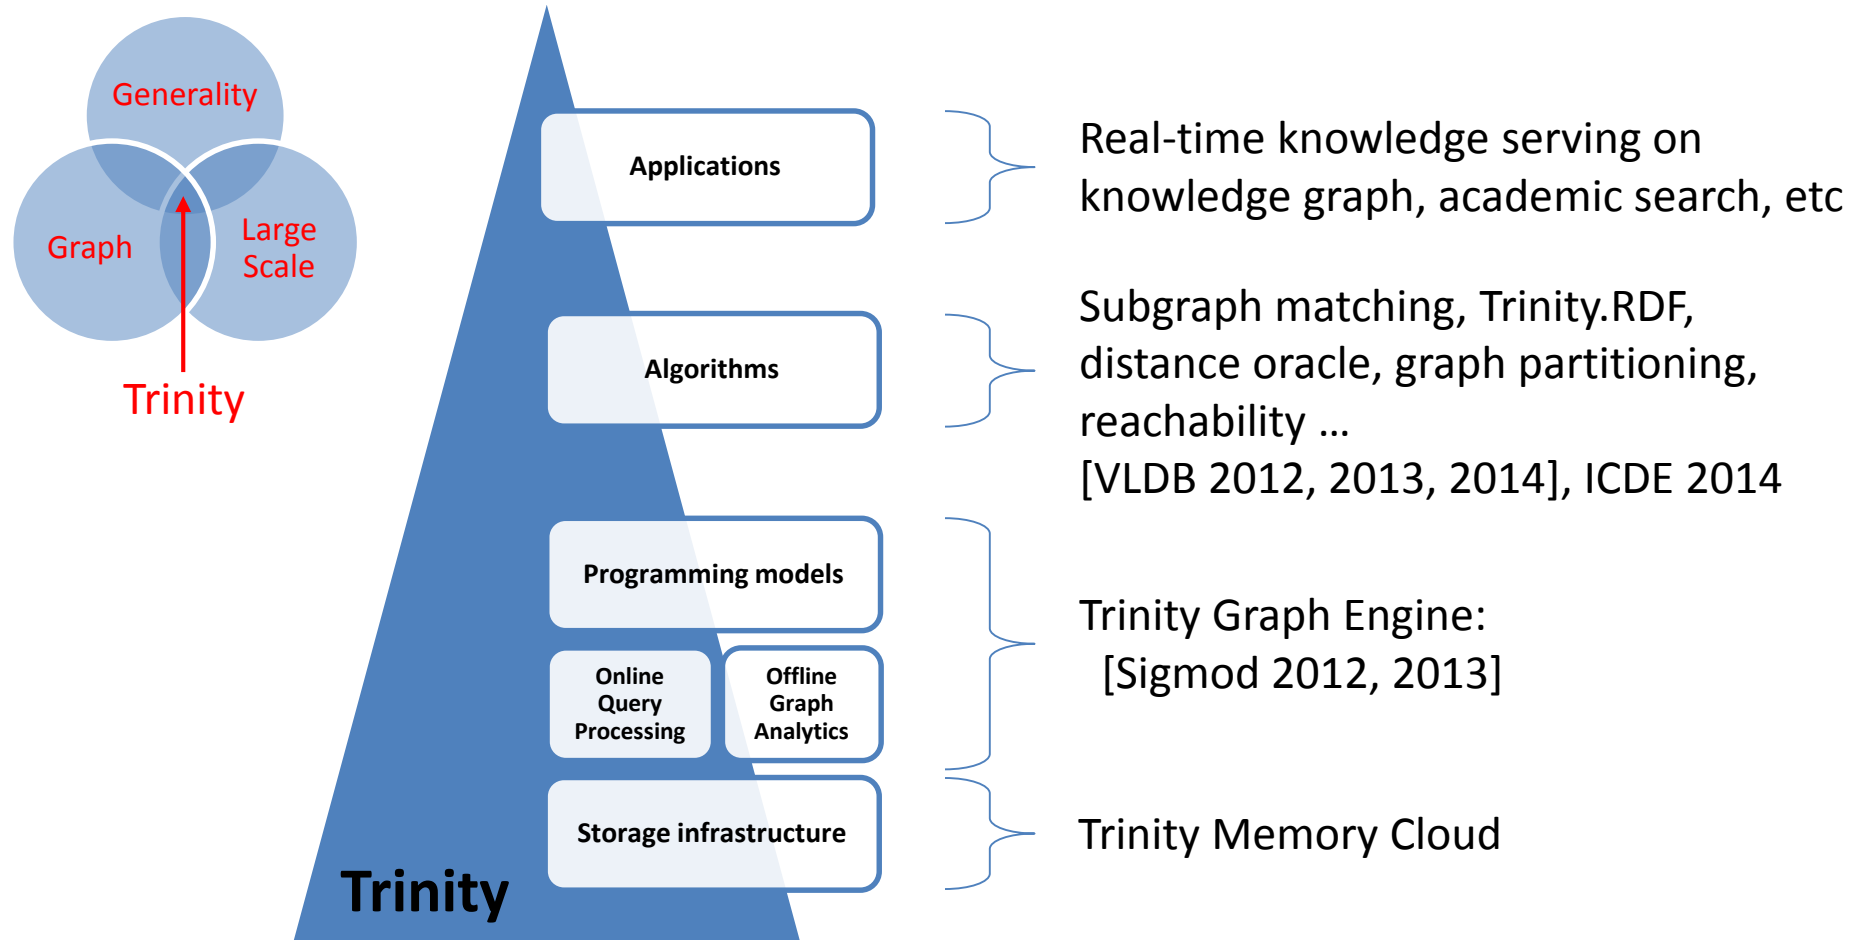
# Challenge III: The **Scale** of Graphs

Makes most graph algorithms in textbooks ineffective!





# Trinity Research Roadmap



# Design Philosophy

**Not a one-size-fits-all graph system, but a graph engine**

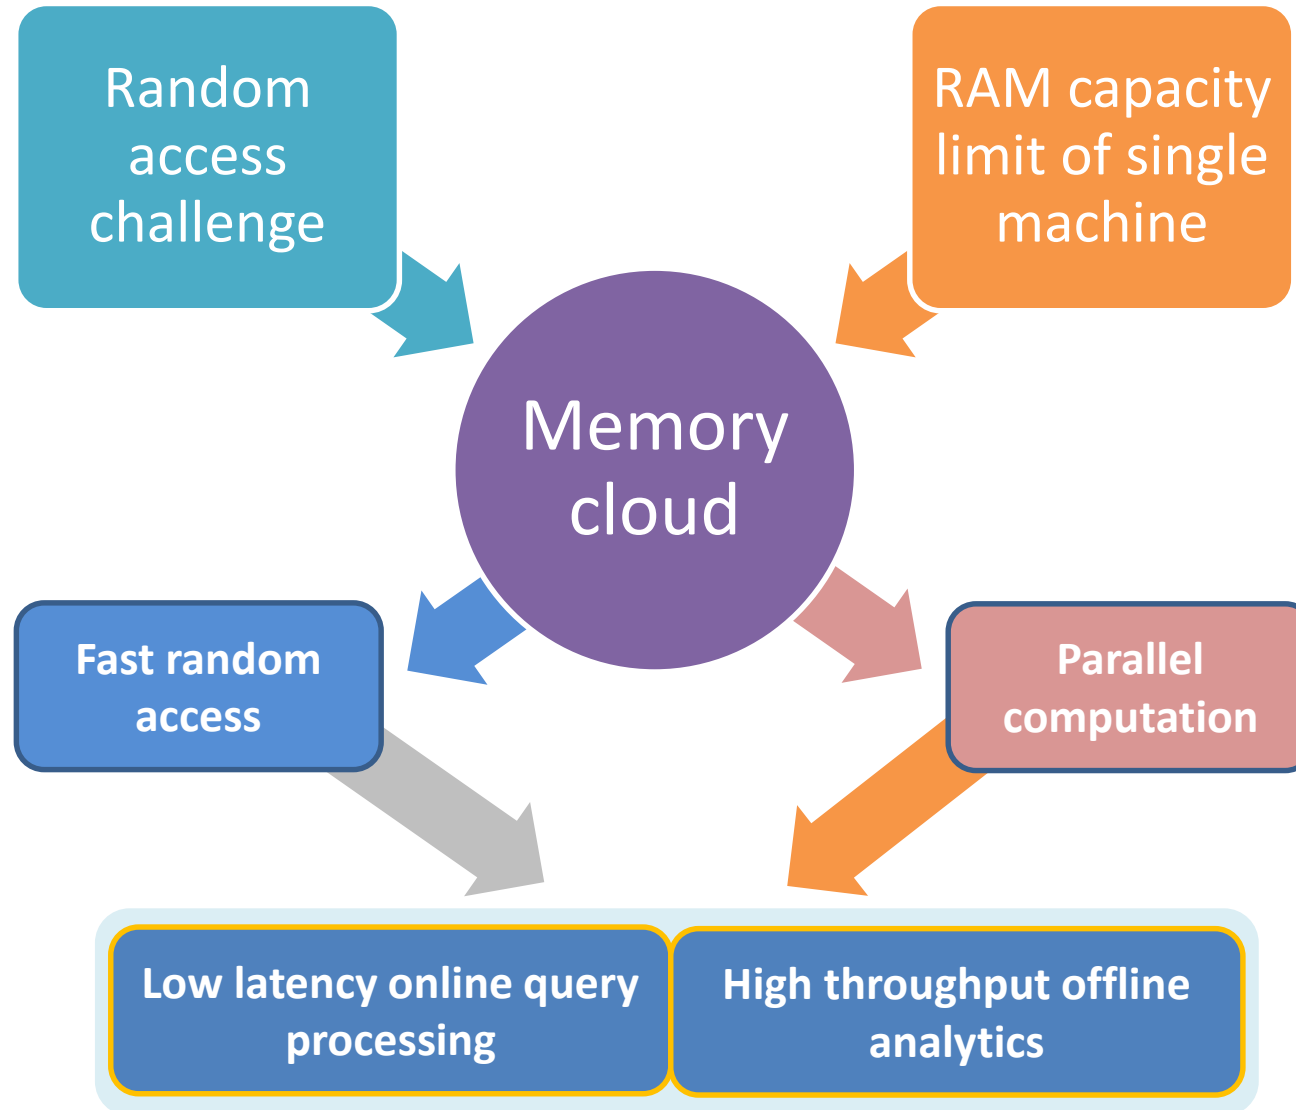
Flexible data and computation modeling capability



Trinity can morph into  
a large variety of graph processing systems

***Trinity* = Graph Modeling Tools +  
Distributed In-memory Data Store +  
Declarative Programming Model**

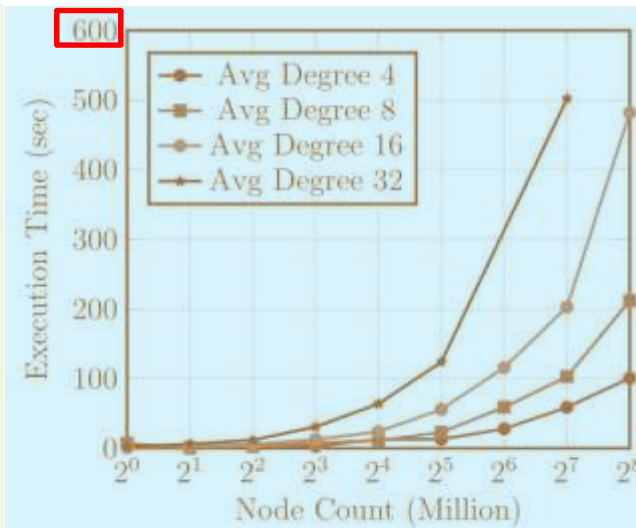

# Design Rationale of Memory Cloud





# System Stack

Graph APIs GetInlinks(), Outlinks.Foreach(...), etc	
Graph Model	
Trinity Specification Language	
Memory Cloud (Distributed Key-Value Store)	
Distributed Memory Storage	Message Passing Framework

# One Byte Counts (Trinity vs. PBGL)



Execution Time

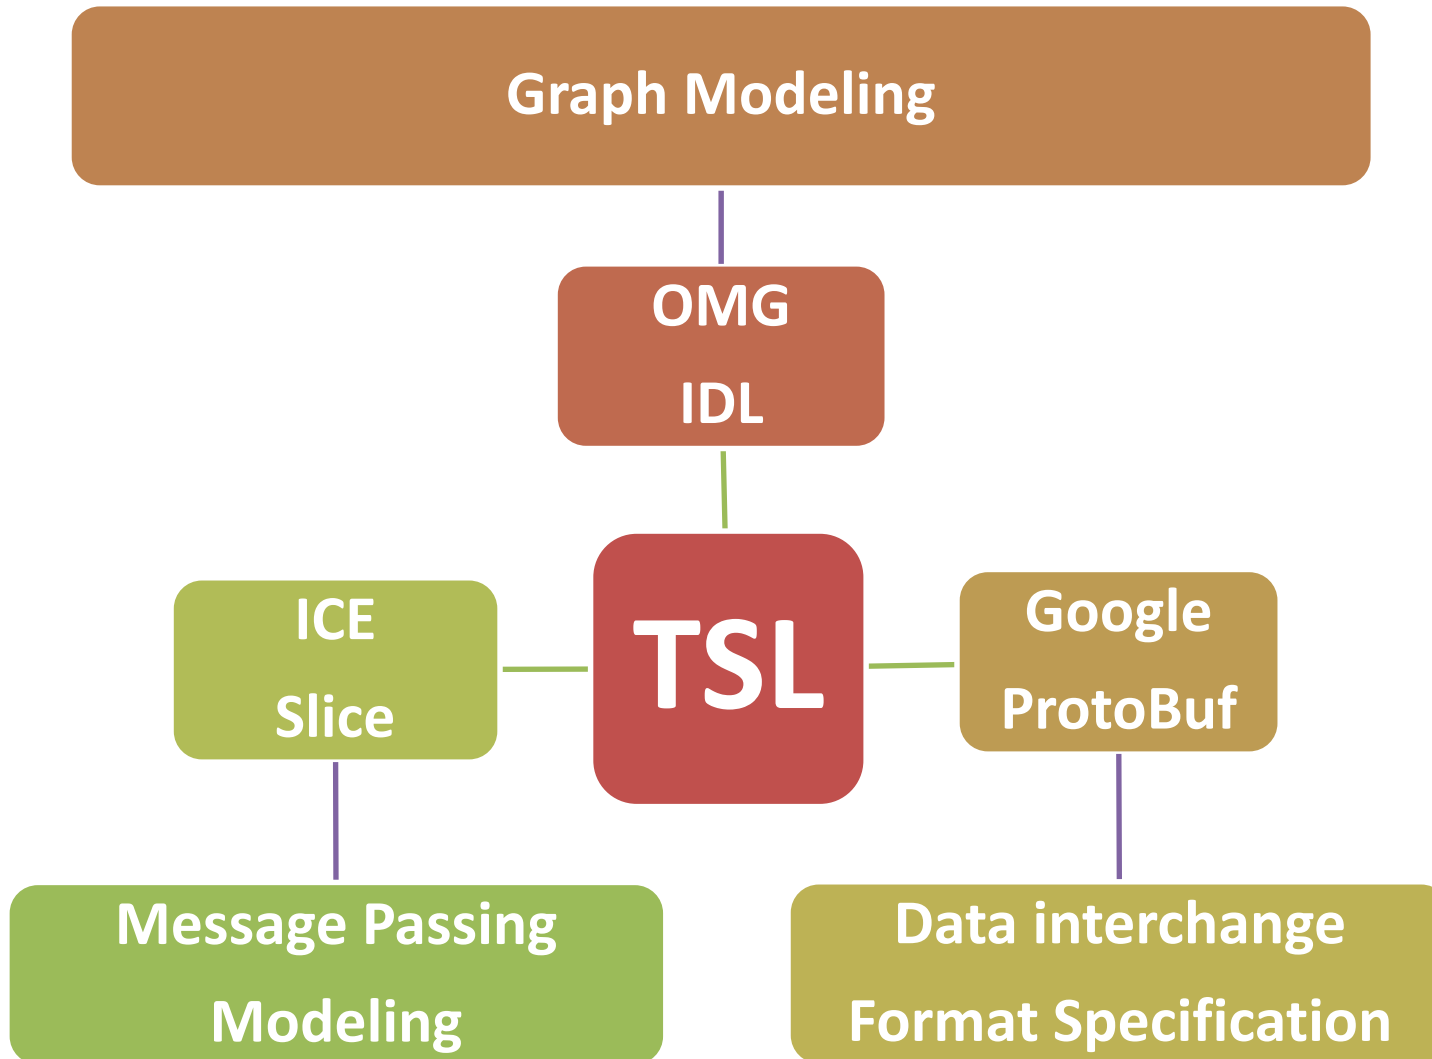


Memory Usage

PBGL

Trinity

# Trinity Specification Language



# Why TSL?

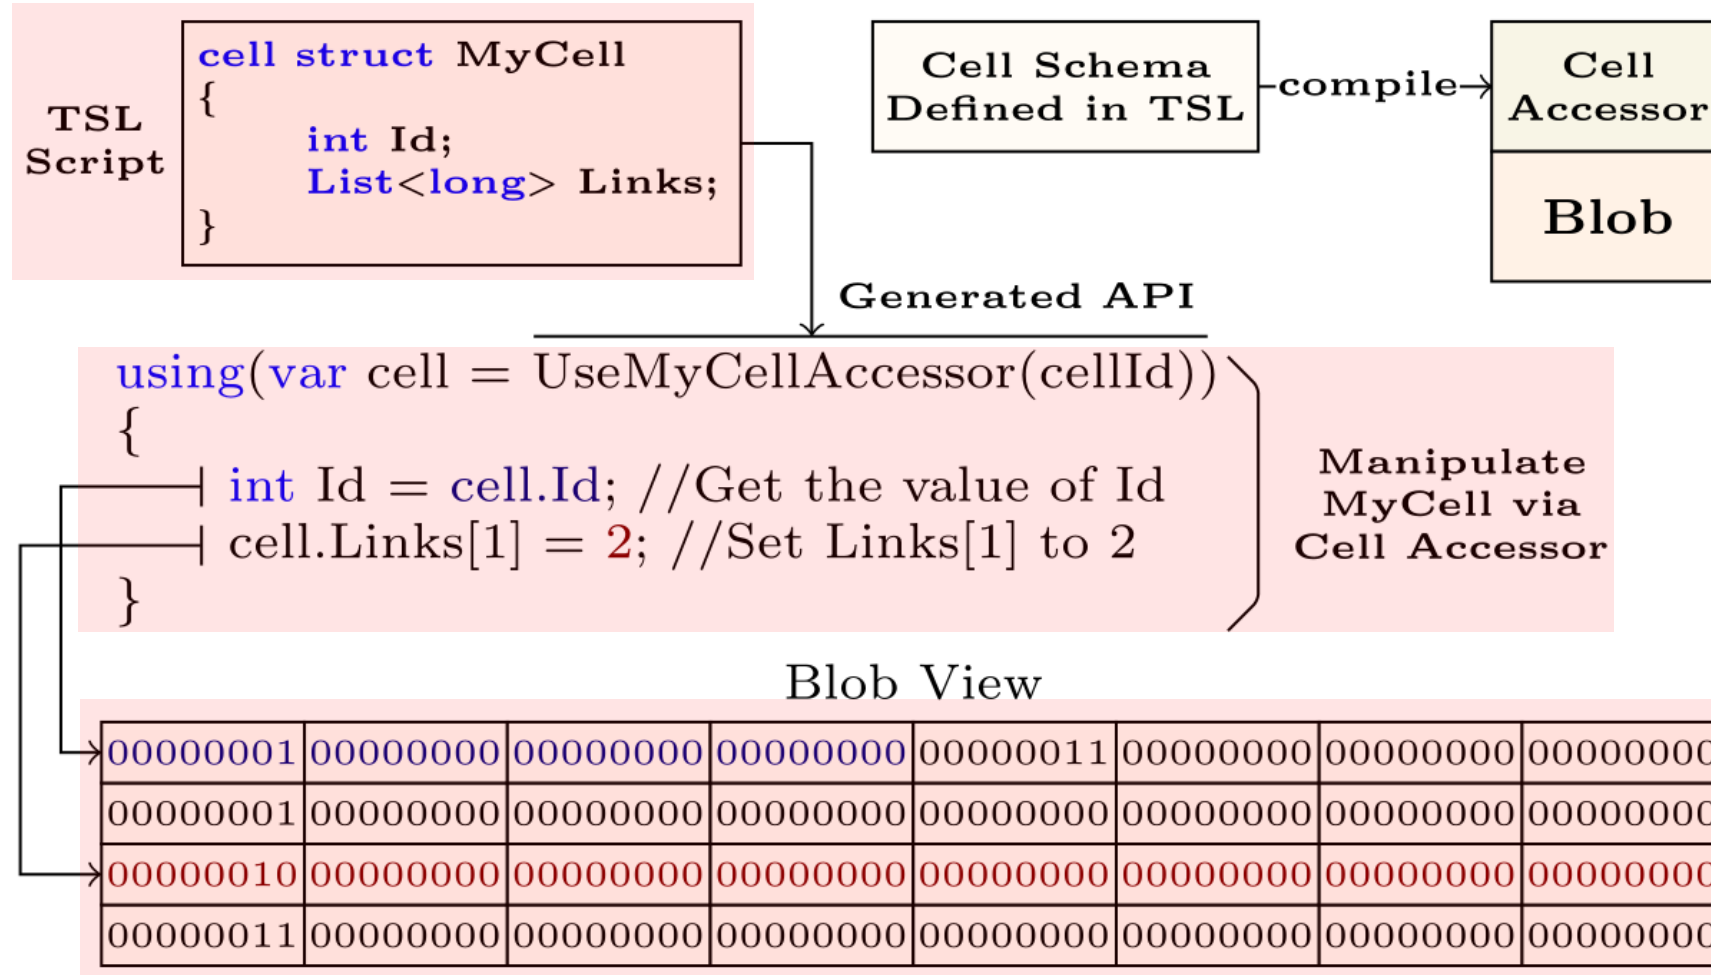
- TSL allows users to define graph schemata, and communication protocols through declarative interfaces.
- TSL makes Trinity memory cloud beyond a key-value store
  - Users are allowed to freely define the data schema
  - TSL makes message passing programming ever so easy

# Modeling a Movie and Actor Graph

```
[CellType: NodeCell]
cell struct Movie
{
    string Name;
    [EdgeType: SimpleEdge, ReferencedCell: Actor]
    List<long> Actors;
}
[CellType: NodeCell]
cell struct Actor
{
    string Name;
    [EdgeType: SimpleEdge, ReferencedCell: Movie]
    List<long> Movies;
}
```



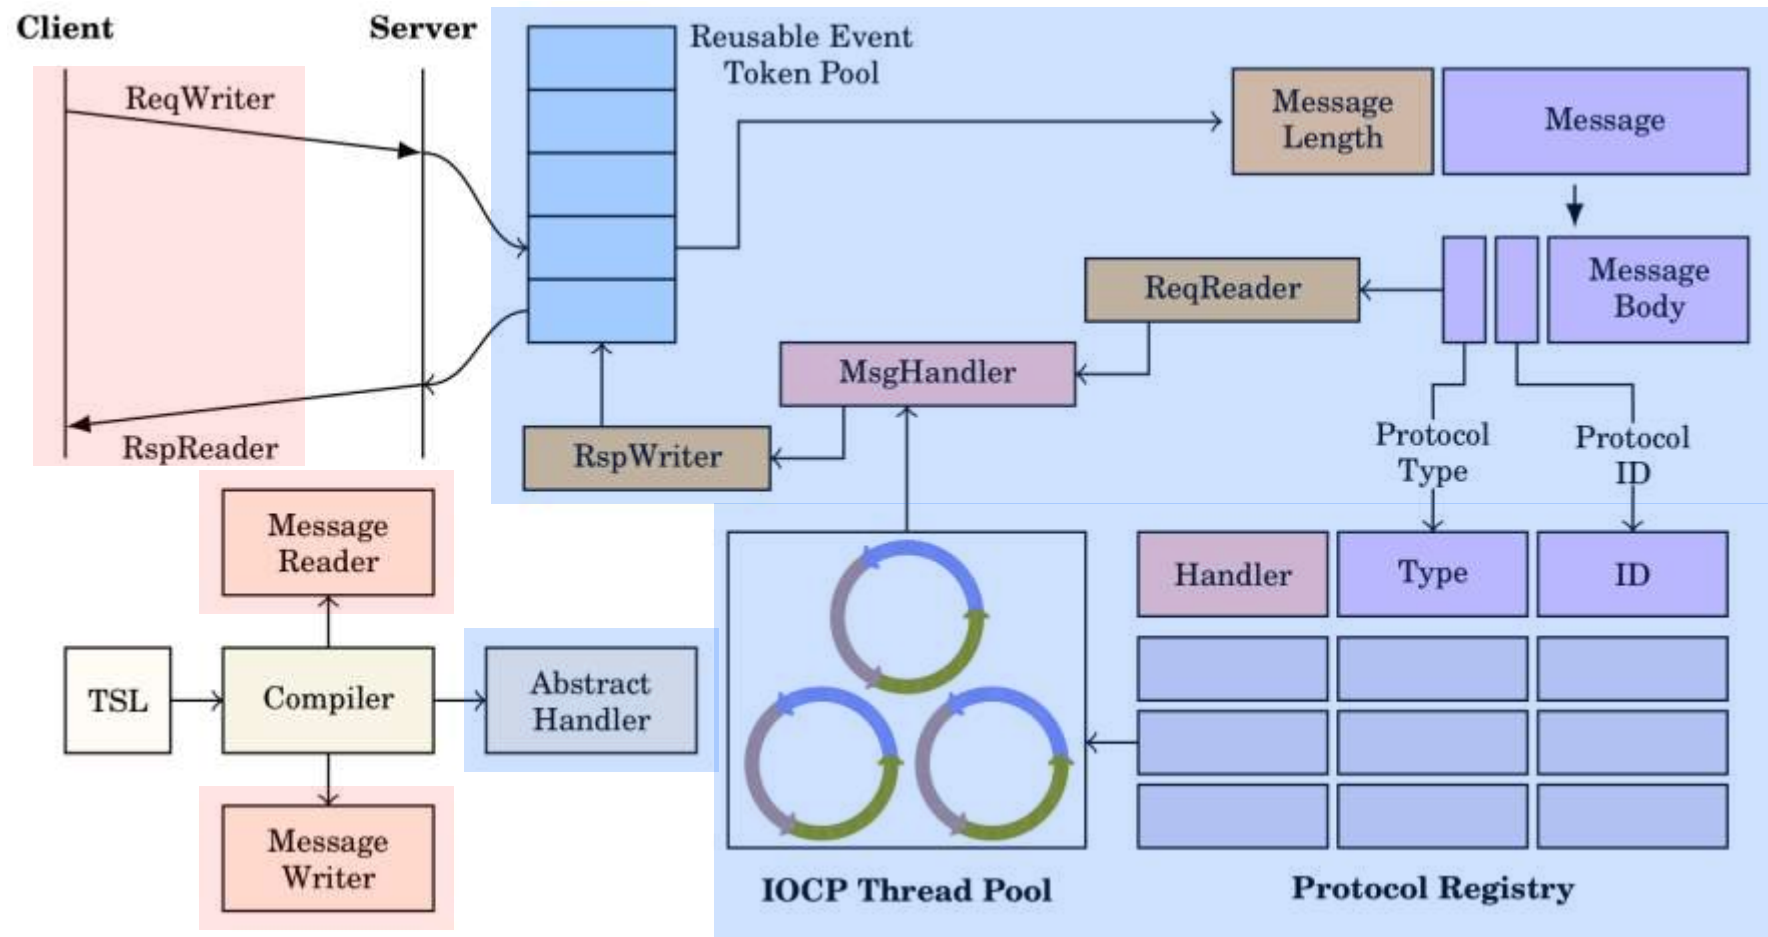
# TSL-enabled Cell Accessor: Efficient and User-friendly



# Modeling Message Passing

```
struct MyMessage
{
    string Text;
}
protocol Echo
{
    Type: Syn;
    Request: MyMessage;
    Response: MyMessage;
}
```

# TSL-Powered Message Passing

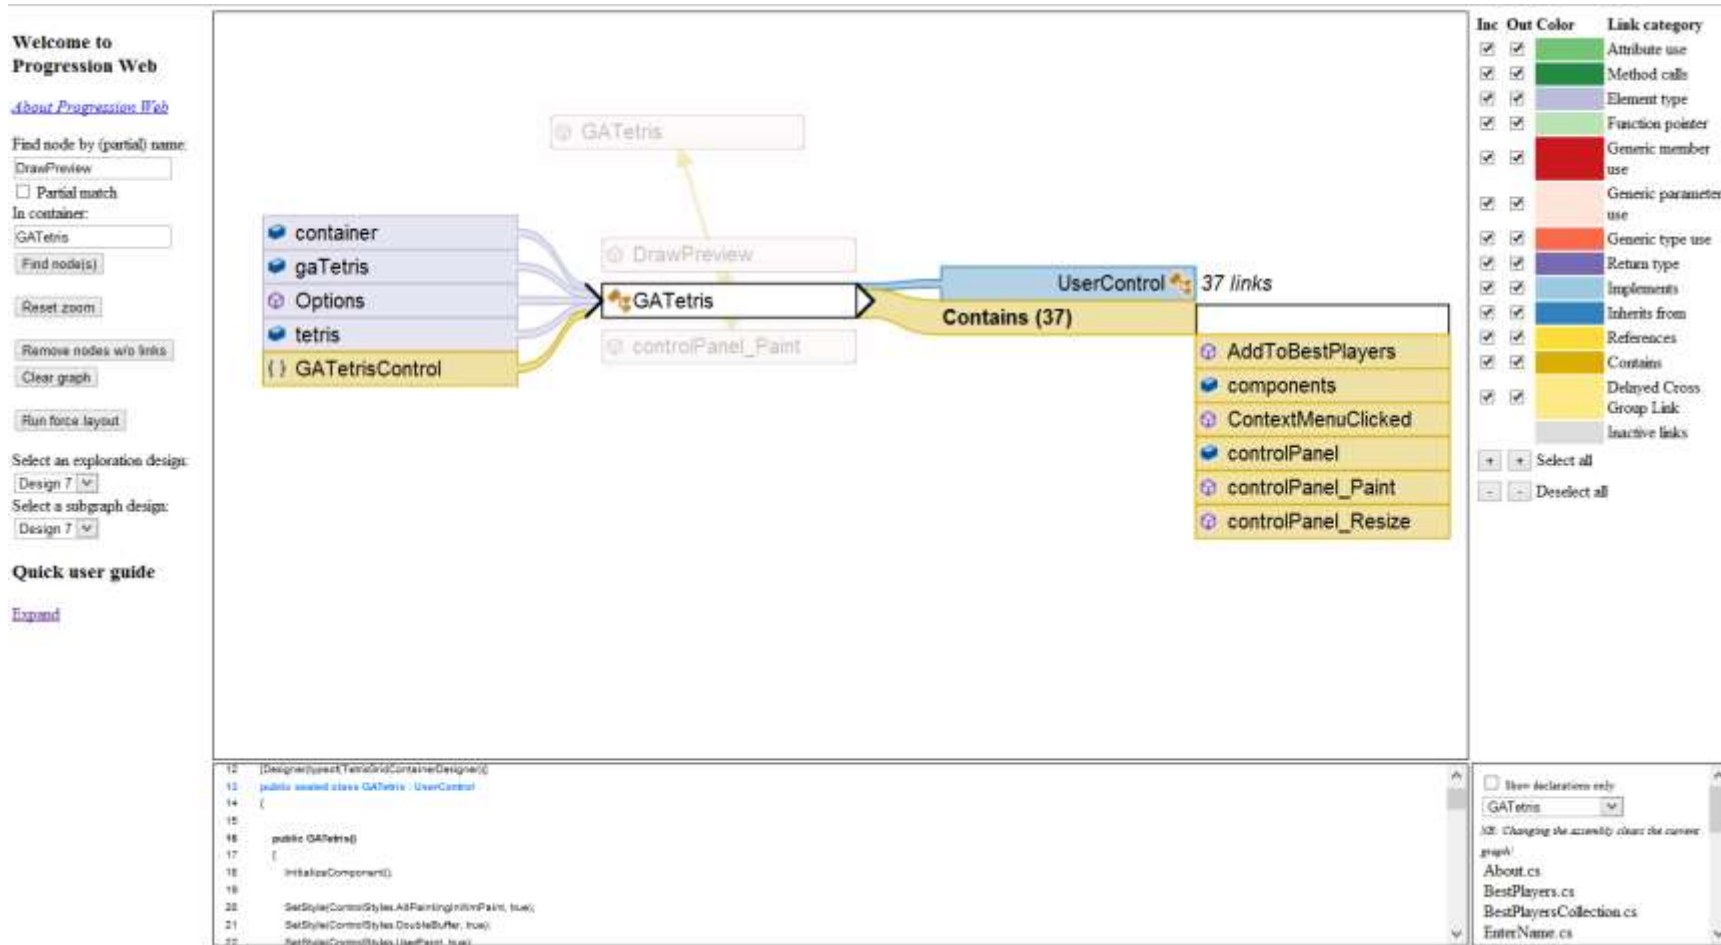


# Trinity-enabled Graph Computation Paradigms

- **Vertex-centric graph analytics**
  - Prosperous since Pregel, e.g. Giraph, GraphChi
- **Approximate graph computation based on local sampling**
  - Enabled by randomly partitioned in-memory graph
  - Fast approximate computation with minimum communication costs
  - Application: distance oracle [VLDB 2014]
- **Index-free real-time online query processing**
  - Enabled by fast in-memory distributed graph exploration
  - Examples, subgraph match (vldb 2012) and Trinity.RDF (vldb 2013)

# Trinity Applications

# Source Code Graph (Visual Studio)



# ACADEMIC SERACH

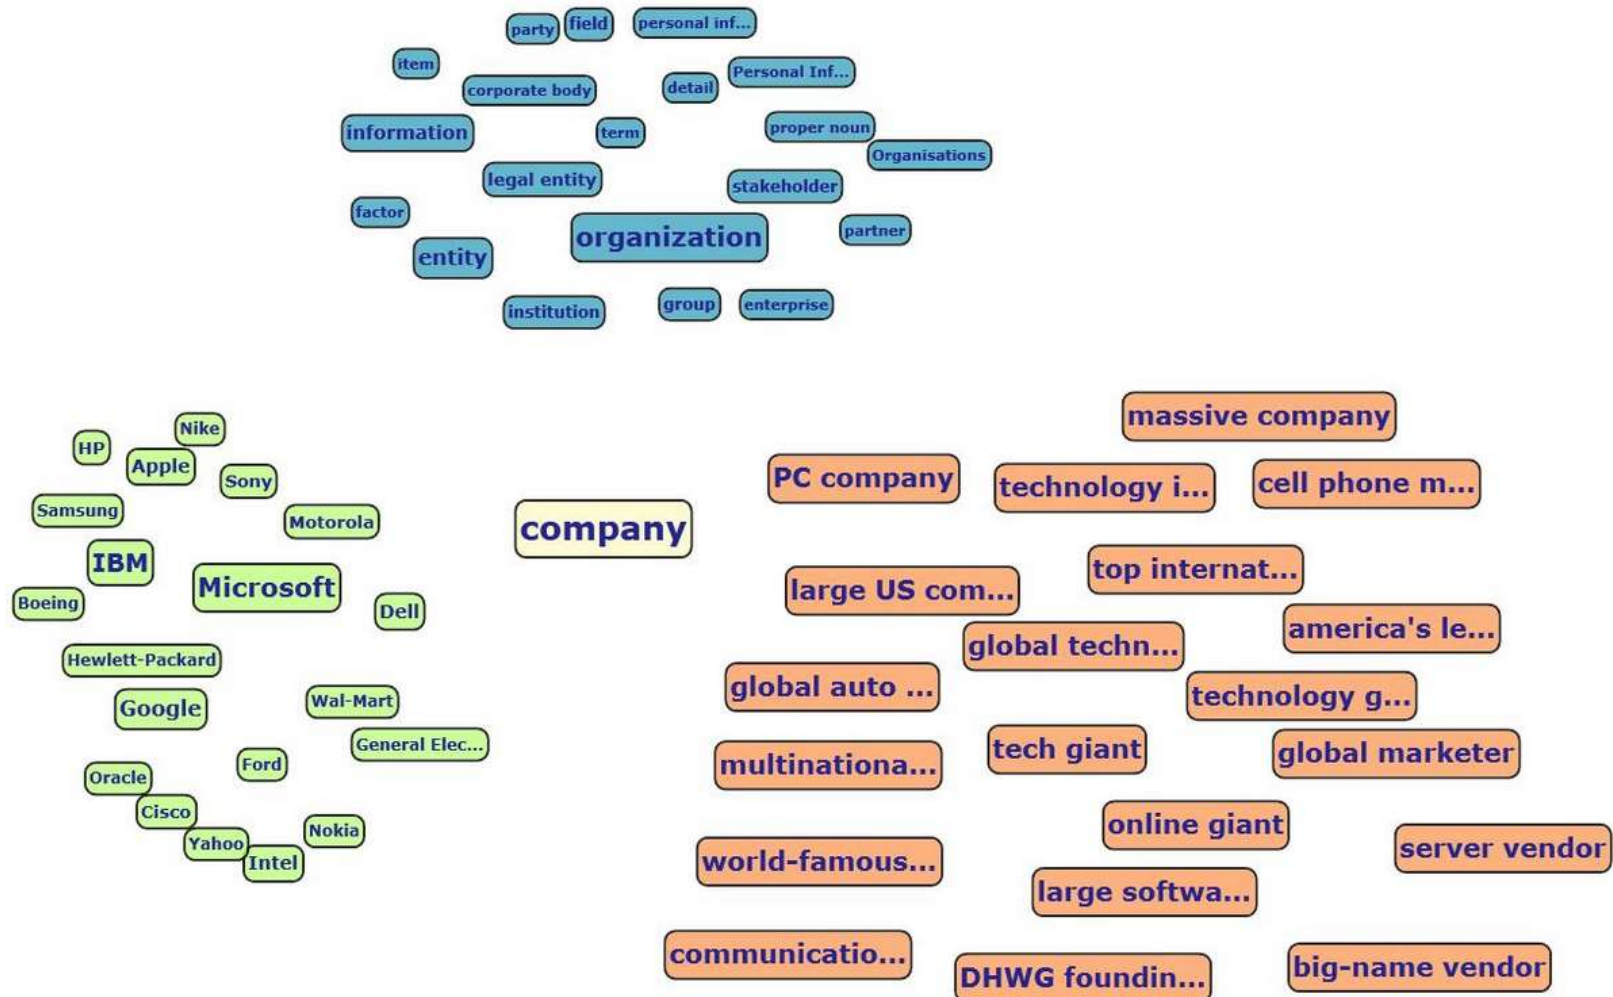
[Demo](#)[Example](#)

```
FROM a in {"Author.FullName='Leslie Lamport'"} MATCH a-->b(PaperAuthorOrganization)-->c(Paper) SELECT a.FullName,c.Title
```


Query Result

a.FullName	c.Title
Leslie Lamport	Composition: A Way to Make Proofs Harder
Leslie Lamport	A Formal Basis for the Specification of Concurrent Systems
Leslie Lamport	The Operators of TLAC
Leslie Lamport	The Synchronization of Independent Processes
Leslie Lamport	Corrigendum: "A New Approach to Proving the Correctness of Multiprocess Programs"
Leslie Lamport	Comment on Bell's quadratic quotient method for hash coded searching
Leslie Lamport	SIFT: Design and analysis of a fault-tolerant computer for aircraft control
Leslie Lamport	Latex: a document preparation system
Leslie Lamport	Constructing digital signatures from a one-v~ray function
Leslie Lamport	Specifying
1 2 3 4 5 6 7 8 9 10 ... >>	

# Knowledge Graph







# Thanks!

[\*http://research.microsoft.com/trinity/\*](http://research.microsoft.com/trinity/)